

Arithmetic landscapes I: the gap series

Kentaro Amauchi*

August 15, 2026

Abstract

Let A be a finite set of odd positive integers, σ the sum, and

$$r_A(n) = \#\{S \subseteq A : \sigma(S) = n\}$$

the number of representations of n as a subset sum of A . For $d \geq 0$ split A at $2d$: let $I_d = \{a \in A : a \leq 2d\}$ be the small elements and $B_d = A \setminus I_d$ the rest. This paper is about the quantity

$$\text{lm}_A(n) = r_A(n) + \sum_{d \geq 1} \left(r_{B_d}(n+d) + r_{B_d}(n-d-\sigma(I_d)) \right),$$

a weighted count of representations by the *truncations* of A . The sum is finite: every term vanishes once $2d \geq \max A$.

We show that the ratio $\text{lm}_A(n)/r_A(n)$ converges, *for every target n* , to a quantity that can be read off A by counting, with no subset sum computed: writing $M = \max A$ and $N_A(d) = \#\{a \in A : a \leq 2d\}$,

$$\Gamma(A) = 1 + 2 \sum_{d=1}^{(M-1)/2} 2^{-N_A(d)} \in \mathbb{Q},$$

equivalently a dyadically weighted sum of the gaps of A , dominated by its smallest elements. For the odd primes $\Gamma = 5.3492879\dots$, and the small initial gaps — the twin pairs — are what make it small.

The quantity $\text{lm}_A(n)$ has a second description, which is where the question came from and which the notation records. Give the subsets S the energy $|\sigma(S) - n|$ and call two subsets adjacent when they differ in one element; then $\text{lm}_A(n)$ counts the S from which every single-element change increases the energy, that is, the strict local minima of the subset-sum landscape, and $r_A(n)$ counts its ground states. That the two descriptions agree is an identity between integers, not an approximation, and it is one of the theorems below. A reader who wants no landscape may take the displayed formula as the definition and read the rest as a statement about representation counts of truncations.

We prove a classification theorem for those local minima; a window identity $W_D(A) = \Gamma(A) + (2D+1-M)2^{-k}$ exhibiting Γ as the total measure of the windows in which a configuration can be a local minimum; sharp bounds $3 - 2^{1-(M-1)/2} \leq \Gamma(A) \leq M$, with equality on the left exactly for $A = \{1, 3, \dots, M\}$ and on the right exactly for $A = \{M\}$; and an exact stratification — the identity displayed above, proved as an equality of integers. If the representation counts r_{B_d} are *flat* to within a factor $1 + \varepsilon$ on an explicit finite window, then

$$\frac{W_D(A)}{1 + \varepsilon} \leq \frac{\text{lm}_A(n)}{r_A(n)} \leq (1 + \varepsilon) W_D(A).$$

*Independent researcher. amkn.sub03@gmail.com, <https://mathlab0911.github.io/>, ORCID <https://orcid.org/0009-0000-0890-4395>

This replaces the annealed (independence) approximation standard in the statistical-mechanics literature by a finite, checkable hypothesis.

We complement it with a lower bound. Writing $F_B(z) = \prod_{a \in B} (1 + z^a)$ for the generating polynomial of the representation counts, we show that $|F_B(i)| = 2^{|B|/2}$ exactly for *every* odd sequence B — the modulus-4 character never degenerates, whereas the modulus-6 one vanishes as soon as $3 \in B$. Consequently the representation counts are never equidistributed modulo 4: the spread of the residue-class totals is determined exactly by the parity of $|B|$, and the relative spread is always at least $2\sqrt{2} \cdot 2^{-|B|/2}$. We also determine which congruence is responsible: the geometric mean $M(q)$ of $|\cos(\pi r/q)|$ over r coprime to q equals $\frac{1}{2}|\Phi_q(-1)|^{1/\varphi(q)}$, and $\sup_{q \geq 3} M(q) = \sqrt{3}/2$ is attained at $q = 6$ alone. The primes, confined modulo 6 to the two residues realising this maximum, sit at rate $\sqrt{3}/2 \approx 0.866$ — above the universal floor $1/\sqrt{2} \approx 0.707$, so they are among the *least* flat odd sequences rather than the flattest.

We compute ε exactly for $k \leq 32$: for the odd primes it decays geometrically at rate $\approx \sqrt{3}/2$, matching the modulus-6 prediction to within 1%, while for random odd sequences it decays strictly faster. Over 100 random sequences per $k \in \{8, 12, 16, 18, 20\}$ we find $(\ln r)/\Gamma \rightarrow 1$ with standard deviation 0.0067 at $k = 20$ and $R^2 = 0.999$. A partial-correlation analysis shows that Γ subsumes the coefficient of variation of the weights, the statistic previously proposed for this role by Alyahya and Rowe.

Every theorem in the structural part of the paper — the classification, the window identity, the exact stratification and the sandwich, together with the lemmas they rest on — is formally verified in Lean 4 with Mathlib, as are the character computations behind the modulus-4 obstruction; no proof contains `sorry` and no additional axiom is used. One statement is verified only in part: for the extremality of the modulus 6 the finite check is formalised, while the supremum over all q rests on the classical values of $\Phi_q(-1)$. Each formal statement is named in the text and indexed in the appendix. The development is public at <https://github.com/mathlab0911/arithmetic-landscapes>.

2020 MSC. 11B75 (primary); 05A17, 68Q25, 82B44.

Keywords. gap series, subset sums, number partitioning, local minima, energy landscape, primes, formal verification.

Contents

1	Introduction	3
1.1	The gap series	3
1.2	Terminology, in one table	5
1.3	Subset-sum landscapes	5
1.4	Relation to previous work	8
1.5	Formal verification	8
2	Definitions, and the bridge to computation	9
2.1	The landscape	9
2.2	The bridge	9
3	The classification theorem	9
4	Symmetry and parity	10
5	The window identity	11
5.1	Windows	11
5.2	The identity	11
5.3	Why the enumeration form needs the order	12

6	Extremal values	12
7	Exact stratification	13
7.1	Strata	13
7.2	Strata are representation counts	14
7.3	What has been achieved	14
8	Flatness and the sandwich	15
8.1	The abstract lemma	15
8.2	The sandwich	15
8.3	A universal obstruction to flatness	16
8.4	Flatness measured	17
9	Numerical experiments	19
9.1	Setup	19
9.2	$\ln/r$ converges to Γ	19
9.3	Primes against random	19
9.4	An oscillation	20
10	Comparison with the coefficient of variation	20
11	Open problems	22
A	The Lean development	26

1 Introduction

1.1 The gap series

Throughout, $A = \{a_1 < \cdots < a_k\}$ is a finite set of odd positive integers. Write $M = a_k$ for its largest element and

$$N_A(d) = \#\{i : a_i \leq 2d\}, \quad d \geq 1,$$

for its counting function: the number of elements at or below $2d$. Since the elements are distinct and odd, $N_A(d) \leq d$ for every d , with equality exactly when A contains every odd number below $2d$. This one inequality is responsible for most of what follows.

Definition 1.1 (gap series). The *gap series* of A is

$$\Gamma(A) = 1 + 2 \sum_{d=1}^{(M-1)/2} 2^{-N_A(d)} \in \mathbb{Q}.$$

The sum stops at the last index for which the layer $\{a \in A : a > 2d\}$ is non-empty; beyond it every term is 2^{-k} and carries no information about A . Two rewritings say what the definition is.

Proposition 1.2 (gap form). $\Gamma(A) = a_1 + \sum_{j=1}^{k-1} (a_{j+1} - a_j) 2^{-j}$.

Proof. $N_A(d) = j$ exactly for $(a_j + 1)/2 \leq d \leq (a_{j+1} - 1)/2$, which is $(a_{j+1} - a_j)/2$ values of d , and $N_A(d) = 0$ for $d \leq (a_1 - 1)/2$, which is $(a_1 - 1)/2$ values. Collecting equal terms, $\Gamma(A) = 1 + (a_1 - 1) + \sum_{j=1}^{k-1} (a_{j+1} - a_j) 2^{-j}$. $\square$

Proposition 1.3 (Γ is Schur-concave). *[proved; the proof is Abel summation followed by the Schur–Ostrowski criterion. Measured for confirmation only: over 4000 majorization-comparable pairs of sets of distinct odd integers, Γ was larger on the more equal set in every case that was not a tie; schur_r149.]* Regard Γ as a function of the multiset. If A majorizes B — same sum, A the more spread out — then $\Gamma(A) \leq \Gamma(B)$.

Proof. Abel summation applied to Proposition 1.2 gives $\Gamma(A) = \sum_{j=1}^k w_j a_j$ with $a_1 < \dots < a_k$ and

$$(w_1, \dots, w_k) = (\tfrac{1}{2}, \tfrac{1}{4}, \dots, 2^{-(k-1)}, 2^{-(k-1)}),$$

positive and *non-increasing*. So Γ is an ordered weighted average taken along the ascending order. It is symmetric as a function of the multiset, and on the open region where the entries are distinct $\partial\Gamma/\partial a_i = w_i$; if $a_i > a_j$ then $i > j$, hence $w_i \leq w_j$ and $(a_i - a_j)(\partial_i\Gamma - \partial_j\Gamma) \leq 0$. That is the Schur–Ostrowski condition for Schur-concavity, and Γ is continuous, so it holds on the closure. $\square$

Remark 1.4 (what Schur-concavity buys, and what it explains). Two things, one technical and one conceptual.

Technical. Γ now sits inside the Hardy–Littlewood–Pólya framework, so Karamata’s inequality and the standard majorization toolkit apply to it without further work. Any statement of the form “ Γ is extremal on the most/least spread set with a given sum” is a corollary rather than a computation.

Conceptual. Schur-concavity says exactly that Γ *rewards equality and punishes spread*. That is the precise version of the informal description used throughout this series — that Γ is a *bottom* statistic, decided by the smallest elements — and it is why Γ separates sets that a second moment cannot: variance is Schur-convex, so it moves the other way. The two are not competing measures of the same thing; they have opposite monotonicity under the same partial order. *[derived from Proposition 1.3; the remark about variance is the standard fact that $\sum(a_i - \bar{a})^2$ is Schur-convex.]*

Proposition 1.5 (enumeration form). $\Gamma(A) = \sum_{j=1}^k \frac{a_j}{2^j} + \frac{a_k}{2^k}$.

Proof. Abel summation applied to Proposition 1.2. $\square$

[Proposition 1.5 is Theorem 5.1 at $D = (M - 1)/2$, and is Lean-verified in that form (windowSeries_eq_gapSeries); Proposition 1.2 then follows from it by the same summation by parts. The bounds of §6 are verified as well.]

Three remarks fix the character of the quantity.

First, Γ is an **invariant of the set**. Definition 1.1 never mentions an ordering: N_A counts. This is what one wants of a quantity that is to describe the subset-sum landscape of A , since that landscape depends on the multiset alone. It also distinguishes Γ from the statistics ordinarily attached to a set of weights — mean, variance, coefficient of variation — not by being order-sensitive, but by being sensitive to *where the elements sit* rather than to their spread.

Second, the name is now exact: by Proposition 1.2, Γ is a dyadically weighted sum of the *gaps* of A , dominated by the earliest ones. For the odd primes 3, 5, 7, 11, ... the small initial gaps — in particular the twin pairs (3, 5), (5, 7), (11, 13), (17, 19) — force Γ to be small: $\Gamma(3, 5, \dots, 73) = 1402281/262144 \approx 5.34928$, against ≈ 8.09 on average for a random odd sequence of the same length drawn from the same range.

Third, by Proposition 1.5 the enumeration series $\sum_j a_j 2^{-j}$ — which does depend on an ordering, and which is how this invariant was first written — agrees with Γ up to $a_k 2^{-k}$, so

the two have the same limit along any sequence with $a_k 2^{-k} \rightarrow 0$. The choice between them is therefore a choice of finite representative, and Definition 1.1 is the better one twice over: it is manifestly a set invariant, and it is the closer approximant — for the odd primes at $k = 20$ it lies $1.04 \cdot 10^{-5}$ from the limit against $8.01 \cdot 10^{-5}$ for the enumeration form.¹

The purpose of this paper is to show that this innocuous-looking quantity controls a genuinely combinatorial object: the local structure of the subset-sum problem for A .

Remark 1.6 (Terminology). The name is meant to record the second formula above: Γ is a dyadically weighted series in the gaps of A . It should not be confused with the *gap-sum sequence* of an integer sequence in the sense of Barry [Ba21], which is formed from the integers *absent* between consecutive terms; that construction and ours share no content beyond the word “gap”. Nor is it related to the *gap series* of complex analysis — a power series whose exponents have Hadamard gaps, also called a lacunary series — where the gaps are in the exponents rather than in the coefficients, and the object of study is continuation across the circle of convergence. Readers who prefer an unambiguous name may read Γ as the *dyadic gap series*, or simply as the generating function of A evaluated at $x = 1/2$.

1.2 Terminology, in one table

This programme names a few objects that additive number theory has no name for, and borrows a few names from statistical physics for objects it does. A reader who meets one of them cold can look it up here; each is defined again where it is used.

Two of these are the point of the paper and the rest are scaffolding. The gap series is a new object; the identity between the valley count and the representation counts of the truncations is a new theorem. Everything else is vocabulary for stating them.

1.3 Subset-sum landscapes

Fix a target $n \in \mathbb{N}$. For $S \subseteq \{1, \dots, k\}$ put $\sigma(S) = \sum_{i \in S} a_i$ and $E(S) = |\sigma(S) - n|$. Two configurations are adjacent if they differ in a single index. This is the standard 1-bit-flip landscape of the number partitioning problem (NPP); the case $n = \sigma(A)/2$ is the balanced partition problem. We study

$$r_A(n) = \#\{S : E(S) = 0\}, \quad \text{lm}_A(n) = \#\{S : E(S) < E(S') \text{ for all neighbours } S'\}.$$

Here lm counts the *valleys* of the landscape and is the standard measure of its ruggedness.

Our main results are the following; full statements appear in the sections indicated.

Theorem A (Classification; §3). Suppose all $a_i > 0$ and let $d = E(S)$. If $\sigma(S) \geq n$, then S is a strict local minimum if and only if $a_i > 2d$ for every $i \in S$. If $\sigma(S) \leq n$, then S is a strict local minimum if and only if $a_i > 2d$ for every $i \notin S$.

Thus membership in the set of local minima is decided by a condition involving only the elements of size at most $2d$; everything larger is irrelevant. This is the engine of the whole paper.

Theorem B (Window identity; §5). Let A be non-decreasing with all a_i odd and $a_i \leq 2D + 1$. Put $N_A(d) = \#\{i : a_i \leq 2d\}$ and $W_D(A) = 1 + 2 \sum_{d=1}^D 2^{-N_A(d)}$. Then

$$W_D(A) = \Gamma(A) + \frac{2D + 1}{2^k}.$$

¹An earlier circulated version of this paper defined Γ by the enumeration series. The two differ by $a_k 2^{-k}$ and have the same limit; every asymptotic statement in this series of papers is unaffected. That version also carried the title *The gap series of an integer sequence*; the present title places it in the series, and no content was removed in the change.

term	what it denotes, in ordinary language	
representation count $r_A(n)$	$\#\{S \subseteq A : \sigma(S) = n\}$. Standard. Every count of this kind in this series carries the letter r .	
truncation B_d	$\{a \in A : a > 2d\}$, the set with its elements up to $2d$ removed; I_d is the removed part.	
gap series $\Gamma(A)$	the rational number $1 + 2 \sum_d 2^{-\#\{a \in A : a \leq 2d\}}$, read off A by counting. Equivalently a dyadically weighted sum of the gaps of A ; hence the name.	
subset-sum landscape	the set of subsets of A with the function $S \mapsto \sigma(S) - n $ and single-element changes as moves. A way of speaking, not an object: every statement about it in this paper is restated as a statement about representation counts.	
energy, ground state	$ \sigma(S) - n $, and a subset with energy 0. Borrowed from statistical physics, where this landscape is the number partitioning problem.	
valley, $\text{lm}_A(n)$	a subset from which every single-element change increases the energy, and the number of them. Equal, as an identity between integers, to a weighted sum of representation counts of the truncations (Theorem 7.3).	
window series $W_D(A)$	$1 + 2 \sum_{d \leq D} 2^{-\#\{a \in A : a \leq 2d\}}$, the same sum stopped at D . Equals Γ at the last informative D . The <i>window identity</i> is that equality, Theorem 5.1; earlier versions also called W_D the <i>window measure</i> , which is the same object.	
stratum	the set of valleys at a fixed distance d from the target. The <i>stratification</i> is the decomposition of $\text{lm}_A(n)$ into strata.	
flatness ε_d	the relative oscillation $\max r_{B_d}(m)/r_{B_d}(m') - 1 $ of a representation count over a bounded window of targets. Small ε_d is what makes the strata comparable.	
ripple	the periodic oscillation that a sub-peak of $ F_A $ contributes to the representation counts as the target varies. Its size is what ε_d measures; Part II computes it for the primes.	
the sandwich	the two-sided bound of Theorem 8.3: if ε_d is small then $\text{lm}_A(n)/r_A(n)$ lies within a factor $1 + \varepsilon$ of $W_D(A)$.	

The quantity W_D is the total measure of the “windows” in which a configuration at error d can be a local minimum, computed as if the elements were placed independently. Theorem B says that this measure is exactly the gap series, up to a tail term negligible for the sequences of interest ($73/2^{20} \approx 7 \times 10^{-5}$ for the first twenty odd primes).

Theorem C (Exact stratification; §7). Without any approximation,

$$\text{lm}_A(n) = r_A(n) + \sum_{d \geq 1} (\text{over}_A(n, d) + \text{under}_A(n, d)),$$

where the strata vanish once $2d$ exceeds $\max_i a_i$, and each stratum equals a representation count of a truncated subsequence,

$$\text{over}_A(n, d) = r_{B_d}(n + d), \quad \text{under}_A(n, d) = r_{B_d}(n - d - \sigma(I_d)),$$

with $I_d = \{i : a_i \leq 2d\}$ and $B_d = A|_{I_d^c}$. Moreover $r_A(n) = \sum_{J \subseteq I_d} r_{B_d}(n - \sigma(J))$.

Theorem D (Sandwich; §8). Suppose that for every $d \leq D$ the representation counts r_{B_d} vary by at most a factor $1 + \varepsilon$ across the window $\{n - \sigma(J) : J \subseteq I_d\} \cup \{n + d, n - d - \sigma(I_d)\}$. Then

$$\frac{W_D(A)}{1 + \varepsilon} \leq \frac{\text{lm}_A(n)}{r_A(n)} \leq (1 + \varepsilon) W_D(A).$$

Combining Theorems B and D, lm/r is pinned to $\Gamma(A)$ with a multiplicative error controlled by a single, explicitly computable number ε .

The flatness ε cannot be made to vanish arbitrarily fast. Write $F_B(z) = \prod_{a \in B} (1 + z^a)$, so the representation counts r_B are the coefficients of F_B , and let $R_c = \sum_{m \equiv c(4)} r_B(m)$.

Theorem E (Universal obstruction; §8.3). For every finite sequence B of odd positive integers, $|F_B(i)| = 2^{|B|/2}$ exactly, whatever the residues of the elements. Consequently the counts R_c are never all equal; their spread is determined by the parity of $|B|$,

$$\max_c R_c - \min_c R_c = 2^{|B|/2} \quad (|B| \text{ even}), \quad 2^{(|B|-1)/2} \quad (|B| \text{ odd}),$$

so the relative spread is always at least $2\sqrt{2} \cdot 2^{-|B|/2}$.

This is a lower bound, in contrast with Theorems B–D, and it is elementary: $1 + i^a = 1 \pm i$ has modulus $\sqrt{2}$ for *every* odd a , so the modulus-4 character never degenerates. The modulus-6 character does: if $3 \in B$ then $F_B(\zeta_6) = 0$.

The obstruction of Theorem E is a statement about the distribution of r_B across residue classes modulo 4, not about the pointwise flatness ε_d of Theorem D; we do not derive a lower bound for the latter, and §8.3 explains why the transfer needs an input we do not have. Numerically, however, the two sit at the same scale for random odd sequences, whereas for the primes ε_2 exceeds the modulus-4 scale by two orders of magnitude at $k = 32$: their obstruction decays at rate $\sqrt{3}/2 \approx 0.866$, against $1/\sqrt{2} \approx 0.707$. In that sense the primes are not the flattest odd sequences but among the least flat.

We emphasise what Theorem D does *not* assume. The statistical-mechanics treatment of the NPP proceeds through the annealed approximation, in which configurations are treated as independent. Theorems C and D dispense with this: the decomposition is an identity, and the only input is the flatness of an explicit arithmetic function on an explicit finite window. Flatness is not an idealisation — it is a quantity one can compute for any given A , and we do so in §8.4.

1.4 Relation to previous work

The NPP has been studied intensively from the statistical-mechanics side. Ferreira and Fontanari [FF98] analysed the problem for i.i.d. weights uniform on $[0, 1]$, identifying it with the ground-state problem for an infinite-range antiferromagnetic Ising model, and computed within the annealed approximation the fraction of metastable states — configurations stable under all single spin flips — showing that it vanishes like $N^{-3/2}$. Mertens [Me01] surveys the physicist’s approach; Stadler, Hordijk and Fontanari [SHF03] study landscape statistics across the phase transition; and Borgs, Chayes and Pittel [BCP01] give a rigorous treatment of the transition and its finite-size scaling.

We stress that *the order of magnitude $2^k k^{-3/2}$ for the number of local minima of a random instance is not new here*. It is exactly the content of the annealed computation of [FF98]. Our contribution lies elsewhere, in two directions.

Deterministic sequences. The literature treats the weights as random and asks for typical behaviour. We fix a deterministic sequence — for instance the primes — and ask which of its *arithmetic* features control the landscape. The answer, Theorems B–D, is the gap series, an invariant with no probabilistic content.

Exactness. The annealed approximation is replaced by an identity plus a checkable hypothesis. To our knowledge Theorem C, and the reduction of lm/r to the flatness of truncated representation counts, are new. That sentence is a literature pass and not a proof of novelty: it records the absence of evidence over the searches we made, and the searches we made are ordinary ones. [\[literature pass, not a proof of novelty: absence of evidence over the searches made, at a database of the size and date stated\]](#) The sequence searches were made and are recorded. On 2026-08-11, against a database of 398 208 sequences, none of four queries to the OEIS returned a match: the numerators 3, 11, 29, 69, 151, $\dots$ of $2^k \Gamma(P_k)$; the digits of $\Gamma(P)$; the number of strict local minima of the first k odd primes at the central target; and the number of ground states of the same instances. A positive control run first — the decimal digits of $\sqrt{2}$ — returned A002193 at the head of its results, so the route does find a sequence that is there. The full record, with the exact query strings and the database timestamp, is `lean/pnp/oeis_r119.log`. The database grows; if Γ or any of these is catalogued now, we would be glad to be told.

The closest precursor is Alyahya and Rowe [AR14], who observed empirically that the number of local optima in the 1-bit-flip NPP landscape is strongly *negatively* correlated with the coefficient of variation CV of the weights, and proposed an estimator based on CV . Their observation is correct and we reproduce it (§10). But CV is a global second moment, whereas Theorem B shows it is Γ — a dyadically weighted sum of gaps, dominated by the smallest elements — that equals the window series W_D . The two can be decoupled, and when they are, the landscape follows Γ : pairs of sets with *exactly* equal CV but Γ a factor 2.1 apart have lm/r a factor 2.1 apart, while pairs with exactly equal Γ and CV a factor 1.6 apart have lm/r within 8% (Table 6). In this precise sense Γ subsumes CV .

1.5 Formal verification

Every theorem stated in §§2–8 has been formalised in Lean 4 against Mathlib. Each formal statement is accompanied in the text by the name of the corresponding Lean declaration, and Appendix A lists the files, the declarations and the reproduction procedure. We have verified with `#print axioms` that no declaration depends on `sorryAx` or on any axiom beyond Lean’s standard `propext`, `Classical.choice` and `Quot.sound`.

Two remarks on what this did and did not buy us. It did not merely confirm proofs already written on paper: formalisation of Theorem D revealed that the exponent could be improved from $(1+\varepsilon)^{\pm 2}$ to $(1+\varepsilon)^{\pm 1}$, since numerator and denominator may be controlled by

the same flatness estimate. On the other hand, formalisation forced us to state hypotheses that are easy to leave implicit — most notably the non-decreasing hypothesis in Theorem B, whose necessity we exhibit by a two-element counterexample in §5.3.

2 Definitions, and the bridge to computation

2.1 The landscape

Throughout, A denotes a finite sequence of positive integers, presented either as a list $(a_1, \dots, a_k)$ or as a function $A : \mathbb{F}_k \rightarrow \mathbb{N}$ on the index type $\mathbb{F}_k = \{0, \dots, k-1\}$; §2.2 shows the two presentations agree. We write $\sigma(S) = \sum_{i \in S} a_i$ for $S \subseteq \mathbb{F}_k$, and $T = \sigma(A)$. For a target n ,

$$E_n(S) = |\sigma(S) - n|, \quad \text{flip}(S, i) = \begin{cases} S \setminus \{i\} & i \in S, \\ S \cup \{i\} & i \notin S, \end{cases}$$

and S is a *strict local minimum* if $E_n(S) < E_n(\text{flip}(S, i))$ for all i . We set

$$\text{gs}_A(n) = \min_S E_n(S), \quad r_A(n) = \#\{S : \sigma(S) = n\}, \quad \text{lm}_A(n) = \#\{S : S \text{ strict local min}\}.$$

The middle one is the representation count of n as a subset sum of A ; in the landscape reading it is the number of ground states, and earlier versions of this paper wrote it $\text{deg}_A(n)$ for that reason. We use r_A throughout, so that the same kind of object — a representation count — always carries the same symbol, whether the set is A or one of its truncations B_d .

(Lean: `ndist`, `subsetSums`, `gs`, `energy`, `flip`, `IsStrictLocalMin`, `lm`, `degCount`.)

2.2 The bridge

Implementations enumerate subsets by bitmasks over a list; proofs are more convenient over $\text{Finset}(\mathbb{F}_k)$. That these describe the same object is not automatic, and we record it.

Proposition 2.1. *For $A : \mathbb{F}_k \rightarrow \mathbb{N}$ and $s \in \mathbb{N}$, $s \in \text{subsetSums}(\text{offFn } A)$ if and only if there exists $S \subseteq \mathbb{F}_k$ with $\sigma(S) = s$. (Lean: `mem_subsetSums_offFn_iff`.)*

Proposition 2.2. *$\text{gs}_A(n)$ equals $\min_S E_n(S)$, and the minimum is attained. (Lean: `gs_offFn_eq_energy_min`.)*

Proposition 2.3 (Representability). *$\text{gs}_A(n) = 0$ if and only if n is a sum of a sub-sequence of A . (Lean: `gs_eq_zero_iff`.)*

Proposition 2.3 is the translation used throughout: the statement “the landscape has ground energy zero” is the statement “ n is representable”, and $r_A(n)$ is the number of representations, denoted $r_A(n)$ below.

3 The classification theorem

Theorem 3.1 (Classification). *Let $a_i > 0$ for all i , let $S \subseteq \mathbb{F}_k$ and $d = E_n(S)$.*

1. *If $\sigma(S) \geq n$: S is a strict local minimum $\iff a_i > 2d$ for all $i \in S$.*
2. *If $\sigma(S) \leq n$: S is a strict local minimum $\iff a_i > 2d$ for all $i \notin S$.*

(Lean: `isStrictLocalMin_iff_of_ge`, `isStrictLocalMin_iff_of_le`, combined form `isStrictLocalMin_iff_classCond`.)

Proof. Suppose $\sigma(S) = n + d$ with $d \geq 0$. Removing $i \in S$ changes the sum to $n + d - a_i$, whose error is $|d - a_i|$; this is $\geq d$ precisely when $a_i \geq 2d$, and $> d$ precisely when $a_i > 2d$. Adding $i \notin S$ only increases the error, since $a_i > 0$. The deficient case is symmetric. $\square$

We claim no originality for Theorem 3.1: the computation is immediate, and some form of it is presumably known to anyone who has looked at the 1-bit-flip landscape of the NPP, though we have not located an explicit statement in the literature. We record it here because everything below is built on it, and because the formal development requires a precise version. What we do claim is the use made of it in §§5–8.

Corollary 3.2. *A ground state ($d = 0$) is always a strict local minimum.* (Lean: `isStrictLocalMin_of_exact`.)

Corollary 3.3. *$\text{lm}_A(n)$ equals the number of S satisfying the arithmetic condition of Theorem 3.1; in particular lm is a counting function of the sequence and target alone.* (Lean: `lm_eq_lmClassCount`.)

The force of Theorem 3.1 is a *locality* statement: whether S is a valley depends only on the elements of A that are at most $2d$, where d is the error at S . Since a landscape has few configurations at large error, and the constraint at error d involves $N_A(d) = \#\{i : a_i \leq 2d\}$ elements, the count of valleys is governed by how quickly $N_A(d)$ grows — that is, by the early gaps of A . Sections 5–8 make this precise.

4 Symmetry and parity

Throughout this section $T = \sigma(A)$.

Theorem 4.1 (Reflection). *For $n \leq T$: $\text{gs}_A(T - n) = \text{gs}_A(n)$ and $\text{lm}_A(T - n) = \text{lm}_A(n)$.* (Lean: `gs_compl`, `lm_compl`.)

Proof. The map $S \mapsto S^c$ is an involution of the configuration space satisfying $\sigma(S^c) = T - \sigma(S)$, hence $E_{T-n}(S^c) = E_n(S)$; and it commutes with single-element flips, $\text{flip}(S, i)^c = \text{flip}(S^c, i)$. So it carries strict local minima for target n bijectively onto those for target $T - n$. $\square$

Theorem 4.2 (No flat edges). *If every a_i is odd, then $E_n(S) \neq E_n(\text{flip}(S, i))$ for all S and i .* (Lean: `no_flat_edge`.)

Proof. $|x - n| \equiv x + n \pmod{2}$, and flipping i changes $\sigma(S)$ by the odd number a_i , hence changes the parity of the error. $\square$

Corollary 4.3. *For odd A , strict local minima and weak local minima coincide.* (Lean: `isStrictLocalMin_iff_le_of_odd`.)

Corollary 4.3 is the reason we restrict to odd sequences: for such A — and the odd primes are the motivating example — the quantity $\text{lm}_A(n)$ does not depend on which of the two standard definitions of “local minimum” one adopts. For sequences containing even elements the landscape has plateaus and the two counts genuinely differ.

5 The window identity

5.1 Windows

By Theorem 3.1, a configuration at error d is a local minimum precisely when the $N_A(d)$ elements of size $\leq 2d$ all lie on a prescribed side. If the elements were placed independently and uniformly, the proportion of configurations satisfying this would be $2^{-N_A(d)}$. Summing over d , with the excess and deficient sides contributing separately and the ground level ($d = 0$) once, gives the *window series*

$$W_D(A) = 1 + 2 \sum_{d=1}^D 2^{-N_A(d)}, \quad N_A(d) = \#\{i : a_i \leq 2d\}.$$

(Lean: `winCount`, `windowSum`, `windowSeries`; the $\mathbb{F}_k$ version is `winCountFin`, `winSeriesFin`.)

W_D is the annealed prediction for $\text{lm}_A(n)/r_A(n)$. Sections 7–8 show in what exact sense it is correct; the present section evaluates it.

5.2 The identity

Theorem 5.1. *Let A be a non-decreasing sequence of odd positive integers with $a_i \leq 2D + 1$ for all i . Then*

$$W_D(A) = \sum_{j=1}^k \frac{a_j}{2^j} + \frac{2D+1}{2^k} = \Gamma(A) + \frac{2D+1-a_k}{2^k}.$$

In particular $W_D(A) = \Gamma(A)$ at the smallest admissible window, $D = (a_k - 1)/2$: the gap series is the window series evaluated at the last non-empty layer, and Definition 1.1 is that statement taken as the definition. Enlarging D beyond $(a_k - 1)/2$ adds empty layers, each contributing exactly $2 \cdot 2^{-k}$, which is the $2D + 1 - a_k$ above. (Lean: `windowSeries_eq_gapSeries`, in the first form.)

Proof. Induction on the length of A , the inductive step resting on the pointwise identity

$$2^{-N_{a:t}(d)} = \frac{1}{2} \left(2^{-N_t(d)} + \mathbf{1}[2d < a] \right),$$

valid when $a \leq b$ for all $b \in t$ (Lean: `window_term_cons`). Indeed if $a \leq 2d$ then $N_{a:t}(d) = N_t(d) + 1$ and the indicator vanishes; if $2d < a$ then no element of $a : t$ is $\leq 2d$, so both sides equal 1. Summing over $d \in [1, D]$ and using $\#\{d \leq D : 2d < a\} = \lfloor a/2 \rfloor$ (here $a \leq 2D + 1$ is used) gives

$$\sum_{d=1}^D 2^{-N_{a:t}(d)} = \frac{1}{2} \left(\sum_{d=1}^D 2^{-N_t(d)} + \lfloor a/2 \rfloor \right),$$

whence $W_D(a : t) = 1 + \lfloor a/2 \rfloor + \frac{1}{2}(W_D(t) - 1)$. Substituting the inductive hypothesis $W_D(t) = \Gamma(t) + (2D + 1)2^{-|t|}$ and using $\lfloor a/2 \rfloor = (a - 1)/2$ for odd a , together with $\Gamma(a : t) = \frac{a}{2} + \frac{\Gamma(t)}{2}$, gives the claim. The base case $A = \emptyset$ reads $1 + 2D = 0 + (2D + 1)$. $\square$

Note that D may be taken arbitrarily large: increasing D adds terms 2^{-k} on the left, absorbed exactly by the tail term on the right. The natural choice for an increasing sequence is $D = (a_k - 1)/2$, the smallest value for which all elements are $\leq 2D + 1$; then $W(A) = \Gamma(A) + a_k 2^{-k}$, and for the first twenty odd primes the correction is $73/2^{20} \approx 6.96 \times 10^{-5}$.

5.3 Why the enumeration form needs the order

Both sides of Theorem 5.1 as stated in the second form are invariants of the set: W_D counts, and so does Γ . The *enumeration* form $\sum_j a_j 2^{-j}$ does not, and the hypothesis that A be non-decreasing is what the first form of the theorem pays for it. The failure is worth exhibiting, because it is also the reason Definition 1.1 was chosen over the enumeration series.

Example 5.2. Take the sequence $(5, 3)$ and $D = 2$. The counting function does not see the order: $N(1) = 0$, $N(2) = 1$, so $W_2 = 1 + 2(1 + \frac{1}{2}) = 4$, and the gap series of the underlying set $\{3, 5\}$ is $\Gamma = 1 + 2(1 + \frac{1}{2}) = 4$ as well — here $D = (M - 1)/2$, so the two agree exactly. The enumeration series does see the order: read in the given order it is $\frac{5}{2} + \frac{3}{4} = \frac{13}{4}$, and $\frac{13}{4} + \frac{5}{4} = \frac{9}{2} \neq 4$; read in increasing order it is $\frac{3}{2} + \frac{5}{4} = \frac{11}{4}$, and $\frac{11}{4} + \frac{5}{4} = 4$. Both computations are `#eval`-checked in the Lean development.

Two consequences. Inside this paper: the sorting hypothesis in Theorem 5.1 is not removable, so any argument that reaches for the enumeration form must carry it. Outside: no order-blind *spread* statistic can play the role of Γ either, which is the subject of §10. Being an invariant of the set is necessary, and Example 5.2 shows the enumeration series fails it; §10 shows the standard invariants that pass it fail for a different reason.

6 Extremal values

Definition 1.1 makes the range of Γ visible at once. Fix the largest element M and write $D = (M - 1)/2$, so that the sum in Definition 1.1 runs over $1 \leq d \leq D$. Every term $2^{-N_A(d)}$ lies between 2^{-d} and 1, the upper bound because $N_A(d) \geq 0$ and the lower because $N_A(d) \leq d$ — there are only d odd numbers at or below $2d$. The two extremes are attained by the two extreme sets.

Theorem 6.1 (extremal values). *Let A be a finite set of odd positive integers with largest element M , and put $D = (M - 1)/2$. Then*

$$3 - 2^{1-D} \leq \Gamma(A) \leq M.$$

The upper bound is attained if and only if $A = \{M\}$, and the lower bound if and only if $A = \{1, 3, \dots, M\}$. (Lean: `windowSeries_bounds`, `windowSeries_eq_max_iff`, `windowSeries_min_iff_all_odds`.)

Proof. Since $N_A(d) \geq 0$ we have $2^{-N_A(d)} \leq 1$, so $\Gamma(A) \leq 1 + 2D = M$; and equality forces $N_A(d) = 0$ for every $d \leq D$, that is, A has no element at or below $2D = M - 1$, that is, $A = \{M\}$. Conversely $\Gamma(\{M\}) = 1 + 2D = M$.

For the lower bound, the odd numbers at or below $2d$ are exactly $1, 3, \dots, 2d - 1$, so $N_A(d) \leq d$ and $2^{-N_A(d)} \geq 2^{-d}$. Summing the geometric series,

$$\Gamma(A) \geq 1 + 2 \sum_{d=1}^D 2^{-d} = 1 + 2(1 - 2^{-D}) = 3 - 2^{1-D}.$$

Equality forces $N_A(d) = d$ for every $d \leq D$; but $N_A(D) = D$ already suffices, since $N_A(D) = D$ says that A contains all D of the odd numbers $1, 3, \dots, 2D - 1 = M - 2$, and then $N_A(d) = d$ holds for every $d \leq D$ automatically. Adjoining $M = \max A$ gives $A \supseteq \{1, 3, \dots, M\}$, and A contains nothing else, being a set of odd numbers bounded by M . Conversely $N_{\{1, 3, \dots, M\}}(d) = d$ for $d \leq D$, which gives equality. $\square$

Remark 6.2. The lower bound is attained as soon as it is attained in the last layer: a single equation, $N_A(D) = D$, pins A completely. Nothing similar is true of the upper bound, where all D equations $N_A(d) = 0$ are needed — but they are all implied by the last one, $N_A(D) = 0$, for the same reason of monotonicity. Both extremal characterisations are therefore statements about the top layer alone.

The mechanism behind the theorem is monotonicity, which is worth isolating, since it is what licenses reading Γ as a measure of sparsity.

Corollary 6.3 (strict monotonicity). *Let A be a finite set of odd positive integers with largest element M , and let b be odd with $b \notin A$ and $b < M$. Then*

$$\Gamma(A \cup \{b\}) < \Gamma(A).$$

Proof. Adjoining b does not change M , so the two sums run over the same range $1 \leq d \leq D$. Termwise, $N_{A \cup \{b\}}(d) = N_A(d) + \mathbf{1}[b \leq 2d] \geq N_A(d)$, so every term weakly decreases. Since b and M are odd and $b < M$ we have $b \leq M - 2 = 2D - 1$, so the indicator is 1 at $d = D$ and that term strictly decreases. $\square$

Thus Γ decreases strictly as A fills in, from M at the sparsest set $\{M\}$ down to $3 - 2^{1-D}$ at the densest set $\{1, 3, \dots, M\}$; and as $M \rightarrow \infty$ the lower bound tends to 3, so

$$\Gamma(\{1, 3, 5, \dots\}) = 3$$

in the sense that the gap series of the first D odd numbers converges to 3, the odd numbers being the unique minimiser at every finite stage. Against this scale, $\Gamma(P_{20}) = 1402281/262144 \approx 5.349$ for the first twenty odd primes lies near the bottom of the interval $[3 - 2^{-35}, 73]$ available to a set with $M = 73$: on the logarithmic scale of that interval the primes sit at 18%. This is the sense in which the primes are a *dense* set of weights, and the reason their subset-sum landscapes are comparatively smooth.

Remark 6.4 (on the two forms). Definition 1.1 and the enumeration form $\sum_j a_j 2^{-j}$ of Proposition 1.5 differ by $a_k 2^{-k}$ and converge to the same limit whenever $a_k 2^{-k} \rightarrow 0$, which holds for the primes ($a_k \sim k \log k$) and for any sequence of subexponential growth. At finite k they differ, and Definition 1.1 is the closer approximant: for the odd primes at $k = 20$ the layer form lies $1.04 \cdot 10^{-5}$ from the limit against $8.01 \cdot 10^{-5}$ for the enumeration form, a factor of 7.7. The reason is visible in Theorem 5.1: the enumeration form is the window series with the whole tail $(2D + 1)2^{-k}$ still attached, whereas Definition 1.1 is the window series at the last informative layer, with the tail removed exactly.

7 Exact stratification

7.1 Strata

Fix A with $a_i > 0$ and a target n . Following Theorem 3.1, define for $d \geq 0$

$$\begin{aligned} \text{over}_A(n, d) &= \#\{S : \sigma(S) = n + d, \forall i \in S, a_i > 2d\}, \\ \text{under}_A(n, d) &= \#\{S : \sigma(S) + d = n, \forall i \notin S, a_i > 2d\}. \end{aligned}$$

(Lean: `overCount`, `underCount`.)

Lemma 7.1. *The set of strict local minima with $\sigma(S) = s$ has cardinality $\text{over}_A(n, s - n)$ if $s \geq n$, and $\text{under}_A(n, n - s)$ if $s < n$. (Lean: `fiber_over`, `fiber_under`.)*

Lemma 7.2. $\text{over}_A(n, 0) = \text{under}_A(n, 0) = r_A(n)$. If $a_i \leq 2d$ for all i and $d \geq 1$, then $\text{over}_A(n, d) = \text{under}_A(n, d) = 0$. (Lean: `overCount_zero_eq_degCount`, `underCount_zero_eq_degCount`, `overCount_eq_zero`, `underCount_eq_zero`.)

Theorem 7.3 (Exact stratification). For $n \leq T$ and any D with $a_i \leq 2D + 1$ for all i ,

$$\text{lm}_A(n) = r_A(n) + \sum_{d=1}^D (\text{over}_A(n, d) + \text{under}_A(n, d)).$$

(Lean: `lm_eq_sum_strata` in the sum-indexed form; `lm_eq_lmPartial` and `lm_eq_lmPartial_of_le` in the d -indexed form used here.)

Proof. Fibring the set of strict local minima over the value $\sigma(S) \in [0, T]$ and applying Lemma 7.1 gives the sum-indexed identity. Re-indexing the excess part by $d = s - n$ and the deficient part by $d = n - s$ (the latter via reflection of a range sum), then extracting the $d = 0$ term and truncating both sums at D using Lemma 7.2, gives the stated form. $\square$

We stress that Theorem 7.3 is an identity between integers. No independence assumption enters.

7.2 Strata are representation counts

Write $I_d = \{i : a_i \leq 2d\}$, $N(d) = |I_d|$, $\sigma_d = \sigma(I_d)$, and let B_d be the subsequence of A indexed by I_d^c . As in §2.2, $r_B(m)$ denotes the number of subsets of B summing to m .

Theorem 7.4 (Fibres). Let $J \subseteq I$. Then

$$\#\{S : \sigma(S) = n, S \cap I = J\} = r_{A|_{I^c}}(n - \sigma(J)).$$

(Lean: `fiber_card_eq_repCount`; the bijection is $S \mapsto S \setminus I$, with inverse $U \mapsto U \cup J$.)

Corollary 7.5. $r_A(n) = \sum_{J \subseteq I} r_{A|_{I^c}}(n - \sigma(J))$ for any I . (Lean: `degCount_eq_sum_repCount`, combining Theorem 7.4 with `degCount_fiberwise`.)

Corollary 7.6. $\text{over}_A(n, d) = r_{B_d}(n + d)$ and, for $d \leq n$, $\text{under}_A(n, d) = r_{B_d}(n - d - \sigma_d)$. (Lean: `overCount_eq_repCount`, `underCount_eq_repCount`; these are Theorem 7.4 with $J = \emptyset$ and $J = I_d$ respectively.)

7.3 What has been achieved

Combining the above,

$$\frac{\text{lm}_A(n)}{r_A(n)} = 1 + \sum_{d=1}^D \frac{r_{B_d}(n + d) + r_{B_d}(n - d - \sigma_d)}{\sum_{J \subseteq I_d} r_{B_d}(n - \sigma(J))}.$$

Every quantity on the right is a representation count of a truncated subsequence of A . The ratio of valleys to ground states has been expressed *exactly* in terms of the arithmetic function r , with no probabilistic input whatever.

The denominator is a sum of $2^{N(d)}$ values of r_{B_d} at the points $n - \sigma(J)$, $J \subseteq I_d$; the numerator consists of two further values of the same function. If r_{B_d} were constant on the window spanned by these $2^{N(d)} + 2$ points, the d -th summand would equal exactly $2 \cdot 2^{-N(d)}$ and the total would be exactly $W_D(A) = \Gamma(A) + (2D + 1)2^{-k}$ by Theorem 5.1.

The annealed approximation of the physics literature is thus revealed to be exactly the assertion that r_{B_d} is constant on this window. It is not; but it is nearly so, and “nearly” is quantifiable. That is the content of §8.

8 Flatness and the sandwich

8.1 The abstract lemma

Lemma 8.1. *Let s be a finite non-empty index set, $g : s \rightarrow \mathbb{Q}_{>0}$, $v \in \mathbb{Q}_{>0}$ and $\varepsilon \geq 0$. If $g(i) \leq (1 + \varepsilon)v$ and $v \leq (1 + \varepsilon)g(i)$ for all $i \in s$, then*

$$\frac{1}{(1 + \varepsilon)|s|} \leq \frac{v}{\sum_{i \in s} g(i)} \leq \frac{1 + \varepsilon}{|s|}.$$

(Lean: `ratio_bounds_of_flat`.)

Proof. $\sum g \geq |s|v/(1 + \varepsilon)$ gives the upper bound; $\sum g \leq |s|(1 + \varepsilon)v$ gives the lower. $\square$

Elementary as it is, Lemma 8.1 is the whole mechanism: it converts flatness of a family into a bound on the share of any one member.

8.2 The sandwich

Definition 8.2. Say A is ε -flat at (n, d) if $r_{B_d}(m') \leq (1 + \varepsilon)r_{B_d}(m)$ for all $m, m' \in \text{Win}_d$, where

$$\text{Win}_d = \{n - \sigma(J) : J \subseteq I_d\} \cup \{n + d, n - d - \sigma_d\}.$$

Theorem 8.3 (Sandwich). *Let $a_i > 0$ be odd, $n \leq T$, and let D satisfy $a_i \leq 2D + 1$ for all i . Suppose $r_A(n) > 0$ and A is ε -flat at (n, d) for every $1 \leq d \leq D$. Then*

$$\frac{W_D(A)}{1 + \varepsilon} \leq \frac{\text{lm}_A(n)}{r_A(n)} \leq (1 + \varepsilon)W_D(A),$$

and by Theorem 5.1 the outer terms equal $(\Gamma(A) + (2D + 1)2^{-k})(1 + \varepsilon)^{\mp 1}$. (Lean: `stratum_ratio_bounds`, `sandwich_partial`, `sandwich_total`.)

Proof. Fix d . By Corollary 7.5 the denominator $r_A(n)$ is the sum of the $2^{N(d)}$ numbers $r_{B_d}(n - \sigma(J))$, all of which lie within a factor $1 + \varepsilon$ of $\text{over}_A(n, d) = r_{B_d}(n + d)$ by ε -flatness. Lemma 8.1 gives

$$\frac{2^{-N(d)}}{1 + \varepsilon} \leq \frac{\text{over}_A(n, d)}{r_A(n)} \leq (1 + \varepsilon)2^{-N(d)},$$

and the same for under. Summing over $1 \leq d \leq D$ and adding the $d = 0$ term $r/r = 1$, then invoking Theorem 7.3, gives the result; the elementary inequalities $1 \leq 1 + \varepsilon$ and $(1 + \varepsilon)^{-1} \leq 1$ absorb the ground term. $\square$

Remark 8.4. The exponent is $(1 + \varepsilon)^{\pm 1}$, not $(1 + \varepsilon)^{\pm 2}$ as a first analysis suggests, because the same flatness hypothesis controls numerator and denominator simultaneously. We record this because it was found only in the course of formalising the argument.

Remark 8.5. Theorem 8.3 is a finite statement over $\mathbb{Q}$. For any explicit A , n , D and ε all quantities are computable, so the theorem yields certified error bounds for individual instances rather than an asymptotic statement. The asymptotic question — whether ε may be taken to 0 as $k \rightarrow \infty$ — is the subject of §8.4 and §11.

8.3 A universal obstruction to flatness

Before measuring ε we record a lower bound for it, valid for every odd sequence and proved by an entirely elementary character computation. Recall $F_B(z) = \prod_{a \in B} (1 + z^a)$ and set $b = |B|$ and $R_c = \sum_{m \equiv c(4)} r_B(m)$ for $c \in \mathbb{Z}/4$.

Theorem 8.6 (Universal modulus-4 obstruction). *Let B be any finite sequence of odd positive integers, $b = |B|$. Then*

1. $|F_B(i)| = 2^{b/2}$, exactly, independently of the residues of the elements of B ; in particular $F_B(i) \neq 0$. More precisely $F_B(i) = 2^{b/2} e^{i\pi(p-q)/4}$ where $p = \#\{a \in B : a \equiv 1(4)\}$ and $q = \#\{a \in B : a \equiv 3(4)\}$.
2. The spread of the residue-class totals is given exactly by the parity of b :

$$\max_c R_c - \min_c R_c = \begin{cases} 2^{b/2} & b \text{ even,} \\ 2^{(b-1)/2} & b \text{ odd.} \end{cases}$$

3. Since $\sum_c R_c = 2^b$, the relative spread equals $4 \cdot 2^{-b/2}$ for b even and $2\sqrt{2} \cdot 2^{-b/2}$ for b odd; in either case

$$\frac{\max_c R_c - \min_c R_c}{\frac{1}{4} \sum_c R_c} \geq 2\sqrt{2} \cdot 2^{-b/2}.$$

(Lean: `normSq_one_add_I_pow_odd`, `max_abs_cos_sin_sq_ge`, `spread_lower_bound` give (i) and the inequality in (iii); `max_abs_cos_sin_sq`, `max_abs_cos_sin_quarter_even` and `max_abs_cos_sin_quarter_odd` give the parity dichotomy in (ii), the spread being $2^{b/2} \max(|\cos \varphi|, |\sin \varphi|)$ with $\varphi = \pi(p - q)/4$ and $p - q \equiv b \pmod{2}$.)

Proof. (i) For odd a we have $i^a = \pm i$, so $1 + i^a$ is $1 + i$ or $1 - i$ according as $a \equiv 1$ or $3 \pmod{4}$; both have modulus $\sqrt{2}$, and $1 \pm i = \sqrt{2} e^{\pm i\pi/4}$. Multiplying over B gives $F_B(i) = 2^{b/2} e^{i\pi(p-q)/4}$.

(ii) By discrete Fourier inversion on $\mathbb{Z}/4$,

$$R_c = \frac{1}{4} \sum_{j=0}^3 i^{-jc} F_B(i^j) = \frac{1}{4} \left(2^b + 2 \operatorname{Re}(i^{-c} F_B(i)) \right),$$

because $F_B(1) = 2^b$ and $F_B(-1) = \prod_{a \in B} (1 + (-1)^a) = 0$, every a being odd. Writing $\varphi = \pi(p - q)/4$, the quantity $\operatorname{Re}(i^{-c} F_B(i))/2^{b/2}$ runs over $\{\cos \varphi, \sin \varphi, -\cos \varphi, -\sin \varphi\}$ as c runs over $\mathbb{Z}/4$, so the spread is $2^{b/2} \cdot 2 \max(|\cos \varphi|, |\sin \varphi|)/2$. By (i) φ is a multiple of $\pi/4$, and $p + q = b$ forces $p - q \equiv b \pmod{2}$. If b is even then φ is a multiple of $\pi/2$ and $\max(|\cos \varphi|, |\sin \varphi|) = 1$; if b is odd then φ is an odd multiple of $\pi/4$ and the maximum is $1/\sqrt{2}$. This gives the two cases.

(iii) Divide by $\frac{1}{4} \sum_c R_c = 2^{b-2}$. □

The identity in (ii) is worth emphasising: the spread is not merely bounded below, it is *determined*, and only by the parity of b — the residues of the individual elements are irrelevant. We have checked it numerically for the truncated prime sequences and for random odd sequences up to $k = 32$, where the ratio of the measured relative spread to $2\sqrt{2} \cdot 2^{-b/2}$ is exactly 1 when b is odd and exactly $\sqrt{2}$ when b is even, with no intermediate values.

Two remarks. First, the vanishing $F_B(-1) = 0$ used in the proof is Theorem 4.2 (no flat edges) in Fourier form, so the parity structure that makes lm well defined is the same structure that forces the modulus-4 obstruction.

Second, the contrast with the modulus 6 is the point. The factor $1 + \zeta_6^a$ can vanish, namely when $a \equiv 3 \pmod{6}$, and then the whole modulus-6 mode disappears. The factor $1 + i^a$ never vanishes for odd a : the modulus-4 obstruction is present in every odd sequence, and by (ii) its size is the same for all of them.

What Theorem 8.6 does and does not say

We must be careful about which quantity is being bounded, and we state the limitation explicitly because it is easy to overread the theorem.

Theorem 8.6 concerns the *aggregate* counts $R_c = \sum_{m \equiv c(4)} r_B(m)$, sums over full residue classes. The quantity appearing in Theorem 8.3 is different: ε_d is a *pointwise* oscillation over the bounded window Win_d . The former is a statement about the global distribution of r_B modulo 4; the latter about its behaviour on $O(1)$ many points. Neither implies the other formally, and we do not claim a lower bound for ε_d .

What we can say is that the two live at the same scale, and that the numerics bear this out. Write $L(b) = 2\sqrt{2} \cdot 2^{-b/2}$ for the quantity in (iii). For random odd sequences the measured ε_2 satisfies $\varepsilon_2/L(b) \in [0.85, 1.77]$ over $k = 16, \dots, 32$, hovering about 1 — and note that it dips *below* 1 at $k = 24$, which is possible precisely because L bounds the aggregate spread and not ε_2 . For the primes, by contrast, $\varepsilon_2/L(b)$ grows from 11.7 at $k = 16$ to 293.8 at $k = 32$, since the modulus-6 obstruction decays at rate $\sqrt{3}/2 \approx 0.866$ against the modulus-4 rate $1/\sqrt{2} \approx 0.707$.

So the informal reading is: random odd sequences are about as flat as the modulus-4 obstruction permits, and the primes are two orders of magnitude less flat than that, for the arithmetic reason isolated in Proposition 11.5. One cannot, however, upgrade the first half of that sentence to a theorem covering all odd sequences: no window analogue of Theorem 8.6 holds, and Remark 11.2 exhibits sequences with $\varepsilon_2 = 0$ exactly.

8.4 Flatness measured

We computed ε_d exactly, by dynamic programming over $\mathbb{Z}$, for the initial segments of the odd primes and for random odd sequences drawn from the same range, at $k = 8, 10, \dots, 24$ and $n = \lfloor T/2 \rfloor$. Table 1 gives the primes.

k	r	lm	lm / r	Γ	ε_1	ε_2	ε_3	ε_4
14	126	628	4.9841	5.3486	0.0328	0.3125	0.3200	0.5833
16	394	2138	5.4264	5.3492	0.0051	0.1827	0.2386	0.6944
18	1290	6706	5.1984	5.3493	0.0031	0.1518	0.1944	0.4122
20	4344	22808	5.2505	5.3493	0.0007	0.1104	0.1462	0.2987
22	14846	80008	5.3892	5.3493	0.0009	0.0734	0.1112	0.2283
24	51068	270392	5.2947	5.3493	0.0003	0.0570	0.0789	0.1573

Table 1: Exact flatness ε_d for $A =$ first k odd primes, $n = \lfloor T/2 \rfloor$. The fifth column is the window series W_D at the last informative layer $D = (a_k - 1)/2$, which by Theorem 5.1 is exactly Γ .

The flatness decays geometrically; see Figure 1. Log-linear regression on $k \geq 14$ gives the decay rates per unit increase of k shown in Table 2.

Two features stand out.

(i) *The rate matches $\sqrt{3}/2$.* We have $\sqrt{3}/2 = 0.86603$, and the measured ripple rates lie in $[0.857, 0.868]$, i.e. within 1.1%. Here is the heuristic that predicts it. Every prime ≥ 5 is $\equiv \pm 1 \pmod{6}$. Consequently the generating polynomial of B_d evaluated at a primitive sixth root of unity ζ is $\prod_{a \in B_d} (1 + \zeta^a)$ with each factor of modulus $|1 + \zeta^{\pm 1}| = \sqrt{3}$, against 2




Figure 1: Exact flatness $\varepsilon_d(k)$ on a logarithmic scale. The primes (blue) track the reference slope $(\sqrt{3}/2)^k$ predicted by the modulo-6 heuristic of §8.4(i); random odd sequences (red, median over 20 seeds) decay strictly faster.

d	$\lambda(\varepsilon_d)$, primes	R^2	$\lambda(\text{ripple})$	R^2	$\lambda(\varepsilon_d)$, random
2	0.8478	0.986	0.8590	0.989	0.683
3	0.8721	0.996	0.8577	0.988	0.611
4	0.8643	0.929	0.8569	0.990	—
6	—	—	0.8684	0.983	—

Table 2: Decay rate λ of $\varepsilon_d(k)$ and of the modulo-6 ripple amplitude, from log-linear regression on $k \geq 14$. Compare $\sqrt{3}/2 = 0.86603$.

for the trivial character. The non-trivial character therefore contributes a ripple of relative amplitude $(\sqrt{3}/2)^{|B_d|}$ to r_{B_d} , periodic with period 6. Since $|B_d| = k - N(d)$ and $N(d)$ is bounded for fixed d , the amplitude decays like $(\sqrt{3}/2)^k$. The measurement confirms both the rate and the period.

(ii) *Random sequences are flatter.* Random odd sequences show $\lambda \approx 0.68$ and 0.61 , i.e. strictly faster decay, and at $k = 24$ their median ε_2 is 0.0007 against 0.0570 for the primes — a factor of 80. The primes are subject to an arithmetic obstruction that generic sequences are not.

This is the sense in which the primes are simultaneously *smoother* and *less flat* than random sequences: smoother at first order, because Γ is small (§9), but less flat at second order, because of the modulo-6 bias. We isolate the resulting conjecture in §11.

9 Numerical experiments

9.1 Setup

All landscapes below are enumerated exhaustively. For $k \leq 20$ we used a compiled Lean executable enumerating all 2^k configurations (the full 100-seed sweep over $k \in \{8, 12, 16, 18, 20\}$ takes 11.5 s); for $k \leq 24$ we used exact dynamic programming over $\mathbb{Z}$ via Theorem 7.3, cross-validated against the enumeration at $k = 20$ (both give $\text{lm} = 22808$, $r = 4344$). Random sequences are drawn without replacement from the odd integers in $[3, a_k]$, matching the range of the primes of the same length. The target is $n = \lfloor T/2 \rfloor$ throughout.

9.2 lm/r converges to Γ

Put $Q = (\text{lm}/r)/\Gamma$. Theorem 8.3 predicts $Q \in [(1+\varepsilon)^{-1}, 1+\varepsilon]$ up to the negligible tail term. Table 3 and Figure 2 record the behaviour.

k	mean Q	s.d.	slope a (s.e.)	intercept b	R^2	corr(Γ, lm)
8	1.096	0.492	1.904 (0.275)	−4.909	0.329	0.841
12	1.002	0.080	1.050 (0.038)	−0.342	0.884	0.956
16	0.993	0.025	1.033 (0.014)	−0.313	0.982	0.951
18	0.992	0.014	1.028 (0.007)	−0.281	0.995	0.982
20	0.995	0.0067	1.013 (0.003)	−0.142	0.99903	0.981

Table 3: Q over 100 random odd sequences per k ; regression of lm/r on Γ .

The dispersion contracts steadily. Correlation between Γ and lm/r reaches 0.9995 at $k = 20$.

9.3 Primes against random

k	8	12	16	18	20
$z(\text{lm})$	−1.43	−0.96	−1.17	−1.09	−1.14
percentile of lm	10	21	8	5	4
$z(\Gamma)$	−0.76	−1.20	−1.30	−1.09	−1.09

Table 4: Position of the primes within the 100 random instances.

The primes have fewer valleys than a random sequence of the same length and range at every k tested, but the effect is of size ≈ 1 standard deviation and is not individually




Figure 2: Convergence of $Q = (\text{lm}/r)/W_D$. Error bars are one standard deviation over 100 random sequences; they contract steadily, consistent with $\varepsilon \rightarrow 0$. The primes (blue) oscillate about 1 rather than converging monotonically; see §9.4 and Conjecture 11.6.

significant. What Table 4 does show clearly is that $z(\text{lm}) \approx z(\Gamma)$: whatever smoothness the primes possess is accounted for by their gap series, exactly as Theorems 5.1 and 8.3 predict.

9.4 An oscillation

The exact computation at $k \leq 24$ reveals a feature invisible at the resolution of Table 3. For the primes,

$$Q(k) = 1.035, 0.883, 1.053, 0.932, 1.014, 0.972, 0.982, 1.007, 0.990 \quad (k = 8, 10, \dots, 24),$$

oscillating about 1 rather than converging monotonically, whereas the median Q over random instances increases monotonically from below to 0.9967 at $k = 24$. We believe this oscillation is the signature of the modulo-6 ripple of §8.4 beating against the target $n = \lfloor T/2 \rfloor$; see Conjecture 11.6.

10 Comparison with the coefficient of variation

Alyahya and Rowe [AR14] reported that the number of local optima in the 1-bit-flip NPP landscape correlates strongly and negatively with the coefficient of variation $CV = \text{sd}(a)/\text{mean}(a)$ of the weights, and proposed an estimator on that basis. Using the 100-seed data of §9 we first reproduce their finding, then separate the two predictors. This section is empirical throughout: it compares two statistics on a fixed random model and claims no theorem. What it is meant to settle is a question of *which* quantity the theorems of §5 and §8 are about, not how well either quantity predicts.

For lm/r — the quantity our theorems actually control — the separation is sharper still: at $k = 20$, $\text{partial}(\Gamma, \text{lm}/r \mid CV) = +0.9986$ ($p = 2.4 \times 10^{-126}$) while $\text{partial}(CV, \text{lm}/r \mid \Gamma) = +0.306$.

k	$\text{corr}(CV, \text{lm})$	$\text{corr}(\Gamma, \text{lm})$	$\text{corr}(CV, \Gamma)$	$\text{partial}(\Gamma, \text{lm} \mid CV)$	$\text{partial}(CV, \text{lm} \mid \Gamma)$
8	-0.874	+0.841	-0.926	+0.176	-0.466
12	-0.878	+0.956	-0.907	+0.790	-0.092
16	-0.757	+0.951	-0.836	+0.887	+0.225
18	-0.741	+0.982	-0.807	+0.967	+0.451
20	-0.793	+0.981	-0.838	+0.952	+0.279

Table 5: Correlations and partial correlations, 100 random instances per k .

For $k \geq 12$, conditioning on CV leaves the predictive power of Γ essentially intact, whereas conditioning on Γ removes most of that of CV . We read this as: CV is a proxy for Γ , useful because the two are strongly correlated (≈ -0.85) for random sequences, but not the operative quantity.

Correlations on a random model can only ever be suggestive, since the two statistics are themselves correlated there (≈ -0.85). The clean separation is to hold one fixed and vary the other. Both are invariants of the set, so this is a question about which *features* of the set each one reads: CV is a global second moment, while by Proposition 1.2 Γ is a dyadically weighted sum of the gaps, dominated by the bottom of the set. The two are therefore decouplable, and Table 6 decouples them.

k	pair	CV ratio	Γ ratio	lm/r ratio
12	CV -twins	1.000000	2.135	2.051
13	CV -twins	1.000000	2.070	2.181
14	CV -twins	1.000000	2.240	2.325
12	Γ -twins	1.628	1.000000	1.080
13	Γ -twins	1.594	1.000000	1.034
14	Γ -twins	1.600	1.000000	1.047

Table 6: Holding one statistic fixed. A CV -*twin pair* is two sets of k odd numbers with the same cardinality, sum and sum of squares — hence exactly the same mean, variance and CV ; a Γ -*twin pair* is two sets with exactly the same gap series. Each lm/r is a median over the 41 targets n nearest $\sigma(A)/2$, computed by full enumeration of all 2^k configurations. When CV is pinned, the landscape moves with Γ ; when Γ is pinned, a 60% change in CV moves the landscape by at most 8%.

The pairs are exact, not approximate: CV -twins agree in CV to the last bit because they agree in $(k, \sum a_i, \sum a_i^2)$ as integers, and Γ -twins agree in Γ as rationals. For the $k = 14$ CV -twins, for instance, $\{1, 11, 31, \dots, 103\}$ and $\{29, 31, 33, \dots, 119\}$ both have $CV = 0.473014\dots$, while $\Gamma = 14.135$ and 31.658 and the measured lm/r are 13.60 and 31.63 — each within 4% of its own Γ , and a factor 2.3 apart, on data about which CV says the two sets are identical. The Γ -twin rows are the control that could have failed and did not.

This is what Theorem 5.1 predicts, and it is the sense in which the correlation of [AR14] is real but indirect: CV tracks Γ across a random model, and only across it.

A concrete illustration: for the primes, $z(CV)$ is $+0.85, +1.52, +1.70, +1.51, +1.57$ at $k = 8, \dots, 20$ — *positive*, i.e. the primes are more dispersed than random. Both accounts therefore predict “primes are smooth”, but only Γ predicts the magnitude, and only Γ appears in an identity.

11 Open problems

Problem 11.1 (Asymptotic flatness). Show that for the sequences of interest the flatness ε_k of Definition 8.2 tends to 0, so that Theorem 8.3 yields $\text{lm}_A(n)/r_A(n) \rightarrow \Gamma(A)$.

For random odd sequences this appears to be within reach of local central limit theorems together with Erdős–Littlewood–Offord anticoncentration [Er45], which bounds $\max_m r_B(m)$ from above and is exactly the ingredient needed to control the tail hypothesis. For the primes the situation is different. The classical asymptotic for partitions into distinct primes, due to Roth and Szekeres [RS54], gives $\log p(n, \mathcal{P}) \sim \kappa \sqrt{2n/\log n}$; *this is a statement on the logarithmic scale and does not suffice*, since flatness is a statement about *ratios* of r at nearby arguments. What is required is either an asymptotic with the pre-exponential factor for distinct-prime partitions of truncated prime sequences, or a direct exponential-sum estimate of Vinogradov type for r_{B_d} .

Remark 11.2 (There is no window analogue of Theorem 8.6). It is natural to ask whether Theorem 8.6 transfers from the aggregate counts R_c to the bounded windows Win_d , giving a lower bound $\varepsilon_d \gg 2^{-|B_d|/2}$ valid for every odd sequence. It does not: the window flatness can vanish identically. For $d = 2$ and a sequence A with $\min A > 4$ one has $I_2 = \emptyset$, so $\text{Win}_2 = \{n-2, n, n+2\}$ has only three points, and r_{B_2} can be made constant on them. For example, with

$$A = (7, 15, 21, 45, 49, 55, 67, 69, 73, 75, 81, 85, 89, 93, 97, 105, 107, 113, 115, 131, 135, 141, 145, 147, 151, 155, 159, \dots)$$

($k = 28$) all three values equal 377736, so $\varepsilon_2 = 0$ exactly. The agreement is to one part in 3.8×10^5 and is therefore not an artefact of the integrality of r ; sequences with $\varepsilon_2 = 0$ were found by local search at every size tested, with the common value growing from 10 at $k = 12$ to 377736 at $k = 28$.

The reason is simply that Win_d imposes $O(1)$ constraints on a sequence with k degrees of freedom. Theorem 8.6 constrains the *global* distribution of r_B modulo 4 and says nothing about any fixed finite set of points. What survives is the statement for the sequences we care about: for random odd sequences ε_2 sits at the modulus-4 scale (§8.3), and for the primes it sits two orders of magnitude above it — but neither is forced by Theorem 8.6 alone.

Conjecture 11.3 (Ripple amplitude). For the odd primes and fixed $d \geq 2$,

$$\frac{r_{B_d}(m)}{\text{Main}(m)} - 1 = 2 \left(\frac{\sqrt{3}}{2} \right)^{|B_d|+1} \cos \left(\frac{\pi m}{3} - \frac{\pi}{6} (c_1 - c_5) \right) + o((\sqrt{3}/2)^k),$$

where Main is the local-CLT Gaussian and $c_j = \#\{a \in B_d : a \equiv j \pmod{6}\}$; consequently $\varepsilon_d(k) \asymp (\sqrt{3}/2)^k$.

The arithmetic input here is not conjectural. Write $\zeta = e^{2\pi i/6}$ and $F_B(z) = \prod_{a \in B} (1 + z^a)$, so the r_B are the coefficients of F_B .

Proposition 11.4 (Character values). $|1 + \zeta^a|$ depends only on $a \pmod{6}$, taking the values $2, \sqrt{3}, 1, 0, 1, \sqrt{3}$ for $a \equiv 0, \dots, 5$. Hence if every element of B is a prime ≥ 5 then $F_B(\zeta) = 3^{|B|/2} e^{i(\pi/6)(c_1 - c_5)}$, so that $|F_B(\zeta)|/2^{|B|} = (\sqrt{3}/2)^{|B|}$ exactly; if $3 \in B$ then $F_B(\zeta) = 0$; and $F_B(-1) = 0$ for any odd B , which is Theorem 4.2 in Fourier form.

Proof. $|1 + e^{i\pi a/3}|^2 = 2 + 2 \cos(\pi a/3)$ takes the stated values. Every prime ≥ 5 is $\equiv \pm 1 \pmod{6}$ and $1 + \zeta^{\pm 1} = \sqrt{3} e^{\pm i\pi/6}$; multiplying gives the formula. For $a \equiv 3$ the factor $1 + \zeta^3$ vanishes, and $\zeta^3 = -1$ gives the last claim. $\square$

What remains conjectural is only the transfer from this Fourier datum to pointwise ratios. The exponent $|B_d| + 1$ rather than $|B_d|$ comes from the width factor $\sqrt{V_0/V_6} = \sqrt{3}/2$ at the sub-arc, where $V_0 = \frac{1}{4} \sum a^2$ and $V_6 = \frac{1}{4} \sum a^2 \sec^2(\pi a/6) = \frac{1}{3} \sum a^2$. We have tested both parts numerically for $k \leq 32$ by extracting the period-6 Fourier component of $r_{B_d}/\text{Main} - 1$.

The measured phase agrees with $(\pi/6)(c_1 - c_5)$ to four decimal places, including at $k = 16$, where $c_1 - c_5 = -3$ moves the predicted phase to $-\pi/2$.

For the amplitude, the extraction is sensitive to the width of the window over which the Fourier component is taken, because the Gaussian reference degrades in the tails: at $k = 26$ the ratio to $2(\sqrt{3}/2)^{|B_d|+1}$ is 0.972, 0.991, 1.056 for half-widths 24, 90, 180 respectively, while $\sigma = \sqrt{V_0} \approx 147$. Taking a narrow window (half-width 24, about 0.15σ , still eight full periods) the ratio is stable in the window and approaches 1; the corresponding ratios to $2(\sqrt{3}/2)^{|B_d|}$ move *away* from 1, which is what identifies the extra factor $\sqrt{3}/2$.

The approach to 1 splits into two branches according to the value of $c_1 - c_5$, which is also the phase:

$$\begin{array}{ll} c_1 - c_5 = -1 : & 0.965, 0.968, 0.970, 0.972, 0.974, 0.976, 0.981, 0.981 \quad (k = 20, \dots, 30, 36, 38) \quad \text{from below,} \\ c_1 - c_5 = -3 : & 1.025, 1.019, 1.015 \quad (k = 32, 34, 40) \quad \text{from above.} \end{array}$$

Both branches converge to 1, from opposite sides. Since the phase itself is reproduced to within 7×10^{-4} in both branches, the residual is not an error in the phase but a *phase-coupled* correction to the amplitude — which is what one expects from interference between the leading sub-arc term and the next one. We regard this as evidence for Conjecture 11.3 and as a constraint on any refinement of it.

We have also verified the width factor at a second modulus. For $q = 4$ the prediction is $\sqrt{V_0/V_4} = 1/\sqrt{2}$, and measuring the period-4 component on the stratum $d = 1$ of the primes (where $3 \in B_1$ kills the period-6 mode) gives ratios 0.91–0.97 to the prediction *with* the width factor against 0.64–0.69 *without* it; for random odd sequences the corresponding medians are ≈ 1.05 and ≈ 0.75 . The modulus $q = 3$ cannot be tested this way: for any odd sequence coprime to 3 one has $\rho_3 = (1/2)^{|B|}$, which is dominated by $\rho_4 = (1/\sqrt{2})^{|B|}$, so the period-3 component never rises above the period-4 one.

Proposition 11.5 (Extremality of the modulus 6). *Let $M(q)$ denote the geometric mean of $|\cos(\pi r/q)|$ over the r coprime to q . Then*

$$M(q) = \frac{1}{2} |\Phi_q(-1)|^{1/\varphi(q)},$$

where Φ_q is the q -th cyclotomic polynomial, and

$$\sup_{q \geq 3} M(q) = \frac{\sqrt{3}}{2},$$

attained at $q = 6$ and nowhere else. The next largest values are $M(10) = 5^{1/4}/2$, $M(4) = 1/\sqrt{2}$ and $M(14) = 7^{1/6}/2$. (Lean, finite part: `M12_six_maximal`, `M12_gap`.)

Proof. Since $\prod_{r \perp q} (x - \zeta_q^r) = \Phi_q(x)$, evaluating at $x = -1$ and using $|1 + \zeta_q^r| = 2|\cos(\pi r/q)|$ gives $|\Phi_q(-1)| = 2^{\varphi(q)} \prod_{r \perp q} |\cos(\pi r/q)|$, which is the identity. For the supremum use the classical values: $\Phi_q(-1) = p$ when $q = 2p^j$ with p an odd prime, $\Phi_q(-1) = 2$ when $q = 2^j$ with $j \geq 2$, and $\Phi_q(-1) = 1$ for all other $q \geq 3$. Hence $M(q) = 1/2$ except in two families. For $q = 2p^j$ we get $M(q) = \frac{1}{2} p^{1/(p^{j-1}(p-1))}$, which is largest when the exponent is largest, i.e. at $j = 1$; and $p \mapsto p^{1/(p-1)}$ decreases on odd primes ($3^{1/2} > 5^{1/4} > 7^{1/6} > \dots$), so the family is maximised at $q = 6$ with value $\sqrt{3}/2$, the runner-up being $q = 10$. For $q = 2^j$, $j \geq 2$, we get $M(q) = \frac{1}{2} \cdot 2^{1/2^{j-1}}$, maximised at $j = 2$ with value $1/\sqrt{2} < \sqrt{3}/2$. $\square$

We verified the identity numerically for $3 \leq q \leq 40$ to machine precision, obtaining $M(6) = 0.866025$, $M(10) = 0.747674$, $M(4) = 0.707107$, $M(14) = 0.691544$. The finite comparison is also formalised: raising to the twelfth power clears denominators, giving 729, 125, 64, 49, 8 over 4096 for $q = 6, 10, 4, 14, 8$ and $1/4096$ otherwise, so the maximality of $q = 6$ reduces to a comparison of rationals.

The number-theoretic content is this. For a sequence equidistributed among the residues coprime to q , $M(q)$ is the decay rate of $|F_B(\zeta_q)|/2^{|B|}$; Proposition 11.5 then says that $q = 6$ is the slowest mode available to any such sequence. The primes are confined to exactly two residues modulo 6, both of which realise the maximal $|\cos|$, and no other modulus offers a comparable concentration — which is why the bias modulo 6, and no other congruence obstruction, governs the flatness of the primes. The step from equidistribution to the decay rate is the one genuinely analytic ingredient here, and we do not prove it; see Problem 11.1.

Conjecture 11.6 (Oscillation). The oscillation of $Q(k)$ for the primes observed in §9.4 is caused by the ripple of Conjecture 11.3 beating against $n = \lfloor T/2 \rfloor$. Precisely, writing A_d, φ_d for the amplitude and phase in Conjecture 11.3 and $P_d = \prod_{a \in I_d} (1 + e^{-i\pi a/3})$,

$$\frac{\text{lm}_A(n)}{r_A(n)} - W_D(A) \approx \sum_{d \geq 1} 2^{-N(d)} \text{Re} \left[A_d e^{-i\varphi_d} (e^{i\pi(n+d)/3} + e^{i\pi(n-d-\sigma_d)/3} - 2e^{i\pi n/3} P_d 2^{-N(d)}) \right].$$

Evaluating the right-hand side for $k \leq 26$ reproduces the sign of the observed deviation in 9 of the 10 cases (all but $k = 8$) and its magnitude to within a factor 0.5–1.7, with a systematic under-estimate of about 25% that we attribute to higher-order sub-arc corrections. The stratum $d = 2$ dominates, and $d = 1$ contributes exactly zero because $3 \in B_1$ forces $F_{B_1}(\zeta) = 0$ by Proposition 11.4 — which also explains why ε_1 is two orders of magnitude below ε_2 in Table 1.

We warn against the simpler guess that the sign is governed by $T \bmod 6$ alone: for the primes $T \equiv 2 \pmod{6}$ throughout $k = 8, \dots, 14$ and $k = 18, \dots, 26$, yet the sign of the deviation alternates across that range. The dependence on $c_1 - c_5$ and on σ_d in the displayed formula is essential.

Problem 11.7 (Asymptotic smoothness of the primes). Determine whether $\text{lm}_{P_k}(n)/\text{lm}_{R_k}(n) \rightarrow 0$ as $k \rightarrow \infty$, where P_k is the first k odd primes and R_k a random odd sequence of the same length and range. By Theorem 8.3 this reduces to comparing $\Gamma(P_k)$, which converges (to ≈ 5.3493), with $\mathbb{E}[\Gamma(R_k)]$, which we expect to grow like twice the mean gap, i.e. $\asymp \log k$.

Problem 11.8 (Off-centre targets). For $n/T = \rho \neq 1/2$ the factor $2^{-N(d)}$ in the window series is replaced by $(1 - \rho)^{N(d)}$, deforming Γ into the family $\Gamma_x(A) = \sum_j a_j x^j$ evaluated at $x = 1 - \rho$. Develop the corresponding theory, including the large-deviation rate function that governs $\text{lm}/2^k$ away from the centre.

Tool and computational resource disclosure

This section follows recommendation 01 for individual mathematicians of the *Leiden Declaration on Artificial Intelligence and Mathematics* (June 2026, DOI: 10.5281/zenodo.20302944), which the author has signed.

Tools. This work was carried out by the author working with AI language models (Anthropic’s Claude, models `opus-5` and `fable-5`) as tools, under the author’s direction. They were used to fill in Lean proofs, to run and tabulate the numerical experiments, and to draft sections which the author then checked. Formal verification uses Lean 4 (toolchain `lean-prover/lean4:v4.32.2`) with Mathlib, and the kernel replay uses `lean4checker`. Typesetting uses Xe^LTeX; the numerical work uses Python 3 with no libraries beyond the standard one.

Computational resources. All of it runs on a single personal computer. No cluster, no accelerator, and no compute grant were used; the longest single computation in this Part is a kernel replay of the whole Lean development, which takes about three minutes. This is recorded because recommendation 10 of the Declaration asks whether smaller systems suffice, and here the honest answer is mixed: the verification and the experiments are small enough for anyone’s laptop, while the language models are commercial, proprietary and energy-intensive, and the author found no smaller system adequate for that role.

Responsibility and attribution. The choice of problem, the design decisions, the completeness of the citations, and the responsibility for every claim here are the author’s alone; credit belongs to no automated system. Where the author has been unable to attribute an idea satisfactorily, the text says so.

Why the apparatus exists. Work produced this way cannot reasonably be taken on the author’s word, and the apparatus described under Data availability exists for that reason: every settled theorem is machine-checked in Lean and replayed through the kernel by an independent checker that must first reject three deliberately corrupted modules; every number quoted in the text is produced by a committed script that writes a log beside itself; every theorem, proposition and lemma declares its status where it is stated; and a suite of mechanical checks enforces all of this before each commit. The repository records, in full, the mistakes this process has made — which is the part the author would argue for hardest, since those failures are precisely the ones no machine check can see.

Data availability

Every number in this paper is produced by a script that writes a log beside itself. In particular `extremal_r120` recomputes the constants of §6 as exact rationals and re-tests Corollary 6.3 on 20,000 random sets; `cvtwin_r120` constructs the twin pairs of Table 6 and measures their landscapes by full enumeration; and `resweep_r120` regenerates the 100-seed instances of §9 from the recorded seeds and recomputes Tables 3 and 5, reproducing the recorded per-instance values to $5 \cdot 10^{-7}$ before doing so. The scripts, the logs, the Lean development and the checks that enforce all of this are public at <https://github.com/mathlab0911/arithmetic-landscapes>.

References

- [AR14] K. Alyahya and J. E. Rowe, *Local optima and weight distribution in the number partitioning problem*, in: Parallel Problem Solving from Nature — PPSN XIII, Lecture Notes in Computer Science **8672**, Springer, Cham, 2014, pp. 862–871. DOI:10.1007/978-3-319-10762-2_85.
- [Ba21] P. Barry, *On the gap-sum and gap-product sequences of integer sequences*, preprint (2021), [arXiv:2104.05593](https://arxiv.org/abs/2104.05593).
- [BCP01] C. Borgs, J. T. Chayes and B. Pittel, *Phase transition and finite-size scaling for the integer partitioning problem*, Random Structures & Algorithms **19** (2001), no. 3–4, 247–288. DOI:10.1002/rsa.10004.
- [Er45] P. Erdős, *On a lemma of Littlewood and Offord*, Bulletin of the American Mathematical Society **51** (1945), 898–902.
- [FF98] F. F. Ferreira and J. F. Fontanari, *Probabilistic analysis of the number partitioning problem*, Journal of Physics A: Mathematical and General **31** (1998), no. 15, 3417–3428. DOI:10.1088/0305-4470/31/15/007.

- [Me01] S. Mertens, *A physicist's approach to number partitioning*, Theoretical Computer Science **265** (2001), no. 1–2, 79–108. DOI:10.1016/S0304-3975(01)00153-0.
- [RS54] K. F. Roth and G. Szekeres, *Some asymptotic formulae in the theory of partitions*, The Quarterly Journal of Mathematics **5** (1954), no. 1, 241–259. DOI:10.1093/qmath/5.1.241.
- [SHF03] P. F. Stadler, W. Hordijk and J. F. Fontanari, *Phase transition and landscape statistics of the number partitioning problem*, Physical Review E **67** (2003), 056701.

A The Lean development

The formalisation targets Lean 4 (toolchain `leanprover/lean4:v4.32.2`) with Mathlib pinned to the matching tag. Building with `lake build` verifies every declaration below. Each is checked with `#print axioms` to depend only on `propext`, `Classical.choice` and `Quot.sound`; no declaration depends on `sorryAx`.

Paper	Lean declaration	File
Prop. 2.1	<code>mem_subsetSums_ofFn_iff</code>	<code>Theory/Bridge.lean</code>
Prop. 2.2	<code>gs_ofFn_eq_energy_min</code>	<code>Theory/Bridge.lean</code>
Prop. 2.3	<code>gs_eq_zero_iff</code>	<code>Theory/Symmetry.lean</code>
Thm. 3.1	<code>isStrictLocalMin_iff_of_ge</code> etc.	<code>Theory/Landscape.lean</code>
Cor. 3.2	<code>isStrictLocalMin_of_exact</code>	<code>Theory/Landscape.lean</code>
Cor. 3.3	<code>lm_eq_lmClassCount</code>	<code>Theory/Landscape.lean</code>
Thm. 4.1	<code>gs_compl; lm_compl</code>	<code>Theory/Landscape.lean; Theory/Symmetry.lean</code>
Thm. 4.2	<code>no_flat_edge</code>	<code>Theory/Symmetry.lean</code>
Cor. 4.3	<code>isStrictLocalMin_iff_le_of_odd</code>	<code>Theory/Symmetry.lean</code>
Thm. 5.1	<code>windowSeries_eq_gapSeries</code>	<code>Theory/Bridge.lean</code>
Lem. 7.1	<code>fiber_over, fiber_under</code>	<code>Theory/Decomposition.lean</code>
Lem. 7.2	<code>overCount_eq_zero</code> etc.	<code>Theory/Decomposition.lean</code>
Thm. 7.3	<code>lm_eq_sum_strata, lm_eq_lmPartial</code>	<code>Theory/Decomposition.lean, Theory/Total.lean</code>
Thm. 7.4	<code>fiber_card_eq_repCount</code>	<code>Theory/Fiber.lean</code>
Cor. 7.5	<code>degCount_eq_sum_repCount</code>	<code>Theory/Fiber.lean</code>
Cor. 7.6	<code>overCount_eq_repCount</code> etc.	<code>Theory/Sandwich.lean</code>
Lem. 8.1	<code>ratio_bounds_of_flat</code>	<code>Theory/Sandwich.lean</code>
Thm. 8.3	<code>sandwich_partial, sandwich_total</code>	<code>Theory/Sandwich.lean, Theory/Total.lean</code>
Thm. 8.6(i)	<code>normSq_one_add_I_pow_odd</code>	<code>Theory/Ripple.lean</code>
Thm. 8.6(iii)	<code>max_abs_cos_sin_sq_ge, spread_lower_bound</code>	<code>Theory/Ripple.lean</code>
Prop. 11.5 (finite)	<code>M12_six_maximal, M12_gap</code>	<code>Theory/Ripple.lean</code>
Prop. 11.4	<code>normSq_one_add_exp</code> and specialisations	<code>Theory/Ripple.lean</code>

The development also contains machine-checked numerical cross-checks (`#eval`), among them: the identity of Theorem 7.3 for $A = (3, 5, 7, 11, 13, 17)$ at every $n \leq 56$; the coincidence of the strict and weak local-minimum counts for the same A (Corollary 4.3); the window identity for the first twenty odd primes, both sides evaluating to 1402281/262144; and the failure of that identity for the reordering $(5, 3)$ of Example 5.2.