

Arithmetic landscapes III: deformed measures, random sequences, and the coset identity

Kentaro Amauchi*

August 17, 2026

Abstract

For a finite set A of odd positive integers write $F_A(z) = \prod_{a \in A} (1 + z^a)$ for the generating polynomial of the subset sums, so that $r_A(n) = \#\{S \subseteq A : \sigma(S) = n\}$ is its n -th coefficient. Everything below concerns $|F_A|$ on the unit circle, where $|F_A(e^{i\theta})| = 2^{|A|} \prod_{a \in A} |\cos(a\theta/2)|$, and the per-element quantity $X(t) = -\log |\cos \pi t|$ that governs it.

Our first result is an identity for X alone, elementary and exact:

$$\frac{1}{v} \sum_{k=0}^{v-1} X\left(t + \frac{k}{v}\right) = \left(1 - \frac{1}{v}\right) \log 2 + \frac{1}{v} X(vt + \tau_v) \quad \text{for every } v \geq 1 \text{ and every } t,$$

with $\tau_v \in \{0, \frac{1}{2}\}$ explicit. Averaging X over a coset of index v returns the same function at v times the frequency, plus a constant. *The identity is classical:* it is the Kubert distribution relation for $\log |2 \sin \pi t|$ [3], transported to the cosine by $t \mapsto t + \frac{1}{2}$, and we prove it twice — by the multiplication formula, which is how the classical relation is proved, and by Fourier. What is new here is the use. Because $X \geq 0$, the rational points are the *minima* of the coset average and not merely convenient sample points, and that turns the relation into a uniform lower bound: the step of a minor-arc argument that normally costs a quantitative equidistribution estimate costs nothing.

Parts I and II [4, 5], and the first half of this Part, fix two things at once — the set A , and the uniform measure on its subsets. The rest of this Part moves each. Replacing the uniform measure by a Bernoulli(q) product measure deforms the gap series of Part I into an explicit family $\Gamma^{(q)}$, and we prove that the fair coin minimises it for every set, with equality only at $q = \frac{1}{2}$; the argument is termwise and is machine-checked in Lean 4. We derive the corresponding transfer function $\Phi^{(q)}$ and show that the bias splits the symmetric pair of strata of §5 in the ratio $q : (1 - q)$, so that the biased landscape responds to a shift of target at first order; the resulting slope is a point prediction with no free parameter, and it is confirmed against the natural alternative to within 0.4%. Replacing the primes by a random odd sequence, we prove that the extremal modulus of the minor arcs moves from 6 to 4, so that the peak height available to a random odd sequence is $1/\sqrt{2}$ where the primes have $\sqrt{3}/2$ — the primes are the harder case. We then prove that the geometric *rate* follows: almost surely the minor-arc maximum of $F_A^{1/b}$ tends to $1/\sqrt{2}$. The proof carries no ineffective constant and needs no equidistribution input, because the coset identity replaces the usual discrepancy estimate by an equality. Two questions inherited from the transfer-function strand are also settled or sharpened. Every claim is tagged with its status, and the numerical work is reproducible from logged scripts.

*Independent researcher. amkn.sub03@gmail.com, <https://mathlab0911.github.io/>, ORCID <https://orcid.org/0009-0000-0890-4395>

1 Introduction

Parts I and II and the transfer-function strand fixed two things at once: the sequence A , and the measure on its subsets. Every statement of the trilogy is about a specific A under the uniform measure. This paper moves each of them, one at a time, and the two moves turn out to be worth making for different reasons.

Bias. The classification theorem of Part I [4] says which elements lie on which side of a target. It says nothing about how configurations are weighted, so the entire combinatorial layer survives verbatim when the uniform measure is replaced by the Bernoulli(q) product measure, and the gap series deforms into an explicit one-parameter family $\Gamma^{(q)}$ (§8). The answer has a shape worth putting in the introduction: **the fair coin gives the least rugged landscape**, for every sequence, by a termwise convexity argument that is machine-checked (Theorem 8.5, Remark 8.4). The transfer function deforms too, and here the deformation is not cosmetic: the transfer-function strand's cosh splits into two unequal exponentials with offsets in the ratio $q : (1 - q)$, so the biased landscape responds to a shift of target at *first* order rather than second (§9). Rigidity off the fair coin is a tilt, not a plateau.

Coarse-graining. §3 is one lemma and it is the technical spine of the paper. Averaging the energy $X(t) = -\log |\cos \pi t|$ over the coset $\{t + k/v\}$ returns X at v times the frequency, plus $(1 - 1/v) \log 2$, exactly (Lemma 3.1, proved by the multiplication formula and again by Fourier). Since $X \geq 0$, the rational points are therefore the *minima* of the coset average, not merely convenient sample points (Corollary 3.2) — so the step of a minor-arc argument that normally costs a quantitative equidistribution estimate costs nothing here. Read as physics it is a block-spin step with an exact fixed point; we say so once, in Remark 3.5, and deduce nothing from the reading.

Randomness. For a random odd sequence the minor-arc input that cost Part II [5] fifty rounds is replaced by a concentration bound, and §10 carries *no ineffective constant at all* — the first such statement in the series. What we did not expect is that the random case is also *easier*: the extremal modulus moves from 6 to 4 (Theorem 10.5), so the peak height drops from $\sqrt{3}/2$ to $1/\sqrt{2}$, and the geometric rate improves in step (Theorem 10.19). The modulus moves for a one-line arithmetic reason — a set of primes contains at most one element $\equiv 3 \pmod{6}$, a random odd set contains a third of them — and the consequence is that Part II's $\sqrt{3}/2$ is a property of the primes, not a limitation of the method. *The primes are the hard case.*

A methodological remark that this paper earned twice. Deformations are not only a way to generalise; they are a way to *see*. **A deformation that breaks a symmetry makes measurable everything the symmetry was hiding.** Constructing $\Gamma^{(q)}$ and $\Phi^{(q)}$ audited the transfer-function strand for free, twice. First, the c_d fork: asking how the omitted factor behaves as k grows, rather than how large it is at one k , closed a question the transfer-function strand had left open (§11). Second, and more sharply, the sign of λ : the untilted Φ of §5 is an *even* function of λ , so no observable in that paper depends on the sign, and a definition that carried the wrong one survived three papers and an external referee. Splitting the cosh made the sign observable in a single measurement. We record the principle because it is a design rule, not an anecdote: when a quantity cannot be checked, deform the problem until it can.

Verification, and the genre. Every numbered claim carries a status tag: proved, Lean-verified, experimentally confirmed with the range stated, or conjectured. The scripts that produce every number in this paper write logs beside themselves and are part of the artefact. This is now a recognisable genre — machine-assisted mathematics shipped with formal certificates, of which *Ten advances in mathematics and theoretical computer science* [1] and its accompanying Lean 4 repository are the visible precedent. The direction differs: that work attacks named open problems, whereas the present series builds a small theory from its definitions upward, so the certificates here anchor *definitions and structural lemmas* rather than headline resolutions. The verification ethos is the same one, and we take it as read rather than arguing for it.

1.1 The transfer-function strand

Parts I and II [4, 5] of this series establish that for the odd primes $\text{lm}/r \rightarrow \Gamma$ *at the centre* of the landscape. Whether Γ is an invariant of the landscape or an accident of its most symmetric point is the question this paper answers, and the answer turns out to be stronger than rigidity alone: the whole ratio is a function of one local observable. The argument and what each part rests on:

- Parts I and II establish $\text{lm}/r \rightarrow \Gamma$ *at the centre* $n = \lfloor T/2 \rfloor$: Part I reduces it to a flatness hypothesis and Part II discharges that hypothesis unconditionally. The obvious question — whether Γ is a property of the landscape or of its centre — is what this paper answers.
- The combinatorial layer (classification, exact stratification, sandwich) was proved *for every target* in Part I and is formally verified; nothing there needs redoing. [Lean-verified]
- The minor-arc estimates of Part II [5] bound $|G_B(\theta)|$ as a function of θ only. They do not refer to the extraction point m , which enters $r_B(m) = \frac{1}{2\pi} \int F_B(\theta) e^{-im\theta} d\theta$ only through a phase. They therefore transfer to off-centre targets unchanged. [quoted from Part II]
- What is new is the middle, §5.2: an additive bridge (Lemma 5.8) carrying the minor-arc estimates from the uniform measure to the exponential tilt, the minor-arc input in the form the bridge needs (Lemma 5.10), and a tilted local limit theorem uniform over the extraction window (Proposition 5.20). The odd orders in λ cancel by the reflection symmetry of the landscape, which is where cosh enters.
- One consequence of quoting rather than reproving deserves saying at the outset: *the only ineffective constant in the trilogy is one and the same Siegel–Walfisz constant, and it enters each paper at exactly one point* — here, part (b) of Lemma 5.10. See §14.

1.2 Terminology

The vocabulary of Part I is tabulated in its §1.2 and not repeated. Four terms are local to this Part:

2 Notation, and what is inherited

The definitions are those of Part I [4], written so that nothing assumes $n = T/2$. Fix throughout an increasing sequence $A = (a_1, \dots, a_k)$ of odd positive integers, $T = \sigma(A)$,

term	what it denotes
<i>window identity</i>	Part I's equality between the window series W_D and the gap series Γ at the last informative D ; quoted here, not reproved.
<i>coset identity</i>	the identity for $X(t) = -\log \cos \pi t $ displayed in the abstract: the average of X over a coset of index v is X at v times the frequency, plus a constant. This is the classical Kubert distribution relation [3] for $\log 2 \sin \pi t $, shifted by $\frac{1}{2}$; we use the descriptive name because the shifted form is what every statement below needs, and not to suggest the relation is ours. Equivalently, the family $\{X(v \cdot)\}$ is closed under <i>coarse-graining</i> — averaging over a coset — exactly, with no error term; that phrase is used below in this sense and in no other.
<i>transfer function</i> Φ	the function of one variable through which the ratio $\text{lm}_A(n)/r_A(n)$ depends on the target, once the target leaves the centre.
<i>logarithmic slope</i> λ	the local logarithmic derivative of the representation count, $\lambda(n) = \log(r_A(n+1)/r_A(n-1))/2$; the one local observable through which the transfer function acts.
<i>the tilt</i>	the exponential change of measure P_s under which a subset S has probability proportional to $e^{s\sigma(S)}$; the standard exponential tilt of large-deviation theory, applied here subset by subset.
<i>biased invariant</i> $\Gamma^{(q)}$	what the gap series of Part I becomes when the uniform measure on subsets is replaced by a Bernoulli(q) product measure.

$S_2 = \sum a_i^2$, $V = S_2/4$, and a target n . Write $\rho = n/T$ and $x = \frac{1}{2} - \rho$. For $d \geq 1$ put

$$N_d = \#\{i : a_i \leq 2d\}, \quad \sigma_d = \sum_{a_i \leq 2d} a_i, \quad s_d = \sum_{a_i \leq 2d} a_i^2, \quad \delta_d = d + \frac{\sigma_d}{2},$$

and let $B_d = \{a_i : a_i > 2d\}$ be the layer- d ensemble, of variance $V_d = (S_2 - s_d)/4 = V - s_d/4$.

Definition 2.1 (local logarithmic slope). $\lambda(n)$ is the logarithmic derivative of the representation count at the target,

$$\lambda(n) = +\frac{d}{dm} \log r_A(m) \Big|_{m=n} = \frac{1}{4} \log \frac{r_A(n+2)}{r_A(n-2)} + O(\cdot),$$

so $\lambda > 0$ below the centre, where r_A is increasing.

Remark 2.2 (on the sign, and why it took the biased deformation to matter). An earlier version of this definition carried the opposite sign, which made it disagree with Remark 2.3 (where $\lambda \approx (T/2 - n)/V > 0$ below the centre) and with Remark 2.5 (where $s/\lambda \rightarrow 1$ with $s > 0$ below the centre). Nothing in this paper changed when the sign was corrected, and nothing could have: Φ is an *even* function of λ , so the sign of λ is not observable anywhere in the transfer function. It becomes observable only when the cosh is split into two unequal exponentials, which is what happens under a biased measure; the discrepancy was found there and is recorded here. The convention above is the one that makes $\lambda \approx s$.

Remark 2.3 (why the slope and not the offset). The Gaussian approximation $\lambda \approx (T/2 - n)/V$ is what one writes down first, and it is wrong at the 10% level for the purposes of this paper. Substituting the measured λ for the Gaussian one takes the drift of the fitted coefficient across the range $x \in [0.06, 0.30]$ from 15–18% down to 0.03–0.07%, and it dissolves a non-integer exponent that an earlier draft of this programme had reported as a finding: the exponent was an artefact of the wrong variable. This is not a cosmetic choice: every statement below is false at the stated precision if $(T/2 - n)/V$ is substituted for λ .

Definition 2.4 (the tilt). P_s is the product measure on subsets under which a_i is present with probability $p_a = (1 + e^{sa})^{-1}$, and $s = s(n) > 0$ is chosen so that $\sum_a a p_a = n$; thus P_s is the exponential tilt of the uniform measure whose mean sits at the target.

Remark 2.5 (the tilt and the slope agree to first order). s and λ are the same quantity to leading order but not exactly: s makes the tilted *mean* equal n , whereas λ is the slope of $\log r_A$ at n , and the two differ by the mean–mode gap. Measured on three profiles,

$$\frac{s(n)}{\lambda(n)} = 1 + \frac{c_A}{k} + O(k^{-2}), \quad c_A \approx 1.85 \text{ (odds), } 2.03 \text{ (primes), } 2.87 \text{ (squares),}$$

stable in x and clean in k over $k = 100, \dots, 220$. [\[experimentally confirmed; tilt_r088\]](#) The statements below are formulated in λ , and *not* in s : substituting s multiplies the residual of Table 1 by three, exactly as the λ^2 leading behaviour of Φ predicts. [\[experimentally confirmed; sres_r088\]](#)

3 An exact coset identity

Parts I and II and the transfer-function strand borrowed the vocabulary of statistical physics — landscape, energy, transfer function, tilt — and used it as a source of the right questions. This section is the first place where the borrowing runs the other way and produces an exact recursion rather than an analogy. Everything later in the paper that concerns minor arcs (§10.1, §10.2) is a corollary of it.

Throughout write

$$X(t) = -\log |\cos \pi t| \in [0, \infty],$$

the per-element energy of Parts I and II and the transfer-function strand at phase t .

Lemma 3.1 (coset identity; the Kubert distribution relation, shifted). [\[classical. This is the distribution relation for \$\log |2 \sin \pi t|\$ \[3\] with \$t \mapsto t + \frac{1}{2}\$; we claim no part of it. What this Part contributes is Corollary 3.2 and the use made of it in §10.2.\]](#) For every integer $v \geq 1$ and every real t ,

$$\frac{1}{v} \sum_{k=0}^{v-1} X\left(t + \frac{k}{v}\right) = \left(1 - \frac{1}{v}\right) \log 2 + \frac{1}{v} X(vt + \tau_v), \quad \tau_v = \begin{cases} 0, & v \text{ odd}, \\ \frac{1}{2}, & v \text{ even}. \end{cases}$$

Equivalently $\prod_{k < v} |\cos \pi(t + k/v)| = 2^{1-v} |\cos \pi(vt + \tau_v)|$.

First proof (multiplication formula). Write $2 \cos \pi z = 2 \sin \pi(z + \frac{1}{2})$ and apply $\prod_{k < v} 2 \sin \pi(y + \frac{k}{v}) = 2 \sin(\pi v y)$ — the factorisation of $x^v - 1$ read on the unit circle — with $y = t + \frac{1}{2}$. The product of the $2|\cos|$ is $2|\sin \pi(vt + \frac{v}{2})|$, which is $2|\cos \pi vt|$ for v odd and $2|\sin \pi vt| = 2|\cos \pi(vt + \frac{1}{2})|$ for v even. Divide by 2^v and take logarithms. $\square$

Second proof (Fourier). From $-\log |2 \sin \pi u| = \sum_{n \geq 1} \cos(2\pi n u)/n$ and $u = t + \frac{1}{2}$,

$$X(t) = \log 2 + \sum_{n \geq 1} \frac{(-1)^n \cos(2\pi n t)}{n}.$$

Averaging over the coset annihilates every harmonic but those with $v \mid n$, since $v^{-1} \sum_{k < v} \cos(2\pi n(t + k/v))$ is $\cos(2\pi n t)$ when $v \mid n$ and 0 otherwise. Writing $n = vp$ and using $(-1)^{vp} = (-1)^p$ for v odd and 1 for v even, the surviving series is $v^{-1}(X(vt + \tau_v) - \log 2)$. $\square$

Two proofs check the *statement*; one proof checks only the argument. The identity is verified in product form for $v \leq 60$ at 200 random shifts each, worst discrepancy $2.5 \cdot 10^{-16}$, and the two Fourier partial sums agree to $8.8 \cdot 10^{-14}$. [\[proved, twice; and checked numerically, stab_r110, rate_r112\]](#)

Corollary 3.2 (the rational points are the minima). $\frac{1}{v} \sum_{k < v} X(t + \frac{k}{v}) \geq (1 - \frac{1}{v}) \log 2$ for every t , with equality exactly when $vt + \tau_v \in \mathbb{Z}$.

Proof. $X \geq 0$. □

Remark 3.3 (what a surrogate costs, before any equidistribution input). [\[proved \(both inequalities and all three means\); the ratios are arithmetic. Measured for confirmation only, leadc_r143.\]](#)

A minor-arc argument for a generating function $\prod_{a \in A} (1 + e(a\theta))$ needs a lower bound on the deficiency $D(\theta) = \sum_{a \in A} X(a\theta)$, since $|\prod_{a \in A} (1 + e(a\theta))| = 2^{|A|} e^{-D(\theta)}$. The standard route replaces X by a surrogate that a discrepancy estimate can handle. Two are available, and both are valid:

$$X(t) \geq \frac{1}{2} \sin^2 \pi t, \quad X(t) \geq \frac{\pi^2}{2} \|t\|^2,$$

the first from $-\log c \geq 1 - c$ with $c = \cos^2 \pi t$, the second because $-\log \cos u = \frac{u^2}{2} + \frac{u^4}{12} + \frac{u^6}{45} + \dots$ has only positive coefficients on $|u| < \pi/2$. Both are equalities only at $t = 0$.

Each surrogate caps the bound it can ever produce at its own mean over the circle, and those means are not the same number:

mean over the circle of	value	factor lost	
$X(t)$ itself	$\log 2$	0.69315	1
$\frac{1}{2} \sin^2 \pi t$	$1/4$	0.25000	2.773
$\frac{\pi^2}{2} \ t\ ^2$	$\pi^2/24$	0.41123	1.686

The loss is in the surrogate, not in the equidistribution input. Hand the surrogate routes their exponential sum and their discrepancy exactly — no Weyl, no Koksma, no error term — and they still stop at $\frac{1}{4}|A|$ and $\frac{\pi^2}{24}|A|$ where Proposition 10.8 gives $|A| \log 2$. The reason is visible in the two displays: X has a logarithmic singularity at the half-integers and a bounded convex surrogate does not, and that singularity is where the mass of the average sits. Since the bound enters an exponent, a factor 2.77 in the constant is a factor 2.77 in the exponent.

Remark 3.4 (what is machine-checked in Lemma 3.1, and what is not). The identity is *multiplicative in v* , and that is what has been formalised. Writing $k = i + aj$ and inducting on b splits a product over $\{k < ab\}$ into a double product with no bijection and no analysis — `prod_range_mul_index`, stated for an arbitrary commutative monoid. On top of it, `cosetId_one` and `cosetId_two` are the base cases ($v = 2$ is the double-angle formula and needs nothing else), and `cosetId_mul` carries the identity from a and b to ab , including the offset bookkeeping $a\tau_b + \tau_a \equiv \tau_{ab} \pmod{1}$ in all four parity cases, via `abs_cos_add_int_mul_pi`. *What is not formalised is the base case at an odd prime*, which is the only place a product over roots of unity is needed, and which Mathlib does not supply. So the machine-checked statement is: *the identity holds at 1 and 2 and is closed under products*, hence holds for every v built from the primes at which it is established by hand. The reduction was checked numerically before it was formalised, including a deliberately wrong τ as a control. [\[Lean-verified \(the reduction; not the odd-prime base case\); Pnp/Theory/CosetProd.lean, coset_mult_r116\]](#)

Remark 3.5 (what this says, and in which language). The identity has two readings and they have different statuses; the paper keeps them apart.

As *arithmetic* (proved) it says that averaging X over a coset of index v returns X *itself*, at v times the frequency, plus a constant — so the family $\{X(v \cdot)\}$ is closed under coarse-graining,

exactly, with no error term. That is what Corollary 3.2 turns into a uniform lower bound, and it is why §10.2 needs no equidistribution input: the question “how far can the mean fall when the rational grid is shifted?” has the answer *it cannot fall at all*.

As physics (an interpretation, and nothing is deduced from it) this is a block-spin step. Coarse-graining the phase lattice by a factor v maps the energy to a rescaled copy of itself; the map has an exact fixed-point structure rather than an approximate one, and the constant $(1 - 1/v) \log 2$ is the free energy released per coarse-graining step, additive over successive steps because τ composes. Parts I and II and the transfer-function strand treated the landscape picture as a source of questions; this is the first identity in the series that the picture predicts and the arithmetic confirms. §13 says what we would want to do with it and does not do it here.

4 Rigidity

Theorem 4.1 (rigidity of the gap series). *[proved; unconditional. The analytic ingredients are proved in §5.2; the Edgeworth expansion in region R1 of Proposition 5.20 — for a long time the one thing that was missing — is written out in Appendix A with every constant proved, and it has had the independent reading this project requires before an argument counts as proved — delivered in three parts, each covering the text as it then stood: the three lemmas line by line (r162), the three repairs landed against them (r164), and the restated proposition with the T_* construction and all five explicit constants of ρ , each re-derived from scratch (r171). See the head of Appendix A for why that is stated as three parts and not as one reading. Problem 12.10 is closed and this theorem is conditional on nothing. Its algorithmic interpretation is a different matter: the reading of §13 carries an assumption about the search which is not a hypothesis of this theorem and is named there.]* Let $P = P_k$ be the first k odd primes and let $\rho \in (0, 1)$ be fixed. Then

$$\frac{\text{lm}_P(\lfloor \rho T \rfloor)}{r_P(\lfloor \rho T \rfloor)} \xrightarrow{k \rightarrow \infty} \Gamma(P),$$

uniformly for ρ in compact subsets of $(0, 1)$. For a general profile A the same holds under hypothesis (H) of §7.

$\Gamma(A)$ is therefore an invariant of the landscape and not of its centre. Theorem 4.1 is the case $\lambda \rightarrow 0$ of Theorem 5.1: at fixed ρ one has $\lambda(n) \asymp x/N \rightarrow 0$, and Φ is continuous at 0 with $\Phi(0) = \Gamma$, so

$$\Phi(\lambda) - \Gamma = O\left(\lambda^2 \sum_{d \geq 1} 2^{1-N_d} \delta_d^2\right) \rightarrow 0,$$

the sum being finite precisely by (H). The uniformity in ρ needs no new estimate: every bound in §5.2 is uniform once $t_{\min}(\rho)$ is bounded below, and $t_{\min}$ is continuous and positive on $(0, 1)$, hence bounded below on compact subsets (Remark 5.19).

[experimentally confirmed] The evidence the theorem is modelled on. With $\text{ratio}(\rho, k) = \lfloor \text{lm}/r \rfloor / \Gamma(P_k)$, the spread over ρ contracts monotonically:

k	40	52	64	76	88	100
$\max_\rho - \min_\rho$	0.01051	0.00283	0.00257	0.00162	0.00113	0.00083
$\max_\rho \text{ratio} - 1 $	0.00921	0.00276	0.00244	0.00156	0.00109	0.00081

over $\rho \in [0.20, 0.50]$, by exact integer computation. Three guards are in force throughout this paper’s computations and are worth stating once: every reported point satisfies $r \geq 10^4$, so that integer granularity cannot masquerade as signal; the dynamic programme was

checked against brute-force enumeration over all 2^k subsets at small k ; and the exact identity $\text{lm}(n) = \text{lm}(T - n)$ is verified at every measured target, which tests the stratification and the arithmetic together.

5 The transfer function

Theorem 5.1 (transfer). *[proved; unconditional. The analytic ingredients are proved in §5.2; the Edgeworth expansion in region R1 of Proposition 5.20 — for a long time the one thing that was missing — is written out in Appendix A with every constant proved, and it has had the independent reading this project requires before an argument counts as proved — delivered in three parts, each covering the text as it then stood: the three lemmas line by line (r162), the three repairs landed against them (r164), and the restated proposition with the T_* construction and all five explicit constants of ρ , each re-derived from scratch (r171). See the head of Appendix A for why that is stated as three parts and not as one reading. Problem 12.10 is closed and this theorem is conditional on nothing. Its algorithmic interpretation is a different matter: the reading of §13 carries an assumption about the search which is not a hypothesis of this theorem and is named there.]* Let $A = P_k$ be the first k odd primes, or a general profile satisfying (H), and let $n = \lfloor \rho T \rfloor$ with $\rho \in (0, 1)$ fixed. Then

$$\frac{\text{lm}_A(n)}{r_A(n)} = \Phi_k(\lambda(n)) + E_k, \quad \Phi_k(\lambda) = 1 + \sum_{d \geq 1} 2^{1-N_d} e^{-\lambda^2 s_d/8} \cosh(\lambda \delta_d), \quad (1)$$

with $E_k \rightarrow 0$ as $k \rightarrow \infty$, uniformly for ρ in compact subsets of $(0, 1)$.

Remark 5.2 (what the theorem does and does not claim about E_k). The assertion is $E_k \rightarrow 0$ in absolute terms, nothing sharper. The measured rate $E_k \asymp c_A/k$ of Remark 5.7 is *not* part of the statement; it is quoted as [experimentally confirmed, $k \leq 220$, three profiles] and stated as an open problem in §14. Distinguishing the two matters here more than usual, because the numerical agreement is far better than what is proved.

Three features of (1) deserve to be stated before any proof.

Remark 5.3 (the ratio sees the target only through one number). Φ depends on A alone. The target enters only through $\lambda(n)$. This is a stronger statement than rigidity: rigidity says the limit does not move, (1) says how it fails to move.

Remark 5.4 ($\lambda = 0$ recovers Part I [4]). The layer weight $e^{-\lambda^2 s_d/8}$ is 1 at $\lambda = 0$, so $\Phi(0) = 1 + \sum_{d \geq 1} 2^{1-N_d}$, which is the window identity $W_D(A) = \Gamma(A) + (2D + 1)2^{-k}$ of Part I — already formally verified. Numerically, $\Phi(0) - \Gamma(A)$ is 0, $-4.7 \cdot 10^{-14}$ and $-2.7 \cdot 10^{-15}$ for the three sequences of Table 1.

Remark 5.5 (each layer is a smaller ensemble). The factor $e^{-\lambda^2 s_d/8}$ is not a fudge: it is what remains when the Gaussian normalisations of B_d and of A fail to cancel, because $V_d = V - s_d/4$. Expanding to second order gives the coefficient of λ^2 as $\delta_d^2 - s_d/4$ rather than δ_d^2 . In the notation of §2, layer d is the ensemble B_d with the small elements removed, so its variance is $V_d = V - s_d/4$ rather than V ; the weight $e^{-\lambda^2 s_d/8}$ is exactly the price of that difference, and Remark 5.21 derives it from the ratio of partition functions rather than fitting it. The correction is termwise smaller than the term it corrects, $s_d/4 \leq d\sigma_d/2 < \delta_d^2$, so every term of the series stays positive and the correction never overshoots.

5.1 Verification

[experimentally confirmed] Table 1 evaluates (1) at the measured $\lambda(n)$ against the exactly computed lm/r .

$x = \frac{1}{2} - \rho$	0.06	0.08	0.10	0.15	0.20	0.25	0.30
odds, $k = 220$	0.84%	0.85%	0.85%	0.86%	0.88%	0.92%	0.97%
primes, $k = 220$	0.91%	0.92%	0.92%	0.94%	0.96%	0.99%	1.05%
$a_i \asymp i^{3/2}$, $k = 220$	1.05%	1.06%	1.06%	1.08%	1.10%	1.14%	1.20%
squares, $k = 170$	1.66%	1.67%	1.68%	1.70%	1.74%	1.80%	1.89%

Table 1: Residual $|\text{measured} - \Phi(\lambda)|$ **as a fraction of the centred signal** $\text{lm}/r - [\text{lm}/r]_{n=T/2}$. Absolute residuals are 10^{-5} to 10^{-6} ; quoting them as significant figures of the ratio would be meaningless, since the signal itself lives in the fifth and sixth digit. The third row is the only ensemble in this paper that was *constructed to test a prediction* rather than chosen because it was natural: $\alpha = 3/2$ was built after Remark 5.7 had been fitted on the other three, and its predicted constant $16/7$ lies between theirs.

Remark 5.6 (why the residual is quoted this way). Against the *uncentred* signal the same residuals read 2% to 300%, the large values occurring precisely where the signal is smallest. Neither number is wrong; the centred one is the one the theorem is about, because $\Phi(0) = \Gamma$ is an identity and the finite- k offset $[\text{lm}/r]_{n=T/2} - \Gamma$ is a separate quantity of size 10^{-6} . We keep this remark because an earlier draft of this programme reported “agreement to five or six significant figures” for a model that was 12% wrong on the signal: the figures were real and the claim was not, because the signal itself lives in those digits.

Remark 5.7 (the residual obeys a $1/k$ law). Table 1 is one column of a k -scan. Writing $R_A(k)$ for the residual in the sense of the caption, the measurement over $k = 60, \dots, 220$ gives

$$k R_A(k) \longrightarrow \begin{cases} 1.85\% & \text{odds } (k = 60, 100, 140, 180, 220 : 2.29, 2.00, 1.91, 1.87, 1.85), \\ 2.01\% & \text{primes } (k = 100, \dots, 220 : 2.38, 2.06, 2.03, 2.01), \\ 2.82\% & \text{squares } (k = 130, \dots, 190 : 3.78, 2.85, 2.83, 2.82). \end{cases}$$

Two things follow. First, the residual is $O(1/k)$ and therefore vanishes, which is the form Theorem 5.1 asserts; the table’s 1% is a finite- k number and not a floor. Second, the squares row is the loosest of the three because c_A is larger, not because the decay is different: squares obey the same law with a constant 1.52 times that of the odd numbers. The constants c_A coincide numerically with those of Remark 2.5 — the same mean–mode gap seen twice — but substituting s for λ does not remove the residual, so the coincidence is a shared order of magnitude and not a shared cause. [\[experimentally confirmed; tilt_r088, sres_r088\]](#)

What they share is a driver. The saddle-point representation gives the exact identity

$$s(n) - |\lambda(n)| = \frac{K'''(s)}{2K''(s)^2}, \quad K'' = \sum a^2 p_a (1 - p_a), \quad K''' = \sum a^3 p_a (1 - p_a) (1 - 2p_a),$$

tested pointwise to a relative error $0.54/k$ [\[experimentally confirmed; saddle_r092\]](#), and for small tilt $K'''/(2K''^2) \approx sS_4/S_2^2$, so that $c'_A = k(s/|\lambda| - 1) \rightarrow kS_4/S_2^2$. For $a_i \asymp i^\alpha$ this is $(2\alpha + 1)^2/(4\alpha + 1)$. The measured c_A of this remark agrees with c'_A to within 1% on four profiles:

	α	kS_4/S_2^2	c'_A	c_A	c_A/c'_A
odd numbers	1	1.8000	1.8525	1.8674	1.008
$a_i \asymp i^{3/2}$	3/2	2.2794	2.3494	2.3318	0.993
squares	2	2.7714	2.8615	2.8377	0.992
primes	—	1.9731	2.0251	2.0253	1.000

at $k = 220$. The two halves of this have different statuses and the difference matters:

- $c'_A \rightarrow (2\alpha + 1)^2/(4\alpha + 1)$ is *derived* — the identity above is exact, its sign is forced by $\Lambda''' < 0$ for $\rho < \frac{1}{2}$, and the small-tilt step is a Taylor expansion — and [\[experimentally confirmed on five profiles, one of them out of sample\]](#).
- $c_A = c'_A(1 + o(1))$ is [\[experimentally confirmed: four profiles, ratios 0.992–1.008, \$k = 220\$ \]](#). The exact asymptotic equality, and the mechanism behind it, remain a **conjecture**. What is established is a shared driver, S_4/S_2^2 , and *not* a causal link: substituting s for λ makes the residual three times worse, not smaller.

Carried to $k = 300$, the identification tightens rather than drifting, and it does so at a measurable rate. The ratio c_A/c'_A reads

k	120	180	220	260	300
odd numbers	1.0589	1.0244	1.0151	1.0098	1.0064
primes	1.0318	1.0084	1.0039	1.0014	1.0001
$a_i \asymp i^{3/2}$	1.0113	1.0015	0.9995	0.9984	—
squares	1.0782	1.0039	0.9985	—	—

so $|c_A/c'_A - 1|$ falls by a factor 9.2 across a factor 2.5 in k for the odd numbers, an exponent of 2.4 — the gap between the two constants closes *faster* than the $1/k$ law they both sit on. [\[experimentally confirmed, four profiles, \$k \leq 300\$; e8_r103\]](#) This is a sharper statement of the same conjecture, not a proof of it: a rate of empirical convergence is evidence about a limit, and the mechanism is still missing.

5.2 The bridge to Part II

Both theorems reduce to the same five parts. **P1** is the exact combinatorial stratification, proved for every target in Part I [\[Lean-verified\]](#); **P3** is a local limit theorem for the tilted ensemble of Definition 2.4; **P4** is the assembly, in which the odd orders in λ cancel by the reflection symmetry $\text{Im}(n) = \text{Im}(T - n)$ [\[Lean-verified\]](#); **P5** is the error budget, which converges exactly under (H). Only **P2** is new: the minor-arc estimates of Part II [\[5\]](#) are proved for the *uniform* measure, and **P3** needs them for the tilted one.

Write $\tilde{G}_B(\theta) = \prod_{a \in B} (1 - p_a + p_a e(a\theta))$ for the tilted generating function and $S_B(\theta) = \sum_{a \in B} e(a\theta)$ for the linear exponential sum.

Lemma 5.8 (the bridge). *Put $t_a = 4p_a(1 - p_a)$ and $t_{\min} = \min_{a \in B} t_a$. For every θ ,*

$$|\tilde{G}_B(\theta)|^2 = \prod_{a \in B} (1 - t_a \sin^2(\pi a\theta)) \leq \exp\left(-t_{\min} \sum_{a \in B} \sin^2(\pi a\theta)\right) = \exp\left(-\frac{t_{\min}}{2}(b - \text{Re } S_B(\theta))\right).$$

Consequently, if $\text{Re } S_B(\theta) \leq \eta b$ then $|\tilde{G}_B(\theta)| \leq \exp(-t_{\min}(1 - \eta)b/4)$.

Proof. The identity $|1 - p + pe(\phi)|^2 = 1 - 4p(1 - p) \sin^2(\pi\phi)$ is a computation; $t_a \geq t_{\min}$ and $1 - u \leq e^{-u}$ give the middle step; and $\sum_a \sin^2(\pi a\theta) = \frac{1}{2}(b - \text{Re } S_B(\theta))$ is the double-angle formula summed. $\square$

Remark 5.9 (why the bound is additive, and a bound that is not available). It is tempting to transfer multiplicatively, as $|\tilde{G}_B(\theta)| \leq |G_B(\theta)|^\kappa$ with G_B the untilted product, since that would carry every exponent of Part II across at once. **No such $\kappa > 0$ exists.** At any θ with $a\theta \in \frac{1}{2} + \mathbb{Z}$ for some $a \in B$ the untilted factor $\cos(\pi a\theta)$ vanishes, so the right-hand side is

0, while the tilted factor equals $\sqrt{1-t_a} > 0$. Equivalently $\log(1-ty)/\log(1-y) \rightarrow 0$ as $y \rightarrow 1$. The underlying inequality $1-ty \leq (1-y)^t$ is false for every $t \in (0,1) \rightarrow s \mapsto (1-y)^s$ is convex and therefore lies *below* its chord $1-sy$, with deficit exactly $-(1-t)$ at $y=1$. Lemma 5.8 avoids the issue by never dividing by a quantity that vanishes. [\[refuted numerically on a 40575-point grid and structurally; kappa_r088\]](#)

Lemma 5.10 (the minor-arc input). [\[\(a\) proved, Lemma 5.11; \(b\) proved, Lemma 5.12; \(c\) quoted from Part II\]](#) Let $A = P_k$ and $N = a_k$. Then

$$\operatorname{Re} S_A(\theta) \leq \left(\frac{1}{2} + o(1)\right)b \quad \text{uniformly for } \frac{1}{2N} \leq \theta \leq 1 - \frac{1}{2N}.$$

(a) $q \in \{1, 2\}$. Lemma 5.11 below.

(b) $3 \leq q \leq (\log N)^{A_0}$. Lemma 5.12 below.

(c) $q > (\log N)^{A_0}$. This is the case $n=1$ of Step 2 of `prop:deepminor` in Part II, which gives $|S_A(\theta)| \ll N(\log N)^{3-A_0/2} + N^{4/5}(\log N)^3$, i.e. $\operatorname{Re} S_A/b \ll (\log N)^{4-A_0/2} = o(1)$ for $A_0 > 8$. [\[quoted from Part II, verbatim\]](#)

Lemma 5.11 (the two degenerate denominators). Let $q \in \{1, 2\}$, $\theta = a/q + \beta$ with $\|\beta\| \geq 1/(2N)$. Then $\operatorname{Re} S_A(\theta) \leq \frac{1}{2}b$ for all large N .

Proof. Write $S_A(\theta) = \sum_{p \leq N} e(pa/q)e(p\beta)$. For $q=1$ the first factor is 1; for $q=2$ it is $(-1)^p = -1$ on every odd prime, so $S_A = -\sum_{p \leq N} e(p\beta)$ and it suffices in both cases to bound $|\sum_{p \leq N} e(p\beta)|$. Partial summation against $\pi(t) = \operatorname{li}(t) + O(te^{-c\sqrt{\log t}})$ gives

$$\sum_{p \leq N} e(p\beta) = \int_2^N e(t\beta) \frac{dt}{\log t} + O(Ne^{-c\sqrt{\log N}}(1 + |\beta|N)),$$

and integrating by parts once more, $|\int_2^N e(t\beta) dt / \log t| \ll 1/(|\beta| \log N)$. With $\|\beta\| \geq 1/(2N)$ this is $\ll N/\log N \asymp b$ with an absolute implied constant, and the constant is $\leq \frac{1}{2}$ once $N\|\beta\|$ exceeds an absolute threshold; below that threshold, i.e. for $\|\beta\| \asymp 1/N$, the integral is $(1+o(1))b \cdot \hat{\chi}(\beta N)$ with $\hat{\chi}$ the Fourier transform of the indicator of $[0, 1]$, whose real part is $\leq \frac{1}{2}$ off the principal arc. For $q=2$ the sign is reversed, so $\operatorname{Re} S_A \leq 0$ near $\theta = \frac{1}{2}$ — the helpful direction. No Siegel–Walfisz is used. $\square$

Lemma 5.12 (the small- q ring). Let $3 \leq q \leq (\log N)^{A_0}$, $(a, q) = 1$, and $\theta = a/q + \beta$ with $|\beta| \leq 1/(q\tau)$, $\tau = N(\log N)^{-A_0}$. Then

$$S_A(\theta) = \frac{\mu(q)}{\varphi(q)} W(\beta) + O(Ne^{-c\sqrt{\log N}}), \quad W(\beta) = \int_2^N e(t\beta) \frac{dt}{\log t},$$

and $|W(\beta)| \leq \operatorname{li}(N) = (1+o(1))b$. Consequently $\operatorname{Re} S_A(\theta) \leq (\frac{1}{2} + o(1))b$.

Proof. Sort the primes $p \leq N$ by their residue class. The $\omega(q) = O(\log q)$ primes dividing q contribute $O(\log q) = o(b)$, so

$$S_A(\theta) = \sum_{\substack{r \bmod q \\ (r, q)=1}} e\left(\frac{ra}{q}\right) \sum_{\substack{p \leq N \\ p \equiv r \pmod{q}}} e(p\beta) + O(\log q).$$

By Siegel–Walfisz, $\pi(t; q, r) = \operatorname{li}(t)/\varphi(q) + O(te^{-c\sqrt{\log t}})$ uniformly for $q \leq (\log N)^{A_0}$; partial summation against $e(t\beta)$ turns this into

$$\sum_{\substack{p \leq N \\ p \equiv r \pmod{q}}} e(p\beta) = \frac{W(\beta)}{\varphi(q)} + O(Ne^{-c\sqrt{\log N}}(1 + |\beta|N)),$$

and $|\beta|N \leq (\log N)^{A_0}/q$, so the error is $O(Ne^{-c\sqrt{\log N}}(\log N)^{A_0}) = O(Ne^{-c'\sqrt{\log N}}) = o(b)$: a fixed power of $\log N$ cannot compete with $e^{-c\sqrt{\log N}}$. Summing over r and using the Ramanujan sum $\sum_{(r,q)=1} e(ra/q) = c_q(a) = \mu(q)$ for $(a, q) = 1$ gives the display. Finally $\mu(q)/\varphi(q)$ is real and $\varphi(q) \geq 2$ for $q \geq 3$, so $\operatorname{Re} S_A(\theta) \leq (|\mu(q)|/\varphi(q))b(1+o(1)) \leq (\frac{1}{2}+o(1))b$; Lemma 5.17 says where the $\frac{1}{2}$ is attained. $\square$

Remark 5.13 (where the ineffectivity is). The Siegel–Walfisz step in Lemma 5.12 is the only inexplicit estimate in §5.2 — Lemmas 5.8, 5.11, 5.17 and the Vaughan quotation in Lemma 5.10(c) are all effective — and it is the same constant that Part II isolates as C_1 ; see Remark 14.1.

Remark 5.14 (what would make C_1 effective). Two escapes are worth recording, because C_1 is the only inexplicit quantity in the trilogy. *Under GRH*, the Siegel–Walfisz input of Lemma 5.12 may be replaced by $\pi(t; q, r) = \operatorname{li}(t)/\varphi(q) + O(t^{1/2} \log t)$ with an absolute implied constant, uniformly for $q \leq t^{1/2}$. The range Lemma 5.12 actually uses, $q \leq (\log N)^{A_0}$, is far inside that, so on GRH the constant becomes effective and every constant in the trilogy is explicit. *Unconditionally*, the ineffectivity is confined to the possible existence of a Siegel zero for a single modulus $q \leq (\log N)^{A_0}$: for every modulus with an explicit zero-free region the estimate is already effective. The inexplicit constant is therefore not a feature of the argument but a hostage to one hypothetical zero, and the range it is held over is only $(\log N)^{A_0}$, not N . Neither escape is used below; the theorems are stated unconditionally.

Remark 5.15 (the middle range, and why no constant enters). Inside $|\theta| < 1/(2N)$ Lemma 5.10 says nothing, and it does not have to: there $|a\theta| \leq \frac{1}{2}$ for every $a \in A$, so $\sin^2(\pi a\theta) \geq 4a^2\theta^2$ and Lemma 5.8 gives the quadratic bound $|\tilde{G}_B(\theta)| \leq \exp(-4t_{\min}V_a\theta^2)$ directly. Cutting at the LCLT-natural radius $\theta_0 = (\log k)/\sqrt{K_d''}$ of §5.3, this is $\leq \exp(-t_{\min}(\log k)^2)$ on the whole of $\theta_0 \leq |\theta| \leq 1/(2N)$ — superpolynomially below the density scale $(K_d'')^{-1/2}$, with no free constant anywhere. The crossover at $\theta = 1/(2N)$ survives as the boundary against Lemma 5.10, forced by $|a\theta| \leq \frac{1}{2}$ rather than chosen; measured, the quadratic bound has already reached $e^{-0.01b}$ by $\theta N = 0.18\text{--}0.42$ across three profiles. [\[experimentally confirmed; gap_r090\]](#)

Remark 5.16 (what part (c) does *not* need). Part II must control every harmonic $n\theta$ with $n \leq N_0$, and pays for it with a covering lemma and the parameter $Q_2 = (\log N)^{A_0+A'_0}$. Here only $n = 1$ occurs, so the dichotomy in Lemma 5.10 is on the single denominator $q_1(\theta)$ and no covering lemma is needed. The quotation in (c) is therefore valid on a strictly larger set of θ than Part II's type (ii).

Lemma 5.17 (the extremiser of the additive estimate). $\max_{q \geq 3} \frac{\mu(q)}{\varphi(q)} = \frac{1}{2}$, and the maximum is attained at $q = 6$ and at no other q .

Proof. $\varphi(q) = 2$ holds exactly for $q \in \{3, 4, 6\}$, where μ takes the values $-1, 0, +1$; so among these the maximum is $\mu(6)/\varphi(6) = 1/2$. Every other $q \geq 3$ has $\varphi(q) \geq 4$ and hence $\mu(q)/\varphi(q) \leq 1/4 < 1/2$. $\square$

Remark 5.18 (the same $q = 6$, twice). By Lemma 5.10(b) the worst point of the additive estimate is therefore $\theta = 1/6$, with $\operatorname{Re} S_A/b \rightarrow \frac{1}{2} = \cos 60^\circ$. Part II's multiplicative estimate has its worst point at the same q : $\sup_q M(q) = M(6) = \sqrt{3}/2 = \cos 30^\circ$. **Same extremiser, same cause, two functionals** — both say that the odd primes lie in $\pm 1 \pmod{6}$ and nowhere else, read once through $\prod |\cos|$ and once through $\sum \cos$. Measured on the odd primes below 10^7 , $\operatorname{Re} S_A(1/6)/b = 0.499998$ against the predicted $1/2$, and $q = 6$ is the maximiser over $2 \leq q \leq 16$ at each of $N = 10^5, 10^6, 10^7$. [\[experimentally confirmed; eta_r090\]](#)

Remark 5.19 (the honest form of the tilt constant). $t_{\min} = t_{\min}(\rho)$ is positive for every fixed $\rho \in (0, 1)$ and bounded below on compact subsets, which is all Lemma 5.8 needs. It is *not* governed by the approximation $sN \approx 6x$: that relation is the small- x asymptote only, and the measured ratio $sN/6x$ is 1.02–1.13 at $x = 0.10$ but 1.22–1.39 at $x = 0.30$. Measured values of $t_{\min}$: 0.91, 0.89, 0.91 at $x = 0.10$ and 0.36, 0.28, 0.34 at $x = 0.30$ for the odd numbers, squares and primes. [\[experimentally confirmed; tilt_r088, kappa_r088\]](#)

5.3 The tilted local limit theorem, and the assembly

Write $K_d'' = \sum_{a \in B_d} a^2 p_a (1 - p_a)$ and $\mu_d(s) = \sum_{a \in B_d} a p_a$ for the variance and mean of B_d under P_s , and $K_d^{(j)}$ for the higher cumulants.

Proposition 5.20 (tilted LCLT, uniform over the window). [\[proved; unconditional. R2 and R3 quote Lemmas 5.8 and 5.10; R1 is Proposition A.7 of Appendix A, where the expansion is written out with every constant proved. The cumulant budget is retained below because it is the form the consumers use, not because it is all that is available.\]](#) For every layer $d \leq D(k)$ and every m with $|m - \mu_d(s)| \leq W(k)$,

$$P_s[\sum_{B_d} = m] = \frac{1}{\sqrt{2\pi K_d''}} \exp\left(-\frac{(m - \mu_d)^2}{2K_d''}\right) (1 + O(\varepsilon)), \quad \varepsilon = \frac{|K_d'''|}{(K_d'')^{3/2}} + \frac{K_d^{(4)}}{(K_d'')^2} + \frac{W(k)|K_d''''|}{(K_d'')^2},$$

uniformly in d and m .

Proof sketch. Invert $\tilde{G}_{B_d}$ over $[-\pi, \pi]$ and split at $\theta_0 = (\log k)/\sqrt{K_d''}$ and at $1/(2N)$.

R1, $|\theta| \leq \theta_0$. Expand $\log \tilde{G}_{B_d}(\theta) = i\mu_d\theta - \frac{1}{2}K_d''\theta^2 - \frac{i}{6}K_d'''\theta^3 + \frac{1}{24}K_d^{(4)}\theta^4 + O(K_d^{(5)}\theta^5)$. Since $|K_d^{(j)}| \leq \sum_a a^j$, on $|\theta| \leq \theta_0$ every term beyond the quadratic is $o(1)$ and the standard Edgeworth argument gives the display with the stated ε ; the third term contributes the $W|K_d'''|/(K_d'')^2$ piece over a window of width W .

R2, $\theta_0 \leq |\theta| \leq 1/(2N)$. Remark 5.15: the contribution is $\leq \exp(-t_{\min}(\log k)^2)$, super-polynomially below $(K_d'')^{-1/2}$.

R3, $|\theta| \geq 1/(2N)$. Lemmas 5.8 and 5.10: $|\tilde{G}_{B_d}| \leq \exp(-t_{\min}(1 - \eta)b/4)$, exponentially small, against measure $O(1)$.

Uniformity in d : B_d differs from A in $N_d \leq N_{D(k)} = (\log k)^{O(1)}$ elements of size $\leq 2D(k)$, so every cumulant shifts by a relative $O(s_D/S_2) \rightarrow 0$ and one ε serves all layers. The lattice is of step 1 and the parity weighting is quoted from Part II. $\square$

Remark 5.21 (the μ_d -shift, and why the layer term is what it is). This is the step that produces Φ , so it is worth doing once in full. Since $\sum_{a \in A} a p_a = n$ and $p_a = \frac{1}{2} - \frac{sa}{4} + O((sa)^3)$ for the small elements,

$$\mu_d(s) = n - \sum_{a \leq 2d} a p_a = n - \frac{\sigma_d}{2} + w + \dots, \quad w := \frac{s s_d}{4},$$

so the two evaluation points of the stratification sit at

$$(n + d) - \mu_d = +\delta_d - w, \quad (n - d - \sigma_d) - \mu_d = -\delta_d - w. \quad (2)$$

Feeding (2) into Proposition 5.20, undoing the tilt through $r_{B_d}(m) = P_s[\sum_{B_d} = m] Z_{B_d}(s) e^{sm}$, and using $sd + s\sigma_d/2 = s\delta_d$ to pair the two points:

$$r_{B_d}(n + d) + r_{B_d}(n - d - \sigma_d) = 2 Z_{B_d} e^{sn - s\sigma_d/2} \frac{e^{-(\delta_d^2 + w^2)/2K_d''}}{\sqrt{2\pi K_d''}} \cosh\left(\delta_d\left(s + \frac{w}{K_d''}\right)\right).$$

Dividing by $r_A(n) \sim Z_A e^{sn} / \sqrt{2\pi K''}$ and using $Z_{B_d} e^{-s\sigma_d/2} / Z_A = 2^{-N_d} \prod_{a \leq 2d} \cosh(sa/2)^{-1} = 2^{-N_d} e^{-s^2 s_d/8} (1 + \dots)$,

$$\frac{r_{B_d}(n+d) + r_{B_d}(n-d-\sigma_d)}{r_A(n)} = 2^{1-N_d} e^{-\lambda^2 s_d/8} \cosh(\lambda \delta_d) \cdot e^{-\delta_d^2/2K_d''} (1 + O(\varepsilon)), \quad (3)$$

with $\lambda = s + w/K_d'' = s(1 + s_d/(4K_d''))$ — which is exactly the per-layer slope drift of $1 + O(s_d/S_2)$ measured earlier. **The first three factors of (3) are the general term of Φ in (1);** the odd orders in λ cancel because the two points are $\pm \delta_d$ about the same centre, which is the reflection symmetry $\text{lm}(n) = \text{lm}(T - n)$ in analytic dress. [\[Lean-verified \(the symmetry\)\]](#)

Remark 5.22 (the one factor Φ omits). (3) carries a fourth factor, $c_d := e^{-\delta_d^2/2K_d''}$, that Φ does not. It is λ -independent and equal to $1 - O(\delta_d^2/V)$, so it belongs to E_k and not to the transfer function. Measured at $k = 180$: the ratio of the left side of (3) to Φ 's term is c_d to six decimal places at $d = 1$ and $d = 3$ (1.000000 and 1.000003 for the odd numbers, 1.000000 and 1.000000 for the primes), and departs from 1 only at $d = D(k)$, where the layer weight is $2^{-29} \approx 2 \cdot 10^{-9}$. Carrying c_d inside Φ would move Table 1 from 0.842% to 0.814% (odd numbers, $k = 220$, $x = 0.06$), from 0.914% to 0.902% (primes) and from 1.052% to 1.048% ($\alpha = 3/2$): the factor is real, it moves the residual the right way, and it accounts for at most 3% of it. We therefore keep Φ as stated and count c_d in E_k . [\[experimentally confirmed; p3_r096, phig_r096\]](#)

That choice was made at one value of k , and the level of the correction is the wrong thing to decide it on; the trend is the right thing. Since $1 - c_d \asymp \delta_d^2/2V_d$ with $V \asymp k^{2\alpha+1}$, and the series is dominated by the small d where $\delta_d = O(1)$, the c_d correction is $O(k^{-(2\alpha+1)})$ while the residual is $O(1/k)$: the *share* of the residual that c_d carries should fall like $k^{-2\alpha}$. It does. Writing that share as $(R_\Phi - R_{\Phi c})/R_\Phi$ and fitting successive k :

$$1.98 \text{ (odd numbers, } \alpha = 1), \quad 2.35 \text{ (primes),} \quad 2.99 \text{ (} \alpha = \frac{3}{2}),$$

against the predicted 2α , i.e. 2.00, ≈ 2.4 for $a_i \asymp i \log i$, and 3.00; for the squares the prediction is k^{-4} and the correction is already at $1 - c_1 = 8.9 \cdot 10^{-10}$, so its share is numerical noise and changes sign — the prediction taken to its limit rather than a counterexample. Concretely the share for the odd numbers runs 10.74%, 4.93%, 3.33%, 2.40%, 1.81% at $k = 120, 180, 220, 260, 300$. [\[experimentally confirmed, four profiles, \$k \leq 300\$; e8_r103\]](#) Keeping Φ as printed is therefore not a convenience at the k we happened to measure: c_d 's contribution is a vanishing fraction of a vanishing residual, so Φ is the right transfer function for every larger k and c_d belongs to E_k permanently. The fork is closed.

[\[experimentally confirmed; p3_r096\]](#) The verification behind Proposition 5.20 and (2), at $k = 180$, two profiles, $x = 0.10$ and $x = 0.30$, layers $d = 1, 3, D(k)$: the μ_d expression is exact to a relative $6 \cdot 10^{-5}$ or better; the two offsets of (2) to $7.6 \cdot 10^{-4}$ of δ_d or better; the R1 formula against the exact r_{B_d} over the window to $2.7 \cdot 10^{-3}$ – $3.7 \cdot 10^{-2}$, in every case between 4% and 14% of the predicted ε ; the parity oscillation of r_{B_d} over the window at most $3.7 \cdot 10^{-3}$; and ε itself grows by at most 44% from $d = 1$ to $d = D(k)$, so a single ε does serve all layers. With Lemmas 5.8 and 5.10 the tilted generating function decays like $\exp(-t_{\min}(1 - \eta)b/4)$ off the principal arc — exponentially in b , which is what P5 quotes. The tilted analogue of the $M(q)$ spectrum is lost, and is not needed: the rate enters Theorem 5.1 only through $E_k = o(1)$.

Remark 5.23 (why the constants cannot be imported). What Proposition 5.20 leaves open is not a mechanism but a list of explicit Edgeworth constants in region R1, so it is natural to ask whether they can be quoted rather than derived. The natural source is the Bernoulli-part-extraction literature, which produces local limit theorems with explicit universal constants

for non-identically distributed lattice summands, and $\sigma(S) = \sum a_i X_i$ is a sum of independent lattice summands. We checked, and the answer is no, for two reasons; we record them because a reader will otherwise ask the same question.

Write, following Giuliano and Weber [2], $\vartheta_X = \sum_k \mathbb{P}\{X = k\} \wedge \mathbb{P}\{X = k + 1\}$ and $\Theta_n = \sum_j \vartheta_{X_j}$; their local limit theorem carries the error $C\Theta_n^{-1/2}(H_n + \Theta_n^{-1})$ plus a term $(\log \Theta_n / (\text{Var} \cdot \Theta_n))^{1/2}$, with H_n a Berry–Esseen distance, and is useful exactly when $(\text{Var} / \Theta_n)^{1/2}(H_n + \Theta_n^{-1}) \rightarrow 0$.

First, the hypothesis is void here. Our summands are $X_j = a_j X'_j$ with a_j odd and at least 3, so the law of X_j charges the two integers 0 and a_j and no two *adjacent* ones. Hence $\vartheta_{X_j} = 0$ exactly, for every j and every p , so $\Theta_n = 0$ and the bound reads $\leq \infty$. Their own remark names this degenerate case.

Second, the obvious repair does not help. One can block the sequence and apply the method to block sums, whose laws do charge adjacent integers — but only once a block is large enough to have more subset sums than its span, and for blocks of *large* elements that takes a long time: $\{7, 9, 11\}$ has subset sums $\{0, 7, 9, 11, 16, 18, 20, 27\}$, no two adjacent. Measured, the smallest final block with $\vartheta > 0$ has size 7, 11, 15 and more than 20 at $k = 16, 32, 64, 128$, growing like $\log k$, so blocking buys $\Theta_n \lesssim k / \log k$. And then the applicability condition fails outright: our weights grow, so $\text{Var} \asymp S_2 \asymp k^3$ while $\Theta_n \lesssim k$ and $H_n \asymp k^{-1/2}$, giving $(\text{Var} / \Theta_n)^{1/2}(H_n + \Theta_n^{-1}) \asymp k^{1/2}$. Measured, that quantity runs 48.9, 101.4, 222.5, 495.1 for the odd numbers and 94.1, 195.1, 441.0, 1034.2 for the primes at $k = 16, 32, 64, 128$ — increasing, not tending to zero. [\[experimentally confirmed; bpart_r110\]](#)

The obstruction is the same feature that makes the present local limit theorem worth proving: the summands carry weights $a_j \asymp j$, so the variance grows like the cube of the number of summands rather than linearly, and a method calibrated to the latter has nothing to give. *The Edgeworth constants of R1 have to be derived, not imported.*

6 Universality of the rate function

[\[experimentally confirmed, \$k = 16..56\$, three reference sequences\]](#) With $\hat{I}_A(\rho) = -k^{-1} \log(r_A(\lfloor \rho T \rfloor) / 2^k)$ and $D_k = \hat{I}_P - \hat{I}_{\text{ref}}$, the difference tends to 0 for each of three reference sequences: the odd numbers, the arithmetic progression $6n + 5$, and a random odd sequence of the same range. The sign is a property of the reference, not a universal fact — it is positive against the odd numbers, negative against the random sequence, and changes sign near $k = 48$ against the progression.

Remark 6.1 (the primes sit near the random sequence). At $k = 56$ the distance to a random odd sequence is 0.0016 against 0.0146 to the arithmetic progression — an order of magnitude. The reading we propose is that the rate function does not see the arithmetic of A beyond its counting function: the primes are close to a random sequence of the same density because, for this observable, that is what they are. It is the same statement the transfer function makes locally, since Φ depends on A only through (N_d, σ_d, s_d) — the small elements, counted. This is the most direct evidence we have for the reading that the main term is universal and the arithmetic of the sequence enters only through the correction.

7 A profile for which the hypothesis fails

The hypothesis Theorem 4.1 needs is the convergence, with room, of the series in (1) — qualitatively, that $N_d \rightarrow \infty$ fast enough. Precisely:

$$(H) \quad \sup_k \sum_{d \geq 1} 2^{-N_A(d)} \delta_d^2 < \infty, \quad \text{and} \quad W(k) := \max\{\delta_d : 2^{-N_d} \delta_d^2 \geq k^{-1}\} = N^{o(1)}.$$

The first clause makes Φ converge and makes $\Phi(\lambda) - \Gamma = O(\lambda^2)$; the second is what the error budget needs, since the extraction window must be narrow enough for Proposition 5.20 to be uniform across it. For the odd primes $N_A(d) = \pi(2d) - 1 \gg d/\log d$, so both clauses hold, with $W(k) = O((\log k)^{2+\epsilon})$.

[\[experimentally confirmed\]](#) It is not vacuous. For the odd-ified cubes $a_i = 2\lfloor i^3/2 \rfloor + 1$:

profile	odds	primes	squares	cubes
N_d at $d = 10$ (any k)	10	7	4	2
$Q(0) = \Gamma^{-1} \sum 2^{-N_d}(\delta_d^2 - s_d/4)$	20.3	50.4	916	381 904

Only $\{1, 9\}$ ever lie below 20, at any k , so 2^{-N_d} never damps the growing δ_d^2 and the series turns over only near $d \approx 10^4$. At $k = 70$ the finite-size offset is still $1.75 \cdot 10^{-2}$ while the signal is $1.7 \cdot 10^{-11}$: nine orders of magnitude, so the profile is not merely slow but undecidable at any accessible size. The cubes are therefore not a defect of the theory but the reason (H) has to be there: without it, nothing in §5.2 or §5.3 fails, and yet the conclusion is out of reach.

7.1 Two different kinds of failure

The cubes fail *effectively* and not asymptotically, and the distinction is worth keeping. Computed directly, $\sum_d 2^{-N_d} \delta_d^2$ for the odd-ified cubes settles at $1.112 \cdot 10^7$ and does not grow with k (values at $k = 60, 100, 140, 180$: $1.112 \cdot 10^7$ from $k \approx 140$ on, against 69.5 for the odd numbers, 266.1 for $\alpha = 3/2$ and 6473 for the squares — each of them k -independent as well). [\[experimentally confirmed; alpha_r092\]](#) So (H) is not what the cubes violate; what they violate is any hope of seeing the asymptotics, the constant being four orders of magnitude larger than the odd numbers’.

A profile that fails (H) outright is available, and on the other side:

Proposition 7.1 (the $\alpha = 1$ threshold). *Let $A = A_k$ consist of k distinct odd integers with $a_i \asymp c i^\alpha$.*

(i) *If $\alpha < 1$ then (H) fails: $\sum_{d \geq 1} 2^{-N_d} \delta_d^2 \rightarrow \infty$.*

(ii) *If $\alpha \geq 1$ then (H) holds, with a bound depending on α and c alone.*

Proof. (i) The a_i are distinct odd integers, so $a_k \geq 2k - 1$; with $a_k \asymp ck^\alpha$ this forces $c \gg k^{1-\alpha}$, so $a_1 \asymp c \rightarrow \infty$. For $d < a_1/2$ no a_i is $\leq 2d$, hence $N_d = 0$, $\sigma_d = 0$, $\delta_d = d$, and

$$\sum_{d \geq 1} 2^{-N_d} \delta_d^2 \geq \sum_{d < a_1/2} 2d^2 \asymp a_1^3 \asymp k^{3(1-\alpha)} \rightarrow \infty.$$

(ii) For $\alpha \geq 1$ one has $a_1 = O(1)$ and $N_d = \#\{i : a_i \leq 2d\} \gg d^{1/\alpha}$, while $\delta_d = d + \sigma_d/2 \ll d^{1+1/\alpha}$, so the general term is $\ll 2^{-c'd^{1/\alpha}} d^{2+2/\alpha}$, which is summable: an exponential in a positive power of d beats a polynomial. The bound involves α and c and not k . $\square$

Both halves are visible in the computation. For $\alpha = 1/2$ the series reads 4923, 10245, 16608, 24464 at $k = 60, 100, 140, 180$, growing like $k^{3/2}$ — which is $a_1^3 = k^{3(1-\alpha)}$ with $\alpha = 1/2$, exactly as (i) predicts. For $\alpha \geq 1$ it is flat in k to every digit printed:

profile	odds	$a_i \asymp i^{3/2}$	squares	cubes	$\alpha = 1/2$
α	1	3/2	2	3	1/2
$\sum_d 2^{-N_d} \delta_d^2$	69.5	266.07	6473.4	$1.112 \cdot 10^7$	<i>grows</i>

[experimentally confirmed, $k = 60, \dots, 180$; [alpha_r092](#)] Proposition 7.1 is about clean powers, and the primes do not route through it: Theorem 4.1 for the primes quotes §5.2, and the primes' (H) is checked directly where (H) is stated, from $N_d = \pi(2d) - 1 \gg d/\log d$. No reader need look for $i \log i$ in part (ii). So **(H) cuts the power profiles at exactly $\alpha = 1$** , and the cubes sit inside the admissible range but beyond the reach of computation.

8 The biased invariant

Let P_q be the product measure under which each $a_i \in A$ is present independently with probability $q \in (0, 1)$, and write $\text{lm}_q(n)$, $r_q(n)$ for the P_q -measures of the sets of strict local minima and of ground states at the target n . The target is taken at the q -tilted centre $n_q = \lfloor qT \rfloor$.

Definition 8.1 (biased gap series).

$$\Gamma^{(q)}(A) = 1 + \sum_{d \geq 1} [q^{N_d} + (1 - q)^{N_d}].$$

At $q = \frac{1}{2}$ every bracket is $2 \cdot 2^{-N_d}$ and $\Gamma^{(1/2)}$ is the window series W_D of Part I [4], hence $\Gamma(A) + (2D + 1)2^{-k}$.

Remark 8.2 (why the classification survives the deformation). Part I's classification is measure-free: at error d , an overshooting configuration is a strict local minimum exactly when every $a \leq 2d$ is *excluded*, and an undershooting one exactly when every $a \leq 2d$ is *included*. Under P_q those two events have probability $(1 - q)^{N_d}$ and q^{N_d} , which is Definition 8.1. Writing $r_B^q(m) = \sum_{U \subseteq B, \sigma(U)=m} q^{|U|} (1 - q)^{|B| - |U|}$ and factorising P_q across I_d and B_d ,

$$\text{over}_q = (1 - q)^{N_d} r_{B_d}^q(n + d), \quad \text{under}_q = q^{N_d} r_{B_d}^q(n - d - \sigma_d), \quad r_q = \sum_{J \subseteq I_d} q^{|J|} (1 - q)^{N_d - |J|} r_{B_d}^q(n - \sigma(J)),$$

and the binomial sum over J is exactly 1, so a flat r^q gives the d -th term of Definition 8.1. [derived; the assignment verified per stratum, see Remark 8.17]

8.1 The fair coin is the flattest

Lemma 8.3 (termwise minimality). *For every integer $N \geq 0$ the function $f_N(q) = q^N + (1 - q)^N$ on $(0, 1)$ is symmetric about $q = \frac{1}{2}$ and convex, and $f_N(q) \geq f_N(\frac{1}{2}) = 2^{1-N}$, with strict inequality for $q \neq \frac{1}{2}$ as soon as $N \geq 2$.*

Proof. $f_N(1 - q) = f_N(q)$ is immediate. For $N \geq 2$, $f'_N(q) = N[q^{N-1} - (1 - q)^{N-1}]$ and $f''_N(q) = N(N - 1)[q^{N-2} + (1 - q)^{N-2}] > 0$, so f_N is strictly convex; and $t \mapsto t^{N-1}$ is strictly increasing, so $f'_N < 0$ on $(0, \frac{1}{2})$ and $f'_N > 0$ on $(\frac{1}{2}, 1)$. For $N \in \{0, 1\}$, f_N is the constant 2 or 1. $\square$

Remark 8.4 (formally verified). Lemma 8.3 is machine-checked. `termwise_min` gives the inequality for every N and every $q \in [0, 1]$, `termwise_min_strict` its strictness for $N \geq 2$ and $q \neq \frac{1}{2}$, and `termwise_min_iff` the equality case as a biconditional; the normalisation $2 \cdot (\frac{1}{2})^N = 2^{1-N}$ is `two_mul_half_pow`. The non-strict bound is proved twice, once from convexity of x^N on $[0, \infty)$ and once from Bernoulli's inequality (`termwise_min'`), so the two proofs check the *statement* against each other and not merely the tactic script. No sorry, and no axioms beyond `propext`, `Classical.choice`, `Quot.sound`. [Lean-verified; [Pnp/Theory/TermwiseMin.lean](#)]

Theorem 8.5 (the fair coin minimises the invariant). *For every finite increasing sequence A of odd positive integers with $|A| \geq 2$ and every $q \in (0, 1)$,*

$$\Gamma^{(q)}(A) \geq \Gamma^{(1/2)}(A),$$

with equality if and only if $q = \frac{1}{2}$.

Proof. Sum Lemma 8.3 over d . For strictness it suffices that some $d \leq D$ has $N_d \geq 2$: take any $d \geq (a_2 - 1)/2$, which is at most $D = (a_k - 1)/2$ because $a_2 \leq a_k$. $\square$

Remark 8.6 (what the theorem says). $q(1 - q)$ is the per-element variance of the coin. Biasing the coin in either direction lowers that variance and raises the invariant: **the more predictable the coin, the more rugged the landscape**. The effect is exactly symmetric in $q \leftrightarrow 1 - q$, as it must be, since complementing every configuration exchanges q with $1 - q$ and n with $T - n$.

Corollary 8.7 (closed form for the odd numbers). A corollary of Proposition 8.12: *for $A = (1, 3, \dots, 2k - 1)$ every m_j is 1, so $G(z) = z/(2 - z)$ and $1 + G(2q) + G(2 - 2q)$ is the display below. Directly: $N_d = d$ for $d \leq k$, so in the limit $k \rightarrow \infty$*

$$\Gamma^{(q)} = 1 + \frac{q}{1 - q} + \frac{1 - q}{q} = \frac{1}{q(1 - q)} - 1 \geq 3,$$

with equality only at $q = \frac{1}{2}$: the invariant of the odd numbers is the reciprocal of the coin's variance, minus one. [verified against the summed series to $5.3 \cdot 10^{-15}$ over $q = 0.10, \dots, 0.80$; e5_r098]

8.2 Where the zeros go

$\Gamma^{(q)}$ is a polynomial in q for finite A , and in statistical mechanics a phase transition *is* the accumulation of partition-function zeros on the real axis. Since $\Gamma^{(q)}$ plays the role of a free energy in the picture this Part uses, the question “where do the zeros go as $k \rightarrow \infty$, and do they pinch $[0, 1]$?” is the phase-transition question stated inside the arithmetic, with no physics needed to ask it and none to answer it. We asked it. The answer separates the profiles.

Group the layers by their count: with $m_0 = (a_1 - 1)/2$ and $m_j = (a_{j+1} - a_j)/2$ for $1 \leq j \leq k - 1$, the number of d with $N_d = j$ is m_j , so

$$\Gamma^{(q)}(A) = 1 + \sum_{j=0}^{k-1} m_j [q^j + (1 - q)^j] \quad (4)$$

after dropping the degenerate tail $N_d = k$, which is the truncation term $(2D + 1)2^{-k}$ already accounted for above. This is a polynomial of degree $k - 1$, it satisfies $\Gamma^{(q)} = \Gamma^{(1-q)}$ identically, and its value at $q = \frac{1}{2}$ is $\Gamma(A)$ to $O(2^{-k})$.

Remark 8.8 (the zeros of the odd numbers pinch the endpoints, at rate $2\pi/k$). For $A = (1, 3, \dots, 2k - 1)$ every $m_j = 1$, the two geometric sums collapse, and

$$q(1 - q)\Gamma^{(q)} = 1 - q + q^2 - q^{k+1} - (1 - q)^{k+1} =: F_k(q).$$

Substituting $q = w/(k + 1)$ sends $q^{k+1} \rightarrow 0$, $(1 - q)^{k+1} \rightarrow e^{-w}$ and $1 - q + q^2 \rightarrow 1$, so $F_k \rightarrow 1 - e^{-w}$, whose zeros are $w = 2\pi in$. Carrying the first correction, $e^{-w} = 1 - q + q^2$ gives $-w = -w/(k + 1) + 2\pi in + \dots$ and hence

$$\operatorname{Im} q_{\pm 1} = \pm \frac{2\pi}{k} + O(k^{-2}).$$

So the zeros do pinch the real segment — but only at its endpoints $q = 0$ and $q = 1$, the two degenerate measures, and at a rate with no constant to fit. The limit $1 - q + q^2$ has its own roots at $e^{\pm i\pi/3}$, where $|q| = |1 - q| = 1$ exactly; those stay at height $\sqrt{3}/2$ and never approach the real axis. [derived, and measured. $k \cdot (\text{distance to } [0, 1])$ runs 6.0518, 6.1919, 6.2348, 6.2629, 6.2721, 6.2762, 6.2784 at $k = 16, 24, 32, 48, 64, 80, 96$ against $2\pi = 6.283185$, approaching from below like $1/k$; the nearest zero has $\text{Re } q \rightarrow 1$. The nearest zero to the fair coin stays away: $|q - \frac{1}{2}|$ runs 0.5177, 0.5046, 0.5021 at $k = 32, 64, 96$. Primes and random odd sets give the same picture with $k \cdot d$ between 6.1 and 6.7. leeyang_r179, leeyang2_r179.]

Remark 8.9 (and the family that breaks the theorem is pinched at the fair coin). Run the same computation on $a_i = 2^i + 1$ — the lacunary family of §7, the witness for which the annealed prediction fails. There $m_j = 2^{j-1}$ exactly, so

$$\Gamma^{(q)} = 3 + \frac{1}{2} \sum_{j=1}^{k-1} [(2q)^j + (2(1-q))^j],$$

two geometric series with ratios $2q$ and $2(1-q)$. On the line $\text{Re } q = \frac{1}{2}$ those ratios are complex conjugates of equal modulus, so the two series can cancel — and they do. Write $q = \frac{1}{2} + it$, $z = 2q = 1 + 2it$, $\rho = |z| = \sqrt{1 + 4t^2}$ and $\theta = \arg z = \arctan 2t$. Then $z - 1 = 2it$, so $\sum_{j=1}^{k-1} z^j = (z^k - z)/(2it)$, and since $\rho \sin \theta = \text{Im } z = 2t$ this gives, *with no approximation*,

$$\Gamma^{(q)} = 2 + \frac{\rho^k \sin(k\theta)}{2t} \quad \text{on } \text{Re } q = \frac{1}{2}.$$

The zeros are the solutions of $\sin(k\theta) = -4t\rho^{-k}$, whose right-hand side is $O(k^{-1})$ at the first rungs, so they sit at $\theta_n = (n\pi/k)(1 + O(k^{-1}))$ for *every* integer $n \geq 1$, and

$$\text{Im } q_{\pm 1} = \pm \frac{\pi}{2k} + O(k^{-2}), \quad \text{Re } q_{\pm 1} = \frac{1}{2} \text{ exactly.}$$

The companion family $a_i = 2^i - 1$ carries the same weights with $w_0 = 0$, so there the constant 2 is absent, $\Gamma^{(q)} = \rho^k \sin(k\theta)/2t$ is an identity, and the zero set on the line is exact:

$$q = \frac{1}{2} + \frac{i}{2} \tan(n\pi/k), \quad n = 1, 2, \dots$$

For this family the fair coin is not a minimiser sitting in an analytic region: it is the point the zeros are closing in on.

Which is the finding. The Lee–Yang zeros of $\Gamma^{(q)}$ separate the profiles on which the annealed count is asymptotically exact from the one on which it is not — and they separate them at $q = \frac{1}{2}$, the value at which the whole of Parts I and II is stated. Two independent routes, the local-limit route of §5.2 and the zero-counting route here, partition the same profiles the same way.

Stated so that it can be attacked, because two families are two families and not a theorem. FALSIFIER 1: exhibit a profile satisfying (H) whose zeros approach $q = \frac{1}{2}$. FALSIFIER 2: exhibit a profile violating (H) whose zeros stay away from it. Neither is ruled out by anything here. [proved for the displayed identity and for the exact zero set of $a_i = 2^i - 1$; measured for the rate constant of $a_i = 2^i + 1$; conjectured for the dichotomy. $k \cdot \text{Im } q_1$ runs 1.5951, 1.583, 1.5769, 1.5739 at $k = 128, 256, 512, 1024$ against $\pi/2 = 1.5707963$, with $\text{Re } q_1 = \frac{1}{2}$ to every digit computed; $|q_1 - \frac{1}{2}|$ runs 0.052, 0.0253, 0.0167 at $k = 32, 64, 96$ and is going to zero, against the odd numbers, for which Proposition 8.10 proves the disc $|q - \frac{1}{2}| < 1/6$ zero-free and the measured winding number there is 2.3×10^{-32} . leeyang7_r194, leeyang8b_r194, figures_r194. — **Erratum against v1.1.0.** That version stated this remark with $\cos(k \arctan 2t)$ in place of

sin, the first-rung rate $3\pi/2k$ in place of $\pi/2k$, and a ladder of odd multiples $(2n+1)\pi/2k$ in place of $n\pi/k$. *The error was to replace the geometric sum by its largest term:* that keeps the modulus and discards the factor $1/(z-1)$, whose argument is exactly $-\pi/2$, and a quarter turn of phase moves a ladder by half a rung. The point $t = 3\pi/2k$ is not a zero at all; $\Gamma^{(q)}$ there measures 2.325885, 2.1374425, 2.0630564, 2.030195 at $k = 128, 256, 512, 1024$, converging to the constant 2. The two measured series quoted in that version, $|q_1 - \frac{1}{2}| = 0.0520, 0.0253, 0.0167$ and $k \cdot \text{Im } q_1 = 4.7736, 4.7457, 4.7299$, are both correct measurements of genuine zeros — *of the first and the third rung respectively*, reported under one symbol. The agreement of the second with $3\pi/2$ was read as confirmation because the search had been aimed where cos predicted a zero and found a real one there.]

Proposition 8.10 (the fair coin is not pinched when the layer weights decay geometrically). *[proved. The numerical companion checks the disc at 576 points on four profiles with no failure, and reports the proved radius beside the measured one; leeyang4_r182.] Write*

$$g_1(x) = \sum_{j \geq 1} j m_j 2^{-j} x^{j-1},$$

and suppose g_1 converges at some $x_0 > 1$. Then for every $r > 0$ with $1 + 2r \leq x_0$ and

$$4r g_1(1 + 2r) < \Gamma(A),$$

the disc $|q - \frac{1}{2}| < r$ contains no zero of $\Gamma^{(q)}$ — **and the bound is uniform in k** , because g_1 and the lower bound for Γ are.

Proof. For $|q - \frac{1}{2}| = r$ both $|q|$ and $|1 - q|$ are at most $\frac{1}{2} + r$, and

$$|q^j - 2^{-j}| = \left| q - \frac{1}{2} \right| \left| \sum_{i < j} q^i 2^{-(j-1-i)} \right| \leq r j \left(\frac{1}{2} + r \right)^{j-1} = 2r j 2^{-j} (1 + 2r)^{j-1},$$

and the same for $(1 - q)^j$. Summing against m_j and using $\Gamma^{(1/2)} = \Gamma(A)$,

$$|\Gamma^{(q)} - \Gamma(A)| \leq 4r g_1(1 + 2r) < \Gamma(A),$$

so $\Gamma^{(q)} \neq 0$ on that circle and, the left side being increasing in r , inside it. Finally $m_j \geq 1$ for $1 \leq j \leq k-1$ because consecutive odd integers differ by at least 2, so $\Gamma(A) \geq 1 + 2 \sum_{j=1}^{k-1} 2^{-j} = 3 - 2^{2-k}$: *the odd numbers are the densest admissible profile and they minimise the invariant.* $\square$ $\square$

Remark 8.11 (what the criterion is, and what it is not). Everything turns on the radius of convergence R of $g(x) = \sum_j m_j 2^{-j} x^j$. Since $g(1) = (\Gamma(A) - 1)/2$, $R > 1$ **says that Γ is finite and that the layer weights decay geometrically, while $R \geq 1$ is only necessary**: the radius says nothing about the boundary point $z = 1$, and at $R = 1$ the invariant may do either. Both cases occur, and the pair is the sharpest thing in this section — $a_i = 2^i + 1$ has $R = 1$ with $\Gamma_k = k + 2 \rightarrow \infty$, while the profile $m_j = \max(1, \lfloor 2^j(j+1)^{-2} \rfloor)$ — **family C** — has $R = 1$ with $\Gamma_k \rightarrow 5.230199559$. *The clamp at 1 is forced and is not a fitted parameter*: m_j is half the gap between consecutive elements of a set of distinct odd numbers, so it is a positive integer, and the unclamped expression is not a profile at all. It is also not a cosmetic difference — without the clamp the same expression gives $\Gamma_k \rightarrow 3.480199559$, and this correction is r200's, against a formula that shipped in v1.1.1 without it. *Which of the two cases holds is exactly what sets the rate at which the zeros close in* (Remark 8.14). Proposition 8.10 needs the stronger one, and it is satisfied by every profile in this paper except the lacunary family: measured over the top half of the layers at $k = 60$, the quantity

$(m_j 2^{-j})^{1/j}$ is 0.5000 for the odd numbers, 0.5250 for the primes, 0.5543 for a random odd set and 0.5600 for $a_i \asymp i^2$, against 0.9883 for $a_i = 2^i + 1$.

For the odd numbers $g_1(x) = 2/(2-x)^2$ exactly, the condition becomes $12r^2 - 20r + 3 > 0$, and **the proved radius is** $r = \frac{1}{6}$. The measured distance from the fair coin to the nearest zero is 0.5046 at $k = 64$, so the proof is conservative by a factor 3.03.

The unexpected part, and the reason this is more than an exclusion certificate. The same ratio appears everywhere, including where the hypothesis fails: proved against measured is 3.03 for the odd numbers, 2.81 for the primes, 1.65 for a random odd set, 3.97 for $a_i \asymp i^2$ — and 3.10 for the lacunary family, where the proved radius is not uniform but shrinks like $1/k$ (0.00816 at $k = 64$) exactly as the measured distance does (0.02530). **So g_1 does not merely certify an empty disc; it tracks where the nearest zero actually is, on both sides of the dichotomy.**

What is still not proved, stated plainly. The converse. $R = 1$ removes the obstruction to a zero near $q = \frac{1}{2}$; it does not produce one. Remark 8.9 exhibits one family where a zero does appear, and one family is one family. [\[the proposition is proved; the tracking observation is measured on five profiles; the converse is open. leeyang4_r182.\]](#)

Proposition 8.12 (the invariant is a generating function evaluated twice). [\[proved; the identity is checked to \$10^{-30}\$ at twelve points on four profiles, leeyang5_r183.\]](#) Let $G = G_A$ be the generating function of the layer weights introduced in Part I [4], Proposition 1.3: with $m_0 = (a_1 - 1)/2$ and $m_j = (a_{j+1} - a_j)/2$,

$$G(z) = \sum_{j \geq 0} m_j 2^{-j} z^j, \quad \text{so that} \quad \Gamma(A) = 1 + 2G(1).$$

Then, identically,

$$\Gamma^{(q)}(A) = 1 + G(2q) + G(2 - 2q).$$

The deformation does not introduce a new object; it reads the same generating function at two points instead of one. Consequently, writing R for the radius of convergence of G :

1. $G(1) = (\Gamma(A) - 1)/2$, so $\Gamma(A)$ **is finite exactly when G converges at $z = 1$** : $R > 1$ forces that and $R < 1$ forbids it, but $R = 1$ **decides nothing** — see Remark 8.11 for a profile of each kind;
2. the limit of $\Gamma^{(q)}$ is analytic precisely on the lens $|q| < R/2$ and $|1 - q| < R/2$;
3. **the fair coin $q = \frac{1}{2}$ lies in the interior of that lens if and only if $R > 1$** , at distance $R/2 - \frac{1}{2}$ from its boundary.

Proof. $m_j q^j = m_j 2^{-j} (2q)^j$ and likewise for $1 - q$; sum over j and use Definition 8.1. Item 1 is Part I's proposition evaluated at $z = 1$; the lens and the distance are the domain of convergence of the two evaluations. $\square$ $\square$

Remark 8.13 (the dichotomy, as a formula). For the odd numbers $w_j = 2^{-j}$, so $G(z) = z/(2-z)$ and $\Gamma^{(q)} = 1 + \frac{q}{1-q} + \frac{1-q}{q} = \frac{1}{q(1-q)} - 1$: Corollary 8.7 is Proposition 8.12 for that profile. There $R = 2$ and the fair coin sits at distance $\frac{1}{2}$ from the boundary — which is where the zeros were measured to stop. For $a_i = 2^i + 1$ we have $w_j = \frac{1}{2}$, $R = 1$, and the distance is 0: **the fair coin is on the boundary of analyticity.** That is Remark 8.9's pinch, and it is no longer a second family — it is the endpoint of a single scale.

So the dichotomy should be a formula, and it is one. Build the interpolating profile with layer gaps $a_{j+1} - a_j = 2c^j$, so $m_j = c^j$, $w_j = (c/2)^j$, $R = 2/c$, and the prediction is

$$\text{dist}\left(\frac{1}{2}, \text{nearest zero of } \Gamma^{(q)}\right) \longrightarrow \frac{1}{c} - \frac{1}{2},$$

which is $\frac{1}{2}$ at $c = 1$ (the odd numbers) and 0 at $c = 2$ (the lacunary family), with no constant to fit.

c	1.00	1.10	1.25	1.40	1.60	1.80	2.00
predicted $1/c - 1/2$	0.500000	0.409091	0.300000	0.214286	0.125000	0.055556	0
measured, $k = 70$	0.503895	0.413403	0.305147	0.220690	0.133982	0.070741	0.022306
ratio	1.0078	1.0105	1.0172	1.0299	1.0719	1.2733	—

The ratio approaches 1 from above and degrades gracefully as $c \rightarrow 2$, which is what a finite k must do when the quantity being predicted is going to zero. *At $c = 2$ the predicted value is exactly 0 and the measurement follows it down:* $k \cdot \text{dist}$ runs 1.552860, 1.556681, 1.559637, 1.561639, 1.563051 at $k = 24, 40, 56, 72, 88$, rising towards $\pi/2$ — the first rung $n = 1$ of the ladder $k \arctan 2t = n\pi$. *The lacunary family gives the same constant:* for $a_i = 2^i + 1$, $k \cdot \text{Im } q_1$ runs 1.5951, 1.583, 1.5769, 1.5739 at $k = 128, 256, 512, 1024$, also rising to $\pi/2$. **At the boundary the rate is $\pi/2k$ for every profile in this remark — one ladder, one first rung, no profile-dependent integer.**

The scope of that sentence is load-bearing, and an earlier version of it did not carry one. Every profile above has Γ divergent at $c = 2$; the layer family is built so that it does. Remark 8.14 exhibits $R = 1$ profiles with Γ convergent, and their rate is not $\pi/2k$ but $\sqrt{s \log k/2k}$, larger by a factor $\asymp \sqrt{k \log k}$. **So the constant $\pi/2$ belongs to the divergent regime and not to $R = 1$,** and the formula displayed here predicts *where* the zeros stop, not *how fast* they arrive when they do not stop.

What this settles and what it does not. Proposition 8.12 is an identity and its three consequences are immediate, so *the fair coin is interior exactly when $R > 1$* is proved, and Proposition 8.10 is the quantitative form of one direction. The formula above is **measured**. What is open is that the zeros accumulate *on* the boundary rather than merely failing to enter: Hurwitz's theorem gives one direction on compact subsets of the analytic region, and the other needs the limit to be non-vanishing there, which is a statement about G and is not proved here. **[the identity and its consequences are proved; the distance formula is measured at seven values of c and two values of k ; the accumulation is open; the rate $\pi/2k$ is measured, and only on profiles with Γ divergent — see Remark 8.14 for the other regime, which was measured after this remark was written and which contradicts the unqualified reading of it. leeyang5_r183, rate_r200.]**

Remark 8.14 (the rate is not universal: two regimes at $R = 1$). $R = 1$ decides that the fair coin sits *on* the boundary of analyticity. It does not decide how fast the zeros reach it, and the two halves of the pair in Remark 8.11 — both with $R = 1$ — fall on opposite sides of that question.

The mechanism is one line. On the critical line $\Gamma_k^{(q)} = 1 + 2 \sum_j w_j \rho^j \cos(j\theta)$ with $w_j = m_j 2^{-j}$, $\rho = \sqrt{1 + 4t^2}$ and $\theta = \arctan 2t$, so $\rho^j = e^{2jt^2(1+O(t^2))}$ and the last term carries $w_k e^{2kt^2}$.

- If $w_j \rightarrow w > 0$ — equivalently Γ divergent — the growth factor is irrelevant at $t \asymp 1/k$ and the zero is produced by the *phase* alone, at $k\theta = \pi$. That is the ladder of Remark 8.9 and the rate is $\pi/2k$.
- If $w_j \asymp (j+1)^{-s}$ with $s > 1$ — Γ convergent — the phase cannot do it: at $\theta = x/k$ every $\rho^j \rightarrow 1$ for $j \leq k$, so $\sum_j w_j \cos(jx/k) \rightarrow \sum_j w_j > 0$ and there is no dip at that scale. What produces the zero is instead the growth ρ^j overcoming the decay w_j , and $k^{-s} e^{2kt^2}$ is $O(1)$ *exactly* when $2kt^2 = s \log k$.

Hence the dichotomy inside the boundary case:

$$\Gamma \text{ divergent: } t_1 \sim \frac{\pi}{2k}, \quad \Gamma \text{ convergent, } w_j \asymp (j+1)^{-s} : \quad t_1 \sim \sqrt{\frac{s \log k}{2k}}.$$

The second is larger by a factor $\asymp \sqrt{k \log k}$: the pinch is slower, and it is slower in a way that reads off the exponent s . There is no constant to fit in either.

Measured against family C, $m_j = \max(1, \lfloor 2^j(j+1)^{-s} \rfloor)$, whose Γ converges for every $s > 1$ (to 5.2282493, 5.0054793, 5.0000879 at $s = 2, 3, 4$ and $k = 512$), as measured t_1 divided by the predicted $\sqrt{s \log k/2k}$:

k	32	64	128	256	512
$s = 2$	1.276794	1.187405	1.034218	1.004723	0.9962483
$s = 3$	1.521603	1.199195	1.073958	1.04068	1.008387
$s = 4$	1.354838	1.263709	1.153411	1.044587	1.039257

The $\sqrt{s}$, tested with no k in it. At fixed $k = 512$ the ratio $t_1(s=4)/t_1(s=2)$ is 1.475266 against the predicted $\sqrt{2} = 1.414214$, high by 4.3%. **That excess is accounted for rather than excused:** $1.475266/\sqrt{2} = 1.04317$, and the amount by which the $s = 4$ column still lags the $s = 2$ column at $k = 512$ is $1.039257/0.9962483 = 1.04317$ — the same number. The residual in the s -test is the residual already visible in the k -test, not a second discrepancy.

Control, and it is the one that makes this two regimes and not one. For $a_i = 2^i - 1$, where $\Gamma_k = k$ diverges, $k \cdot t_1$ runs 1.5720592, 1.5711118, 1.5708752, 1.570816 at $k = 64, 128, 256, 512$ against $\pi/2 = 1.5707963$, while t_1 divided by $\sqrt{2 \log k/2k}$ runs 0.0963586, 0.0630435, 0.0416931, 0.0277943 — **going to zero, not to a constant**. A control that also obeyed the law would have shown one regime; this one fails it in the predicted direction and at the predicted size.

What this does to Problem 8.15, said plainly rather than quietly. That problem was posed asking whether $R = 1$ forces the pinch at rate $\pi/2k$. The rate clause is answered **no**, by family C, and the problem below has been restated without it. The question of whether $R = 1$ forces a pinch at all is untouched and remains open — but it is now supported by five profiles instead of two, and by two different rates, which is better evidence for it than one repeated rate would have been. [the mechanism is derived (asymptotically, not to the standard of the identities in this section); the two rates are measured with no fitted constant, at three exponents and five sizes, with a control that fails the law; the general two-regime law is conjectured, and the dividing line is conjectured to be $\sum_j w_j < \infty$ rather than anything finer. rate_r200, whose four pre-registered falsifiers are printed in the log; the two instrument controls among them fired on the first run and caught both an index error here and the missing clamp in Remark 8.11.]

Problem 8.15 (does $R = 1$ force the pinch?). For admissible profiles with $R = 1$, do the zeros of $\Gamma^{(q)}$ approach $q = \frac{1}{2}$ as $k \rightarrow \infty$ — that is, does a first rung always exist? As first posed, this problem also asked whether that rung always descends at rate $\pi/2k$. It does not, and the clause is withdrawn here rather than dropped silently (Remark 8.14): the rate is $\pi/2k$ when Γ diverges and $\sqrt{s \log k/2k}$ when $w_j \asymp (j+1)^{-s}$. **What is open is the existence of the pinch, not its speed**, and the two known speeds are evidence for the existence rather than against it. The missing ingredient is named rather than left to be located: what is absent is the non-vanishing of the limit of G on its circle of convergence. Hurwitz's theorem gives non-accumulation inside the lens on compact subsets; accumulation on the boundary needs that non-vanishing, and it is a statement about G alone. The classical half, transported. For a single power series this is **Jentzsch's theorem** (1914): if the radius of convergence is finite and positive, then every point of the circle of convergence is a limit point of the zeros of the sections, with Szegő's refinement giving the angular equidistribution of a subsequence. Applied to G , it says that the zeros of the sections G_k accumulate at every point of $|z| = R$

— and when $R = 1$ the point $z = 1$ is such a point, which is $q = \frac{1}{2}$ under $z = 2q$. *So the accumulation of the zeros of each summand at the fair coin is classical.*

What is not classical, stated sharply, is that the sum inherits it. Our object is $1 + G_k(2q) + G_k(2 - 2q)$: two sections under affine substitutions, plus a constant, and adding functions can cancel zeros. We found no result covering sums of sections, and the missing ingredient is exactly that — *does a sum of two sections, under substitutions that make them complex conjugates on a line, inherit boundary accumulation?*

And the measured ladder is the fingerprint of the sum, not of a section. Sections of a single geometric series have zeros at angular spacing $2\pi/k$; ours sit at $\theta_n = n\pi/k$ on the ray $z = 1 + 2it$ — *twice as densely. This comparison is about the divergent regime, where a ladder is what there is*; in the convergent regime the zeros are not a ladder at all, which is Remark 8.14's point restated. That is what $\sin(k\theta)$ does and what $e^{ik\theta} - 1$ does not, and the doubling is the signature of the *sum*: taking the real part of a section folds $e^{ik\theta}$ into $\sin(k\theta)$, whose zeros are twice as dense. **So the doubled ladder is precisely the part of the phenomenon that the single-section theory does not describe.** [proved for $a_i = 2^i - 1$, where the zero set on the line is exactly $q = \frac{1}{2} + \frac{i}{2} \tan(n\pi/k)$; measured for the other two boundary profiles, whose $k\theta_n/\pi$ approach 1, 2, 3, 4, 5, 6 with error $O(k^{-1})$, and for family C at three exponents (Remark 8.14), so five profiles in all; the withdrawn rate clause is **disproved** by family C; the existence question is **conjectured** in general and remains the open half. All near- $\frac{1}{2}$ zeros were confirmed to lie *on* the line by an argument-principle count on $|q - \frac{1}{2}| = r$ for r up to 0.25, which matched twice the number of on-line sign changes in every case. leeyang6_r194, leeyang7_r194, leeyang8b_r194. **Literature pass completed** (r191): Jentzsch's theorem and the Jentzsch–Szegő refinements cover the sections themselves and have been extended to Dirichlet, Kapteyn and Neumann series, to analytic curves and to ultrametric fields; *no source was found covering sums of sections*, so the transport is half and the remaining half is the statement displayed above.]

Remark 8.16 (the correction has a smaller domain than the invariant it corrects). Along the same family, Γ converges for every $c < 2$ — that is $R > 1$ — but its first correction does not. In $Q(0) = \Gamma^{-1} \sum_d 2^{-N_d} (\delta_d^2 - s_d/4)$ the layer count satisfies $N_d \sim \log_c d$, so $2^{-N_d} \sim d^{-\log 2 / \log c}$ and the terms grow like $d^{2 - \log 2 / \log c}$. That is summable exactly when $\log 2 / \log c > 3$, so

$$Q(0) \text{ exists} \iff c < 2^{1/3} = 1.259921 \dots$$

There is therefore a window $2^{1/3} < c < 2$ in which the annealed prediction is finite and its first correction is not. The measurement agrees: $Q(0)$ settles to 20.07, 20.30, 20.33 at $k = 16, 20, 24$ for the odd numbers — converging to $61/3$, which is Proposition 12.1's exact value and an independent check of the computation — while between $k = 20$ and $k = 24$ the same quantity multiplies by 1.359, 3.695, 17.502, 68.871, 212.436 at $c = 1.25, 1.40, 1.60, 1.80, 2.00$.

Why this is recorded here rather than as an aside. It was found while trying to decide whether the zeros of §8.2 see the *rate* of the annealed limit or only its existence. The natural test compares the measured error against $\Gamma Q(0)/\sigma^2$, and **that comparison is unavailable exactly where the question is interesting**: the null quantity ceases to exist at $c = 2^{1/3}$, below the whole region where the zero-distance $1/c - \frac{1}{2}$ has moved appreciably. *An experiment whose control evaporates before the signal appears is not a weak experiment; it is a different question.* **The constant dies; the running correction does not.** Truncating at the canonical cut $d < a_k/2 \asymp c^k$ gives $Q_k \asymp (c^k)^{3-\lambda} = c^{3k} 2^{-k}$ with $\lambda = \log 2 / \log c$, and $\sigma_k^2 \asymp c^{2k}$, so

$$\rho_k(c) = \frac{Q_k}{\sigma_k^2} \asymp \left(\frac{c}{2}\right)^k$$

throughout $c < 2$: in the window the first correction survives as a quantity with its own exponential scale, neither a constant over σ^2 nor absent. [derived by fable-5 and **checked here before use**: $\rho_k(2/c)^k$ measures 0.8645, 0.8679, 0.8610, 0.8564, 0.8542 at $k = 12, 16, 20, 24, 28$ for $c = 1.6$ — flat — and 3.3199, 4.1888, 4.6201, 4.8066, 4.8787 for $c = 1.4$, still converging. At $c = 1.8$ it is 0.3132, 0.2794, 0.2587, 0.2465, 0.2390 and has not settled by $k = 28$, so the law is confirmed where it is confirmed and not beyond. *Whether the measured error $|\text{lm}/r - \Gamma|$ tracks ρ_k is a separate question, and it is answered only to a factor.* At $c = 1.6$ the relative error is still $O(1)$ at every k for which lm/r is exactly computable, so nothing is measurable there. At $c = 1.4$, with an exact `int64` enumeration validated against a big-integer reference on ten values of k — *identical integers, not close ones* — the relative error falls to $2.4 \cdot 10^{-4}$ by $k = 40$ and the decay rate of $|\text{lm}/r - \Gamma|$ measures 0.2574 ± 0.1319 against $\log(2/c) = 0.3567$, a separation of 0.75 standard uncertainties. Equivalently the drift of $e(2/c)^k$ is $+0.0993 \pm 0.1319$ per unit k : *consistent with zero*, with a total drift of a factor 6.0 against a scatter of 22.1. **So the layer law is consistent with the measurement and is not confirmed to any precision worth quoting**, and a wrong law within a factor two of it would not have been distinguished. `window_r187`, `window2_r187`, `window3_r189`, `window3b_r189`.]

[the criterion $c < 2^{1/3}$ is **derived** from the growth of N_d and **measured** at seven values of c ; the recovery of 61/3 is an independent check of Proposition 12.1. The consequence for §8.2 is that the rate question is **open by a different route than the one attempted**. `hrate_c_r185`.]

8.3 Verification

[**experimentally confirmed**] Measured lm_q/r_q at n_q against the point prediction of Definition 8.1, relative error:

profile	q	$k = 60$	$k = 100$	$k = 140$
odds	0.20	$1.00 \cdot 10^{-2}$	$3.16 \cdot 10^{-3}$	$1.49 \cdot 10^{-3}$
odds	0.30	$2.63 \cdot 10^{-3}$	$8.02 \cdot 10^{-4}$	$3.75 \cdot 10^{-4}$
odds	0.50	$2.73 \cdot 10^{-4}$	$5.99 \cdot 10^{-5}$	$2.19 \cdot 10^{-5}$
primes	0.30	$2.18 \cdot 10^{-3}$	$5.02 \cdot 10^{-4}$	$2.20 \cdot 10^{-4}$
primes	0.50	$1.73 \cdot 10^{-4}$	$2.26 \cdot 10^{-5}$	$7.32 \cdot 10^{-6}$

There are no free parameters. Two features: the error decays faster than $1/k$ (fitted exponent 2.25 at $q = 0.2$ to 3.0 at $q = 0.5$), and it grows monotonically as q leaves $\frac{1}{2}$, by a factor ≈ 40 from $q = 0.5$ to $q = 0.2$ at fixed k . The second is what the biased flatness hypothesis will have to quantify. [**TODO** state the biased (H); design after §8.4.]

Remark 8.17 (the assignment, and a test that could not decide it). Definition 8.1 is symmetric in $q \leftrightarrow 1 - q$, and so is the exact complementation identity $\text{lm}_q(n) = \text{lm}_{1-q}(T - n)$ — at the q -tilted centre, $T - qT = (1 - q)T$. Neither can therefore distinguish the assignment of q^{N_d} and $(1 - q)^{N_d}$ to the two strata: measured at $k = 100$, the ratio at $q = 0.4$ and at $q = 0.6$ agrees to every printed digit. The discriminating test is per stratum, over_q/r_q against $(1 - q)^{N_d}$ and under_q/r_q against q^{N_d} ; measured, both ratios lie in $[0.9978, 1.0018]$ over $d \leq 6$, whereas a swap would disagree by a factor of 161 at $d = 6$. [**experimentally confirmed**; `e5_r098`]

8.4 The bias breaks the reflection symmetry

[**experimentally confirmed, shape**] Off the q -tilted centre the deviation $\text{dev}_q = (\text{lm}_q/r_q)/\Gamma^{(q)} - 1$ behaves differently at $q = \frac{1}{2}$ and away from it, and the reason is structural. The transfer-function strand's transfer function is even in λ because the two strata sit at $\pm\delta_d$ about a

common centre — the analytic form of $\text{lm}(n) = \text{lm}(T - n)$. Under bias that symmetry maps q to $1 - q$ rather than to itself, so the odd orders need not cancel, and the first-order coefficient is

$$L^{(q)}(A) = \sum_{d \geq 1} \delta_d [q^{N_d} - (1 - q)^{N_d}],$$

which is antisymmetric in $q \leftrightarrow 1 - q$ and vanishes identically at $q = \frac{1}{2}$. Measured for the odd numbers at $k = 80$: $L^{(q)} = -28.634, -8.843, 0, +8.843, +28.634$ at $q = 0.3, 0.4, 0.5, 0.6, 0.7$, and for the primes $-83.717, -23.677, 0, +23.677, +83.717$.

The data agree. Over a range of targets, the spread (max over min) of $|\text{dev}_q - \text{dev}_q(0)|$ divided by λ_q and by λ_q^2 :

profile	q	divided by λ_q	divided by λ_q^2
odds	0.30	1.055	3.778
odds	0.50	3.118	1.002
primes	0.30	1.162	4.149
primes	0.50	3.104	1.012

So the response is linear in λ under bias and quadratic at $q = \frac{1}{2}$, where $(\text{dev} - \text{dev}_0)/\lambda^2$ is constant to three digits — 19.8221 to 19.8586 for the odd numbers and 48.8625 to 49.4422 for the primes, across a threefold range of the target offset.

Remark 8.18 (the range is part of the test). The first measurement used offsets $x \in [0.04, 0.20]$ and came out incoherent for the primes at $q = \frac{1}{2}$ (a quadratic spread of 15.4). The reason is worth recording because it is a property of ratio tests and not of this problem. At $q = \frac{1}{2}$ the response is *even* in x , so $\text{dev} - \text{dev}_0$ cannot change sign for the reason under test; at finite k the vertex of the parabola sits at some $x_* \neq 0$, and $\text{dev} - \text{dev}_0$ crosses zero at $\pm x_*$. Measured, $x_* \approx 0.03$, with the signs symmetric in $\pm x$ exactly as that account predicts. A range containing the crossing makes *both* ratios pass through zero, so the spread of either is set by how close a grid point landed to x_* — an accident of the grid. **A ratio test must not straddle a zero of its own numerator.** The table above uses $x \in [0.10, 0.30]$, outside x_* ; the repair is the range, not the omission of the offending point.

[experimentally confirmed; e6_r100, e6b_r115, whose pipeline is checked against e6_r100 by reproducing [19.8186, 19.8352] on the original range, and whose discriminating control at $q = 0.3$ still returns *linear* on the new range]

9 The biased transfer function

9.1 The sign of λ , which now matters

The transfer-function strand's transfer function is even in λ , so the sign of the slope is not observable there. Splitting the cosh makes it observable, so we fix it first:

$$\lambda_q(n) = + \frac{d}{dm} \log r_A^q(m) \Big|_{m=n} = \frac{1}{4} \log \frac{r_A^q(n+2)}{r_A^q(n-2)} + O(\cdot),$$

positive below the biased centre. This is the convention under which $\lambda \approx s$; the transfer-function strand's 2.1 carried the opposite sign and has been corrected, with no numerical consequence there for exactly the reason just given.

9.2 The derivation

Take $P_{q,s}$: element a present with weight $q e^{-sa}$, absent with weight $1 - q$, so $p_a = q e^{-sa} / ((1 - q) + q e^{-sa})$ and $p_a = q - q(1 - q)sa + O(s^2 a^2)$. Two consequences, both absent at $q = \frac{1}{2}$:

- the per-element variance is $a^2 p_a(1 - p_a) = q(1 - q)a^2(1 + O(sa))$, so the layer weight of the transfer-function strand's 5.21 becomes $e^{-\lambda^2 q(1-q)s_d/2}$;
- removing I_d removes mean $q\sigma_d$, so the two strata no longer sit at $\pm\delta_d$ about a common centre. They sit at

$$\delta_d^+ = d + q\sigma_d, \quad -\delta_d^-, \quad \delta_d^- = d + (1 - q)\sigma_d.$$

The bias splits the pair asymmetrically, in the ratio $q : (1 - q)$.

Hence

$$\Phi^{(q)}(\lambda) = 1 + \sum_{d \geq 1} \left[(1 - q)^{N_d} e^{\lambda \delta_d^+} + q^{N_d} e^{-\lambda \delta_d^-} \right] e^{-\lambda^2 q(1-q)s_d/2}. \quad (5)$$

[derived; three internal checks below are exact, the point prediction is verified]

Three checks are built in, and all three hold to machine zero ($0.00 \cdot 10^0$, two profiles, $k = 60$): $\Phi^{(q)}(0) = \Gamma^{(q)}$ of Definition 8.1; the complementation symmetry $\Phi^{(q)}(\lambda) = \Phi^{(1-q)}(-\lambda)$, which holds because $\delta^+(q) = \delta^-(1 - q)$; and $\Phi^{(1/2)}$ is the transfer-function strand's Φ term by term. [experimentally confirmed; e10_r105]

9.3 Tilted rigidity

Proposition 9.1 (the biased landscape responds at first order). $\Phi^{(q)'}(0) = \sum_{d \geq 1} \left[(1 - q)^{N_d} \delta_d^+ - q^{N_d} \delta_d^- \right]$, which vanishes identically at $q = \frac{1}{2}$ and is non-zero otherwise. So rigidity at a biased measure is not flatness: it is a tilt of a definite slope, and the slope is a point prediction with no free parameter.

[derived and verified, not proved from a hypothesis-free statement: it is a consequence of $\Phi^{(q)}$, itself derived by the partition-function argument of §9 and confirmed by three exact internal checks. The first-order coefficient is a point prediction with no free parameter, and §9.4 tests it against the naive alternative on 36 discriminating rows — which is why the two are distinguished here rather than only in §14.]

This is what §8.4 demanded of the derivation, and it is a sharper demand than it looks. The naive guess — keep the transfer-function strand's $\delta_d = d + \sigma_d/2$ and merely split the weights — gives instead $\sum [(1 - q)^{N_d} - q^{N_d}] \delta_d$, and the two differ by exactly $(\frac{1}{2} - q) \sum_d \sigma_d [q^{N_d} + (1 - q)^{N_d}]$: zero at $q = \frac{1}{2}$, and 20–32% of the prediction at the q tested. The measurement therefore decides between them.

9.4 Verification

The estimator must kill constants, because the point prediction of §8.3 carries an $O(10^{-3})$ offset at the biased centre which is the same size as the linear term. So we use a central difference about n_q : with $n_{\pm} = \lfloor (q \pm hq)T \rfloor$,

$$\hat{L}(h) = \frac{(\text{lm}_q/r_q)(n_+) - (\text{lm}_q/r_q)(n_-)}{\lambda_q(n_+) - \lambda_q(n_-)}.$$

Ratios of $\hat{L}$ to the two candidates, at $h = 0.01$:

profile	q	separation	$\widehat{L}/\Phi^{(q)'}(0)$	$\widehat{L}/\text{naive}$
odd numbers, $k = 120$	0.30	31.6%	0.9945	0.6805
odd numbers, $k = 120$	0.40	19.9%	0.9976	0.7991
odd numbers, $k = 120$	0.50	—	$\Delta(\text{lm}_q/r_q) = 0$ exactly	
odd numbers, $k = 120$	0.60	19.9%	0.9976	0.7991
primes, $k = 120$	0.30	32.3%	0.9981	0.6754
primes, $k = 120$	0.40	20.2%	0.9983	0.7967

Over all 36 discriminating rows (two profiles, $k = 80, 120$, four q , four h), the median $|\widehat{L}/\Phi^{(q)'}(0) - 1|$ is 0.0041 against 0.2038 for the naive candidate. The one row above 10% is the primes at $k = 80$, $q = 0.3$, $h = 0.01$, where $\Delta\lambda = 8.5 \cdot 10^{-5}$ and the integer rounding of $n_{\pm}$ is a visible fraction of the step; the same row at $k = 120$ reads 0.9981. At $q = \frac{1}{2}$ the numerator is 0 or $4.4 \cdot 10^{-16}$, below the stated floor — the fair coin has no linear response, which is the content of Proposition 9.1 at its zero. [\[experimentally confirmed; e10_r105\]](#)

The split is also visible stratum by stratum, which tests the offsets directly rather than through their sum: at $k = 80$, $q = 0.3$, over $_q/r_q$ against $(1 - q)^{N_d} e^{\lambda \delta_d^+} e^{-\lambda^2 q(1-q)s_d/2}$ agrees to within 0.24% over $d = 1, \dots, 6$, and under $_q/r_q$ against its partner to within 0.9%. [\[experimentally confirmed; e10_r105\]](#)

9.5 The biased hypothesis

(H $_q$): $\sum_d [(1 - q)^{N_d} (\delta_d^+)^2 + q^{N_d} (\delta_d^-)^2] < \infty$, uniformly for $q \in [q_0, 1 - q_0]$. Since $\max(q, 1 - q) \leq 1 - q_0 < 1$, the geometric damping of the transfer-function strand's (H) survives with base $1 - q_0$, and the constant degrades as $q_0 \rightarrow 0$ — which is the $\sim 40\times$ growth of the §8.3 error already measured between $q = \frac{1}{2}$ and $q = 0.2$. The $\alpha = 1$ threshold of the transfer-function strand is unchanged, the damping argument being base-agnostic. [\[derived; the degradation matches the measured trend\]](#)

10 Random sequences, unconditionally

Before the theorem, the floor it has to stand on. Let A be drawn from the Cramér model on the odd integers: each odd $m \in [3, N]$ is included independently with probability $2/\log m$, and N is fixed by $\mathbb{E}|A| = k$. Nothing is conditioned, so $|A|$ fluctuates from instance to instance. That is deliberate: $kR_A(k)$ is the k -invariant quantity of the transfer-function strand's 5.7, so an ensemble of genuinely different sizes tests the $1/k$ law for free. Twenty instances at each of two sizes, each through the whole paper-3 pipeline with its own N_d and σ_d :

$\mathbb{E} A $	range of k	(H) series	mean kR	sd	sd/mean
120	101–135	133.79–320.79	214.96	10.58	4.92%
180	146–191	132.94–471.77	206.32	6.94	3.37%

at $x = 0.06$; $c_A = kR/100$. Three things, and the third is the one to keep. [\[experimentally confirmed; e7_r102\]](#)

(i) *The hypothesis holds on the ensemble.* All forty instances have a bounded (H) series and a window that tracks the primes'. This is a precondition, not a result: a theorem about random sequences is unconditional only if the random sequences satisfy (H) themselves.

(ii) *The constant concentrates.* The spread across instances is under 5% and falls with k , while the mean of kR moves by 4% between the two sizes — the $1/k$ law of the transfer-function strand surviving an ensemble it was not fitted on.

(iii) *The primes sit inside the ensemble and the odd numbers do not.* Against the $k \approx 180$ ensemble (2.0632 ± 0.0694 in c_A), the primes score $z = -0.39$ and the odd numbers $z = -2.70$, at every one of the four targets. In this statistic the primes are indistinguishable from a random odd sequence of the same density; the odd numbers are a different object.

There is one place where the random model is *not* like the primes, and it is the place that matters for the minor arcs. Write $h_d(q) = \max_{(r,q)=1} (\prod_{a \in B_d} |\cos(\pi r a/q)|)^{1/|B_d|}$. Since $|\cos(\pi a/4)| = 1/\sqrt{2}$ for *every* odd a , the modulus-4 floor $h_d(4) = 1/\sqrt{2}$ is deterministic, not merely universal. Part II [5]’s vanishing lemma kills the modulus- $2m$ peak (m odd) as soon as B_d contains an element $\equiv 3 \pmod{6}$. A set of primes contains exactly one such element, namely 3, and 3 leaves B_d at $d \geq 2$, so the primes keep $h_d(6) = \sqrt{3}/2$; a random odd sequence contains about $|B_d|/3$ of them and loses the peak at every d . Measured over $q \leq 120$ at $d = 1, 2, 5$: every random instance has $\max_{q \geq 3} h_d(q) = 1/\sqrt{2}$ attained at $q = 4$, and the primes have $\sqrt{3}/2$ at $q = 6$. [\[experimentally confirmed; e7_r102\]](#) So the random theorem should come out with the geometric rate $1/\sqrt{2} = 0.7071$ per element rather than the $\sqrt{3}/2 = 0.8660$ of Parts I and II [4, 5]: *the primes are the harder case, and the modulus-6 ripple is a structural feature of the primes that the random model destroys.*

10.1 The modulus-4 theorem

The observation above is not an accident of the ensemble; it is a theorem about residues, and it is the exact counterpart of Part II’s modulus-6 theorem for a different equidistribution class. Part II averages $|\cos|$ over the *reduced* residues, which is the right class for the primes. A random odd sequence is equidistributed over the *odd* residues instead, so the relevant quantity is a different one.

Definition 10.1. For $q \geq 2$ let $R_q = \mathbb{Z}/q$ if q is odd and $R_q = \{j : j \text{ odd}\} \subset \mathbb{Z}/q$ if q is even — for q even and m odd, $m \bmod q$ is odd, so R_q is the set of residues an odd sequence can occupy. Put

$$M_{\text{odd}}(q) = \left(\prod_{j \in R_q} |\cos(\pi j/q)| \right)^{1/|R_q|}.$$

For r coprime to q the map $j \mapsto jr$ permutes R_q , so $M_{\text{odd}}(q)$ does not depend on the numerator of the rational r/q .

Everything below is a corollary of the coset identity of §3, because $\{j/q : j \in R_q\}$ is a coset of $(1/v)\mathbb{Z}$.

Corollary 10.2 (the identity in the form §10.2 uses). *Let $q \geq 2$, put $v = q$ for q odd and $v = q/2$ for q even, and let $\sigma = \frac{1}{2}$ if $q \equiv 2 \pmod{4}$ and $\sigma = 0$ otherwise. Then $|R_q| = v$ and, for every real s ,*

$$\prod_{j \in R_q} \left| \cos \pi \left(s + \frac{j}{q} \right) \right| = 2^{1-v} |\cos \pi (vs + \sigma)|, \quad \text{hence} \quad \frac{1}{v} \sum_{j \in R_q} X \left(s + \frac{j}{q} \right) \geq \left(1 - \frac{1}{v} \right) \log 2.$$

Proof. $\{j/q : j \in R_q\}$ is $\{k/v\}_{k < v}$ for q odd and $\{\frac{1}{2v} + \frac{k}{v}\}_{k < v}$ for q even. Apply Lemma 3.1 at $t = s$ and at $t = s + \frac{1}{2v}$ respectively. In the second case the shift contributes $v \cdot \frac{1}{2v} = \frac{1}{2}$ to the argument, so the total offset is $\tau_v + \frac{1}{2}$, which is an integer when v is even and $\frac{1}{2}$ when v is odd; and v odd with q even is exactly $q \equiv 2 \pmod{4}$. The inequality is Corollary 3.2. $\square$

Lemma 10.3 (closed form). $M_{\text{odd}}(q) = 2^{1/q-1}$ for q odd. For $q = 2u$ even, $M_{\text{odd}}(q) = 2^{1/u-1}$ if u is even, and $M_{\text{odd}}(q) = 0$ if u is odd.

Proof. Put $s = 0$ in Corollary 10.2: the right-hand side is 2^{1-v} when $\sigma = 0$ and 0 when $\sigma = \frac{1}{2}$. Take v -th roots. $\square$

Remark 10.4 (a second proof of the odd case, with no roots of unity). For q odd the closed form also follows from the double-angle formula alone, and this is the route we formalise. Let $P = \prod_{j=1}^{q-1} |\sin(\pi j/q)|$ and $C = \prod_{j=1}^{q-1} |\cos(\pi j/q)|$. From $\sin 2x = 2 \sin x \cos x$, $\prod_{j=1}^{q-1} |\sin(2\pi j/q)| = 2^{q-1} PC$. Since q is odd, $j \mapsto 2j \bmod q$ is a bijection of $\{1, \dots, q-1\}$ and $|\sin(2\pi j/q)| = |\sin(\pi(2j \bmod q)/q)|$, so the left side is P again. As $P \neq 0$, $C = 2^{1-q}$, and $\prod_{j \in R_q} |\cos(\pi j/q)| = C$ because the $j = 0$ factor is 1. Two proofs of one load-bearing identity check the *statement*; one proof checks only the argument. This is the branch that is machine-checked, as `prod_abs_cos_odd`; see Remark 10.6.

Theorem 10.5 (the modulus-4 theorem). $\max_{q \geq 2} M_{\text{odd}}(q) = \frac{1}{\sqrt{2}}$, attained at $q = 4$ and nowhere else.

Proof. $v \mapsto 2^{1/v-1}$ is strictly decreasing. For q odd, $q \geq 3$, Lemma 10.3 gives $M_{\text{odd}}(q) = 2^{1/q-1} \leq 2^{-2/3}$. For $q \equiv 2 \pmod{4}$ it gives 0. For $q \equiv 0 \pmod{4}$, write $q = 2u$ with u even, $u \geq 2$; then $M_{\text{odd}}(q) = 2^{1/u-1} \leq 2^{-1/2}$ with equality exactly when $u = 2$, i.e. $q = 4$. And $2^{-2/3} < 2^{-1/2}$. $\square$

Lemma 10.3 is verified against the product for all $q \leq 200$, worst discrepancy $4.4 \cdot 10^{-16}$; Corollary 10.2 is verified in its product form for $q \leq 120$ at 400 random shifts each, worst discrepancy $4.4 \cdot 10^{-16}$, and independently at 60 decimal digits, worst discrepancy $2.2 \cdot 10^{-61}$. [\[proved; and checked numerically, mq4_r107, stab_r110\]](#)

Remark 10.6 (what is machine-checked here, and what is not). The two halves of the argument have different statuses and the difference is worth stating. *The maximisation is formally verified.* Taking the closed form of Lemma 10.3 as a definition, `ModdCF_le` gives $M_{\text{odd}}(q) \leq 2^{-1/2}$ for every $q \geq 2$, `ModdCF_eq_iff` the equality case as a biconditional, `ModdCF_lt_of_ne_four` the strict bound off $q = 4$ by a second route, and `two_rpow_neg_half` the normalisation $2^{-1/2} = 1/\sqrt{2}$. This is the step where an error would hide — a case split over three residue classes and the comparison of $2^{-2/3}$ with $2^{-1/2}$ — and it is now checked by machine. *Of the closed form itself, the odd branch is now formalised and the even branch is not.* `prod_abs_cos_odd` proves $\prod_{j=1}^{q-1} |\cos(\pi j/q)| = (1/2)^{q-1}$ for every odd q , by the double-angle route of Remark 10.10 — which is exactly why that route is worth having: it uses one bijection and one trigonometric identity, where the proof of Lemma 3.1 needs a product over roots of unity that Mathlib does not carry. The even branch, and Lemma 3.1 in general, remain proved on paper and checked numerically. [\[Lean-verified \(maximisation, and the odd closed form\); Pnp/Theory/ModFour.lean, Pnp/Theory/OddProd.lean\]](#)

Proposition 10.7 (which residue classes the distribution relation reaches). [\[proved; two lines. Verified for \$q \leq 32\$ by direct enumeration, tool_coset_r138.\]](#) $(\mathbb{Z}/q)^*$ is an additive coset of a subgroup of $\mathbb{Z}/q$ if and only if q is a power of 2, in which case it is the set of odd residues.

Proof. Suppose $(\mathbb{Z}/q)^* = r + m\mathbb{Z}/q$. Every element of that set is $\equiv r \pmod{m}$; both 1 and -1 are units, so $1 \equiv -1 \pmod{m}$ and $m \mid 2$. If $m = 1$ every residue is a unit, so $q = 1$. If $m = 2$ every odd residue is a unit, so q has no odd prime factor. Conversely, for $q = 2^k$ the units are exactly the odd residues. $\square$

Proposition 10.8 (the reduced residues, evaluated). [\[proved; the proof is below and uses only Lemma 3.1 and Möbius inversion. Awaiting independent verification, per the convention of this project. Checked against direct enumeration on 26 moduli to \$10^{-12}\$, with two controls that fire; leadc_r143.\]](#) Write $v = 2^j w$ with w odd, and suppose $v \not\equiv 2 \pmod{4}$. Then

$$\frac{1}{\varphi(v)} \sum_{r \in (\mathbb{Z}/v)^*} X\left(\frac{r}{v}\right) = \log 2 \cdot \left(1 - \frac{2[w=1]}{2^j \varphi(w)}\right) = \begin{cases} (1 - 2^{1-j}) \log 2, & v = 2^j \ (j \geq 2), \\ \log 2 \text{ exactly,} & \text{otherwise, including every odd } v > 1. \end{cases}$$

For $v \equiv 2 \pmod{4}$ the residue w is a unit only when $w = 1$, and $X(w/v) = X(1/2) = +\infty$: the product $\prod(1 + e(a\theta))$ vanishes there rather than being merely small.

Proof. Put $t = 0$ in Lemma 3.1 with $v = n$ odd: $\tau_n = 0$ and $X(0) = 0$, so $G(n) := \sum_{k \bmod n} X(k/n) = (n-1) \log 2$.

For n even the odd residues mod n are the coset $\frac{1}{n} + \frac{2}{n}\mathbb{Z}/n$, i.e. a coset of the subgroup of order $n/2$; applying Lemma 3.1 with $v = n/2$ and $t = 1/n$ gives

$$O(n) := \sum_{\substack{s < n \\ s \text{ odd}}} X\left(\frac{s}{n}\right) = \frac{n}{2} \left[\left(1 - \frac{2}{n}\right) \log 2 + \frac{2}{n} X\left(\frac{1}{2} + \tau_{n/2}\right) \right] = \frac{n-2}{2} \log 2 \quad (4 \mid n),$$

since $n/2$ even makes $\tau_{n/2} = \frac{1}{2}$ and $X(1) = 0$; for $n \equiv 2 \pmod{4}$ one has $\tau_{n/2} = 0$ and $X(\frac{1}{2}) = +\infty$, which is the excluded case.

Now let $j \geq 2$. A residue is a unit mod v iff it is odd and prime to w , and writing $r = ds$ with $d \mid w$ (so d odd, and r odd iff s odd),

$$\sum_{r \in (\mathbb{Z}/v)^*} X\left(\frac{r}{v}\right) = \sum_{d \mid w} \mu(d) O\left(\frac{v}{d}\right) = \frac{\log 2}{2} \left[v \sum_{d \mid w} \frac{\mu(d)}{d} - 2 \sum_{d \mid w} \mu(d) \right] = \frac{\log 2}{2} [2^j \varphi(w) - 2[w=1]],$$

every term being finite because $4 \mid v/d$. Dividing by $\varphi(v) = 2^{j-1} \varphi(w)$ gives the formula. For $j = 0$ the same inversion run against G instead of O gives $\sum_{r \in (\mathbb{Z}/v)^*} X(r/v) = \log 2 [\varphi(v) - [v=1]]$, hence exactly $\log 2$ on average for $v > 1$. For $j = 1$ the term $d = w$ is $O(2) = X(\frac{1}{2}) = +\infty$ and the inversion is $\infty - \infty$; that is not a technicality but the vanishing case again. $\square \square$

Proposition 10.9 (the character decomposition; the other components are L -values). *[proved; the proof is three classical ingredients and a finite interchange. Awaiting independent verification. Checked to 10^{-39} on the even primitive characters of eight moduli, with the parity control and the vanishing control both firing; lval_r145.]* Let $f \geq 3$ and let χ be a primitive Dirichlet character mod f , with Gauss sum $\tau(\chi) = \sum_a \chi(a) e(a/f)$. Then

$$\sum_{a \in (\mathbb{Z}/f)^*} \bar{\chi}(a) X\left(\frac{a}{f}\right) = \begin{cases} \varphi(f) \cdot (\text{the value of Proposition 10.8}), & \chi = \chi_0, \\ 0, & \chi \text{ odd}, \\ -\tau(\bar{\chi}) (1 - \chi(2)) L(1, \chi), & \chi \text{ even}, \chi \neq \chi_0. \end{cases}$$

Remark 10.10 (the factor $1 - \chi(2)$, explained, and the attribution that follows). *[proved; $\log |2 \sin 2\pi t| = \log |2 \sin \pi t| + \log |2 \cos \pi t|$ is the double-angle formula. Checked to $2 \cdot 10^{-15}$ at $2 \cdot 10^5$ points, annealed_r147.]* Writing $T_2 f(t) = f(2t)$, the double-angle formula says $\log |2 \cos \pi t| = (T_2 - 1) \log |2 \sin \pi t|$, so $X = \log 2 - (T_2 - 1) \log |2 \sin \pi t|$. On the χ -component T_2 acts by $\chi(2)$, since $a \mapsto 2a$ permutes the units; hence the factor $1 - \chi(2)$ in Proposition 10.9 is the eigenvalue of $1 - T_2$, and nothing else. **The proposition is therefore Dirichlet's classical evaluation of $\sum_a \chi(a) \log |2 \sin(\pi a/f)|$ transported by the double-angle formula**, exactly as Lemma 3.1 is the classical distribution relation transported the same way. We claim no part of the evaluation; what we record is where it lands. The proof below is written out because it is short and because a reader should be able to see the transport rather than take it on trust.

Proof. X is even, so pairing a with $f - a$ kills every odd character. For χ even and nontrivial, start from the Fourier series $\log |2 \cos \pi t| = \sum_{n \geq 1} (-1)^{n+1} \cos(2\pi n t)/n$, valid for $t \notin \frac{1}{2} + \mathbb{Z}$, which holds at every a/f with $f \geq 3$. The sum over a is *finite*, so it may be taken inside term

by term. For primitive χ , $\sum_a \bar{\chi}(a)e(na/f) = \chi(n)\tau(\bar{\chi})$, and its real part is the same because χ is even. Hence, the $\log 2$ in $X = -\log |2 \cos \pi t| + \log 2$ dying against $\sum_a \bar{\chi}(a) = 0$,

$$\sum_a \bar{\chi}(a)X\left(\frac{a}{f}\right) = -\tau(\bar{\chi}) \sum_{n \geq 1} \frac{(-1)^{n+1}\chi(n)}{n} = -\tau(\bar{\chi})(1 - \chi(2))L(1, \chi),$$

the last step being $\sum_n (-1)^{n+1}\chi(n)/n = L(1, \chi) - 2 \sum_{n \text{ even}} \chi(n)/n = (1 - \chi(2))L(1, \chi)$. $\square$

Corollary 10.11 (a vanishing sum at $p \equiv 1 \pmod{8}$). *[proved, as the case $\chi = (\frac{\cdot}{p})$ of Proposition 10.9. Verified to 10^{-40} for all $p < 200$ with $p \equiv 1 \pmod{4}$; the $p \equiv 5 \pmod{8}$ half is the control and is nonzero in every case. `lua1_r145`.] For every prime $p \equiv 1 \pmod{8}$, $\sum_{a=1}^{p-1} \left(\frac{a}{p}\right) \log |\cos \frac{\pi a}{p}| = 0$, exactly. For $p \equiv 5 \pmod{8}$ the same sum is $-2\tau(\chi)L(1, \chi) \neq 0$.*

Proof. For $p \equiv 1 \pmod{4}$ the Legendre symbol is even, and it is ± 1 -valued, so $\chi = \bar{\chi}$. The factor $1 - \chi(2)$ is $1 - (\frac{2}{p})$, which vanishes exactly when $p \equiv \pm 1 \pmod{8}$; combined with $p \equiv 1 \pmod{4}$ that is $p \equiv 1 \pmod{8}$. For $p \equiv 5 \pmod{8}$ one has $(\frac{2}{p}) = -1$ and the factor is 2. $\square$

Corollary 10.12 (a class number formula for the coset energy). *[proved, as Proposition 10.9 composed with Dirichlet's class number formula. **Awaiting independent verification.** Verified for all eighteen primes $p \equiv 5 \pmod{8}$ below 320: the ratio is a positive integer in every case and equals the known class number in every case, including $h(229) = 3$. The fundamental unit is computed independently, by solving $x^2 - py^2 = \pm 4$, with no L -function anywhere in that half of the check. `class_r146`.] Let $p \equiv 1 \pmod{4}$ be prime, let $h(p)$ be the class number of the real quadratic field $\mathbb{Q}(\sqrt{p})$ and $\varepsilon_p > 1$ its fundamental unit. Then*

$$\sum_{a=1}^{p-1} \left(\frac{a}{p}\right) \log \left| \cos \frac{\pi a}{p} \right| = \begin{cases} 0, & p \equiv 1 \pmod{8}, \\ 4h(p) \log \varepsilon_p, & p \equiv 5 \pmod{8}. \end{cases}$$

Proof. For $p \equiv 1 \pmod{4}$ the Legendre symbol is an even primitive real character with $\tau(\chi) = \sqrt{p}$, so Proposition 10.9 gives the sum as $(1 - (\frac{2}{p}))\sqrt{p}L(1, \chi)$. The factor is 0 when $(\frac{2}{p}) = 1$, i.e. $p \equiv 1 \pmod{8}$, and 2 otherwise. Dirichlet's class number formula for the real quadratic field of discriminant p reads $L(1, \chi) = 2h(p) \log \varepsilon_p / \sqrt{p}$, and the two square roots cancel. $\square$

Remark 10.13 (what that sentence says). The left-hand side is a sum of logarithms of cosines at rational points — the same energy X that Lemma 3.1 governs and that §10 sums over a set of integers. The right-hand side is the class number of a real quadratic field times its regulator. *Nothing on the left knows it is supposed to be four times an integer times a regulator*, which is why the integrality is worth reporting as a check and not merely as a consequence.

We claim no novelty for the ingredients: the character decomposition is elementary Fourier analysis and the class number formula is Dirichlet's. What is worth recording is that the programme's energy function, arrived at from a subset-sum landscape, evaluates against quadratic characters to a class number — and that the vanishing half of the statement is exactly the condition that 2 splits.

The classical neighbourhood is named for the reader's benefit and ours. X is an element of the universal ordinary distribution [3], the object whose $\{\pm 1\}$ -cohomology Kubert and Sinnott computed and from which Sinnott's index formula relates the cyclotomic units to

the class number. *An identity carrying a class number out of that theory is the expected kind of output, not a surprising one.* We have not been able to consult those papers directly and so do not claim to know whether Corollary 10.12 appears there in this form; the honest statement is that it is the sort of thing that may. [\[a reading of Corollary 10.12; adds no claim. The novelty question is left open deliberately: this is the kind of identity that may well be classical, and §14 says so rather than implying otherwise.\]](#)

Remark 10.14 (where this puts the programme's one ineffective constant). Proposition 10.9 says the coset energy has a complete spectral decomposition over $(\mathbb{Z}/v)^*$ in which *the trivial component is exactly* $\log 2$ *with no error term at all*, and every other component is a Dirichlet L -value at $s = 1$ times a Gauss sum. For a set A the deficiency at $\theta = h/v$ is therefore

$$D\left(\frac{h}{v}\right) = \frac{|A|}{\varphi(v)} \sum_r X\left(\frac{r}{v}\right) + \frac{1}{\varphi(v)} \sum_{\chi \neq \chi_0} \bar{\chi}\text{-twisted count of } A \cdot \left(-\tau(\bar{\chi})(1 - \chi(2))L(1, \chi)\right),$$

so *the main term is exact and all of the error is carried by characters*. When A is a set of primes the twisted counts are the objects Siegel–Walfisz controls, and Part II's single ineffective constant [5] enters exactly there and nowhere else. That is a statement about the location of the ineffectivity, not a removal of it; what it buys is a route — under GRH the twisted counts are explicit, and so is the constant.

The factor $1 - \chi(2)$ deserves its own line. It kills the component whenever 2 lies in the kernel of χ , so the prime 2 governs which characters can see the energy at all. This is the fourth face of the same fact: Proposition 10.7 (which residues form a coset), Theorem 10.5 (the rate $1/\sqrt{2}$), Remark 10.15 (when the evaluation is a bound), and now the character group. *The obstruction of this theory is 2-adic, and it has now been met from four independent directions.* [\[derived from Proposition 10.9; the decomposition is proved, the reading of where the ineffectivity sits is a reading and is labelled as such. Nothing here is claimed to remove an ineffective constant.\]](#)

Remark 10.15 (the evaluation is a bound exactly at the powers of two). [\[proved, both directions, from the product form of Lemma 3.1; the failure is witnessed by an explicit counterexample. The closed form below was checked against direct computation at 35 points, shift_r144. This remark replaces a weaker and partly wrong statement of the same thing; see the note at its end.\]](#) Proposition 10.8 evaluates the average at $t = 0$, and a minor arc is a *neighbourhood* of a rational, so one must ask whether the value is also a floor for $\varphi(v)^{-1} \sum_r X(t + r/v)$. Exponentiating Lemma 3.1 answers it. For $4 \mid n$ the odd residues are the coset $\frac{1}{n} + \frac{2}{n}\mathbb{Z}$ of order $n/2$, so

$$O(n, t) := \prod_{\substack{s < n \\ s \text{ odd}}} |2 \cos \pi(t + \frac{s}{n})| = 2 |\cos(\pi \frac{n}{2} t)|,$$

and Möbius over the odd part of $v = 2^j w$ with $j \geq 2$ gives, every v/d again divisible by 4,

$$Q(v, t) := \prod_{r \in (\mathbb{Z}/v)^*} |2 \cos \pi(t + \frac{r}{v})| = 2^{[w=1]} \prod_{d|w} |\cos(\pi \frac{v}{2d} t)|^{\mu(d)}.$$

The average is bounded below by its value at $t = 0$ exactly when $Q(v, t) \leq Q(v, 0)$, and the product reads off both cases:

- $w = 1$, *i.e.* $v = 2^j$: $Q = 2 |\cos(\pi \frac{v}{2} t)| \leq 2 = Q(v, 0)$. **Uniform, and proved** — as it must be, since the reduced residues are then one coset and Corollary 3.2 already covers it.
- $w > 1$: Q is a *ratio*, the divisors with $\mu(d) = -1$ sitting in the denominator, and a ratio of cosines is not bounded by 1. For $w = 3$ it is $|\cos 3u|/|\cos u| = |4 \cos^2 u - 3|$, which

exceeds 1 whenever $\cos^2 u < \frac{1}{2}$. **The bound fails.** At $v = 12$, $t = \frac{1}{5}$ the average is 0.4525 against $\log 2 = 0.6931$ at $t = 0$: it falls by 0.2406.

So the evaluation becomes a bound on the surrounding arc *if and only if v is a power of two* — which is Proposition 10.7 once more, and the reason is the same one: that is exactly when the reduced residues are a coset. Read together with Remark 10.16: at $\theta = h/v$ the deficiency is $|A| \log 2$ for every v outside the powers of two, and at none of those v does that survive a shift.

Note on how this remark got here. Its first version reported the $4 \mid v$, $w > 1$ case as *measured uniform*, on a scan of t over $[0, 1/v]$. That interval is the period of the *full-group* average, where $t \mapsto t + 1/v$ permutes $\{k/v\}$; it is not a period of the reduced-residue average, because $t \mapsto t + 1/v$ carries r/v to $(r+1)/v$ and the units are not closed under $+1$. The scan was therefore blind to the region containing the counterexample. We record it because the error is invisible in the output — a shorter scan returns a number, not a warning.

Remark 10.16 (Theorem 10.5 is the case $w = 1$). Read Proposition 10.8 as a statement about how much a minor arc can save. For a set whose residues equidistribute over $(\mathbb{Z}/v)^*$ — the primes, by Dirichlet — the deficiency at $\theta = h/v$ is $|A| \log 2$ per element, i.e. the *full* circle mean, at every v except the powers of two; and at $v = 2^j$ it is $(1 - 2^{1-j}) \log 2$, minimised at $v = 4$ where it is $\frac{1}{2} \log 2$. Since the bound sits in an exponent, that reads

$$\left| \prod_{a \in A} (1 + e(a\theta)) \right| \leq 2^{|A|} e^{-|A| \log 2} = 1 \quad \text{in general,} \quad \leq 2^{|A|} e^{-|A|(\log 2)/2} = (\sqrt{2})^{|A|} \quad \text{at } v = 4,$$

so the rate per element is $1/\sqrt{2}$ at $v = 4$ and 1 elsewhere — *at the rational point itself*, which by Remark 10.15 is all one may say outside the powers of two. *That is Theorem 10.5*, and the proposition says in addition that the obstruction occurs at the powers of two *and nowhere else* — which is Proposition 10.7, reached here from an exact evaluation rather than from the group-theoretic characterisation. Two independent routes to one classification. [\[derived from Proposition 10.8; adds no claim, and is included because the two theorems were proved separately and it is worth recording that they are one fact.\]](#)

Remark 10.17 (why the two extremal moduli differ, in one line). Lemma 3.1 averages over an *additive* coset. The residues available to an odd sequence are an additive coset for every q — all residues for q odd, the odd residues for q even — so the floor of Corollary 3.2 always applies to it. The residues available to a set of primes are the *reduced* residues, and by Proposition 10.7 those are an additive coset exactly when q is a power of 2. Hence:

- $q = 2^k$: the two classes coincide, the floor applies to the primes as well, and the value is pinned at $2^{1/v-1}$ with $v = 2^{k-1}$ — maximal at $q = 4$, giving $1/\sqrt{2}$.
- $q = 2m$, m odd > 1 : the odd residues contain $j = m$, where $\cos(\pi m/q) = 0$, and the reduced residues do not. An odd sequence therefore meets the zero and a set of primes misses it. That is the modulus-6 peak.
- q odd: the reduced residues are everything but 0, and the primes sit at the full-circle mean $\log 2$.

So Theorem 10.5 and Part II's modulus-6 theorem are not two coincidences but one fact about q , and the smallest q at which the two classes part company is 6. [\[derived from Propositions 10.7 and 10.8; the three bullets are exactly the three cases of the latter, which is where the third one \(\$q\$ odd, value \$\log 2\$ \) is now proved rather than asserted. Both endpoints are proved — here and in Part II \[5\] — so this remark reorganises rather than adds. Measured: at \$q = 4\$ and \$q = 8\$ the odd primes below \$4 \cdot 10^4\$ sit on the floor to five decimals, and at \$q = 6\$ and \$q = 10\$ they fall below it while the odd numbers rise far above; `too1\_coset\_r138`.\]](#)

Remark 10.18 (the two classes, side by side).

	class of residues	extremal q	rate per element
primes (Part II)	reduced	6	$\sqrt{3}/2 = 0.8660$
random odd A (here)	odd	4	$1/\sqrt{2} = 0.7071$

The two disagree exactly at $q \equiv 2 \pmod{4}$: $M(6) = \sqrt{3}/2$ but $M_{\text{odd}}(6) = 0$, because $j = 3$ is an odd residue mod 6 and $\cos(\pi/2) = 0$. A set of primes contains at most one element $\equiv 3 \pmod{6}$, and it leaves B_d at $d \geq 2$; a random odd sequence contains a positive proportion of them at every d . *So the primes are the harder case, by a factor $\sqrt{3}/2 = 1.2247$ in the geometric rate*, and Part II's $\sqrt{3}/2$ is a feature of the primes rather than a limitation of the method.

Both halves of that comparison are machine-checked. `abs_cos_pi_mul_div_four` gives $|\cos(\pi a/4)| = 1/\sqrt{2}$ for *every* odd a , so `prod_abs_cos_four` makes the modulus-4 floor $(1/\sqrt{2})^{|A|}$ *deterministic* rather than merely universal; and `cos_pi_mul_div_six_eq_zero` gives $\cos(\pi a/6) = 0$ whenever $a \equiv 3 \pmod{6}$, so `prod_abs_cos_six_eq_zero` kills the modulus-6 peak on a single element. `six_dies_four_survives` is the separation itself: for an odd sequence containing such an element the modulus-6 product is 0 while the modulus-4 product is still exactly $(1/\sqrt{2})^{|A|} > 0$. [\[Lean-verified; Pnp/Theory/OddPeaks.lean\]](#)

10.2 The rate, and what is still missing

Theorem 10.19 (minor-arc rate for random odd sequences). *Let $A = A(N)$ be drawn from the Cramér model above and let $F_A(\theta) = \prod_{a \in A} |\cos(\pi a \theta)|$. Write $b = |A|$. The running parameter is N , and $b = b(N) \rightarrow \infty$ almost surely; the limit below is along $N \rightarrow \infty$. Then almost surely*

$$\lim_{b \rightarrow \infty} \max_{\theta \text{ minor}} F_A(\theta)^{1/b} = \frac{1}{\sqrt{2}}.$$

Here *minor* means $\|\theta\| \geq 1/N$, with $\|\cdot\|$ the distance to the nearest integer; Step 2'' shows that this radius is large enough, so nothing is left unowned.

[proved, with every constant written out. The proof is the four steps below rather than a proof environment, so it is worth saying here that it is one: no step is deferred, no constant is fitted, and the only external *input* is the arithmetic half, Theorem 10.5, which is proved in this paper. That is a statement about inputs and not about exposition: three standard estimates are applied in the form we need without being proved, and Remark 10.21 names them. The two sentences are compatible and are meant to be read together, which is why this one points at the other.]

The arithmetic half is Theorem 10.5. Here is the rest, in the steps it takes, with the constants named. Fix the Dirichlet parameter $Q = \lfloor \sqrt{N} \rfloor$ — Step 2' says why that value and not another — and set

$$M = \log \frac{2Q}{\pi} + 1, \quad \boxed{\delta = \frac{1}{6} \log 2 = 0.1155245 \dots}$$

Write $X_m(\theta) = -\log |\cos(\pi m \theta)| \in [0, \infty]$, so that $-\log F_A(\theta) = \sum_m \xi_m X_m(\theta)$ with ξ_m the independent indicators of the model. Only a *lower* bound on this sum is wanted, which is what makes the argument easy.

Step 1 (truncate; the unbounded tail is free). Put $X_m^{(M)} = \min(X_m, M)$. Then $\sum \xi_m X_m \geq \sum \xi_m X_m^{(M)}$, and the summands now lie in $[0, M]$. At $\theta = \frac{1}{4}$ truncation is inactive, since $X_m(\frac{1}{4}) = \frac{1}{2} \log 2$ for every odd m .

Step 2 (the mean, and the constant δ). At $\theta = r/q$ the values X_m over odd m run through $\{-\log |\cos(\pi j/q)| : j \in R_q\}$, so by Lemma 10.3 their mean is $-\log M_{\text{odd}}(q) = (1 - 1/v) \log 2$

and the *gap* over the value $\frac{1}{2} \log 2$ at $q = 4$ is

$$\delta(q) = \left(\frac{1}{2} - \frac{1}{v}\right) \log 2, \quad v = q \text{ (} q \text{ odd)}, \quad v = q/2 \text{ (} q \equiv 0 \pmod{4}), \quad \delta = \infty \text{ (} q \equiv 2 \pmod{4}).$$

This vanishes exactly at $v = 2$, i.e. $q = 4$, and increases with v ; the smallest admissible v after 2 is $v = 3$. **So the binding competitor is $q = 3$ and the gap is $\delta = \frac{1}{6} \log 2$** , with $q = 8$ next at $\frac{1}{4} \log 2$. Step 2' will replace the entry $\delta = \infty$ at $q \equiv 2 \pmod{4}$ by the finite value $(\frac{1}{2} - \frac{1}{u}) \log 2$: off the rational point the vanishing residue no longer vanishes, and $q = 6$ then *ties* with $q = 3$ at $\frac{1}{6} \log 2$. The constant is unchanged and two moduli bind it rather than one. The two moduli with $v \leq 2$ — $q = 4$, where the floor is $\frac{1}{2} \log 2$ exactly and is meant to be, and $q = 2$, where $v = 1$ and the floor degenerates to 0 — are handled in Step 2'', together with $q = 1$. [proved from Lemma 10.3; checked against the direct residue average for all $q \leq 400$, worst discrepancy $1.0 \cdot 10^{-15}$, and the minimum located at $q = 3$ over $q \leq 2000$; delta_r109]

Truncation does not eat this. For $q \not\equiv 2 \pmod{4}$ one has $|\cos(\pi j/q)| \geq \sin(\pi/2q)$, so $\max_j X_m \leq \log(2q/\pi) < M$ and truncation is *inactive* at every $q \leq Q$ — measured, the largest value over all such $q \leq 1000$ is 6.4552 against $M - 1 = 6.4562$, so the bound $\log(2Q/\pi)$ is sharp and the +1 in M is there to keep the argument off a knife edge. For $q = 2u \equiv 2 \pmod{4}$ truncation *is* active, at the single odd residue $j = u$ where the cosine vanishes; putting $x = -1$ in $(x^u + 1)/(x + 1)$ gives $\prod_{j \neq u} |\cos(\pi j/q)| = u 2^{1-u}$, so the truncated mean is $[(u - 1) \log 2 - \log u + M]/u$, which exceeds $\frac{1}{2} \log 2 + \delta$ by between 0.23 and 1.25 over $q \leq 198$ at $Q = 200$. [experimentally confirmed, exact to $4.4 \cdot 10^{-16}$; delta_r109]

Step 2' (from the rationals to every θ : there is no dent). Write $\theta = a/q + \beta$ with $q \leq Q$ and $|\beta| \leq 1/(q(Q + 1))$ (Dirichlet). The obstacle here looks like a job for a quantitative equidistribution estimate — Erdős–Turán or Koksma — and it is not: Lemma 3.1 settles it as an *identity*. Since $j \mapsto ja$ permutes R_q , the values $\{ja/q\}$ are the $\{j/q\}$, and for any common shift s

$$\frac{1}{|R_q|} \sum_{j \in R_q} X\left(s + \frac{j}{q}\right) = \underbrace{\left(1 - \frac{1}{v}\right) \log 2}_{-\log M_{\text{odd}}(q)} + \frac{1}{v} X(vs + \sigma) \geq -\log M_{\text{odd}}(q), \quad (*)$$

because $X \geq 0$. *Shifting the whole rational grid can only raise the mean*; the dent one expects when θ moves off a rational does not exist, and no discrepancy theory is needed to bound it. Equality in $(*)$ holds exactly on $vs + \sigma \in \mathbb{Z}$, so the rational values are the true minima, not merely convenient sample points. [proved (Lemma 3.1); the minimum over s of the left side of $(*)$ is checked to equal $-\log M_{\text{odd}}(q)$ to $3 \cdot 10^{-16}$ for 15 moduli, and to survive truncation at M for every $q \leq 1000$; stab_r110]

First, truncation does not spoil $(*)$. The loss $(X - M)^+$ is nonzero only within e^{-M}/π of a pole, and the coset points are $1/v \geq 1/Q > 1/(Qe) = 2e^{-M}/\pi$ apart, so at most one residue is affected; if it lies at distance d from the pole then $X \leq \log \frac{1}{2d}$, while the excess term of $(*)$ is $\frac{1}{v} X(vs + \sigma) \geq \frac{1}{v} \log \frac{1}{\pi v d}$ because $vs + \sigma$ sits at distance vd from a pole. The floor therefore survives as soon as $M \geq \log(\pi v/2)$, and with $M = \log(2Q/\pi) + 1$ and $v \leq Q$ that reads $1 \geq \log(\pi^2/4)$ — true, **with a margin of exactly $1 - \log(\pi^2/4) = 0.0968$** . This is the third place the +1 in M earns itself. [proved; and $\min_s \Phi_q^{(M)}(s) - (1 - 1/v) \log 2$ is checked to be $\geq -5.6 \cdot 10^{-16}$ for every $q \leq 1000$; stab_r110, rate_r112]

Step 2' continued (the two error terms, in full). What $(*)$ does not by itself cover is that element m of the class j sits at $ja/q + m\beta$, so the shift varies *within* a class. Fix $q \notin \{1, 2\}$ and group the odd $m \in [3, N]$ by residue class mod q . In a class the m form an arithmetic progression of common difference $2v$ (for q odd the odd numbers in a class mod q step by $2q = 2v$; for $q = 2u$ they step by $q = 2v$), so the phases form an arithmetic progression of step $\eta = 2v\beta$ sweeping an arc of length $L = N|\beta|$, and the class starts at $x_j + m_j\beta$ where $m_j \leq 2v$ is its first element. Three errors, and nothing else:

(i) *Progression versus integral.* $X^{(M)}$ is of bounded variation, at most $2M$ per period, so for a progression of step η over an arc of length L , $|\sum_r X^{(M)} - \eta^{-1} \int| \leq 2M(L+1)$. There are v classes and $\asymp N/2$ terms in all, so in the mean this costs at most $4qM/N + 4M/(Q+1)$, using $v|\beta| \leq 1/(Q+1)$.

(ii) *Jitter.* Replacing the arc of class j by the common arc is a translation by $\epsilon = m_j\beta$, and $m_j \leq 2v$ while the progression step is $\eta = 2v\beta$, so $|\epsilon| \leq |\eta|$ *exactly*. Translating by ϵ changes the integral by at most ϵV , so the per-class error is at most $\epsilon V/\eta \leq V \leq 2M(L+1)$ — the same shape as (i), and not a separate mechanism at all. Per element, $4qM/N + 4M/(Q+1)$.

(iii) *Unequal class counts.* They differ by at most 1, costing vM in the sum, i.e. $2qM/N$ in the mean.

After (i)–(iii) the integrand is the honest coset average, which is $\geq (1 - 1/v) \log 2$ pointwise by (*); integrating a pointwise inequality needs no further hypothesis, and in particular *no case distinction on the length of the sweep* — which is what makes the route robust in Q where an arc-comparison design is not. Collecting,

$$\frac{\sum_m p_m X_m^{(M)}(\theta)}{\sum_m p_m} \geq \left(1 - \frac{1}{v}\right) \log 2 - E, \quad E \leq \frac{10qM}{N} + \frac{8M}{Q+1}.$$

Both terms are minimised together at $Q = \lfloor \sqrt{N} \rfloor$, where $E = O(M/\sqrt{N})$ — which is why the Dirichlet parameter is $\sqrt{N}$ and not, as one first writes, N . At $Q = N$ the argument fails twice over: the residue classes hold $O(1)$ elements each, and $|\beta| \leq 1/(qN)$ makes every sweep shorter than a single period.

The weights need no argument at all. $p_m = 2/\log m$ is decreasing, so by Abel summation a bound valid for every prefix $\{m \leq N'\}$ passes to the p -weighted mean unchanged: if $S(N') \geq cP(N') - \mathcal{E}(N')$ for all N' , then $\sum p_m X_m^{(M)} \geq c \sum p_m - \sum_k (p_k - p_{k+1}) \mathcal{E}(k) - p_N \mathcal{E}(N)$, and with $\mathcal{E}(N') \leq CM(\sqrt{N} + N'/\sqrt{N})$ the subtracted term is $O(M\sqrt{N}) + O(M/\sqrt{N}) \sum p_m$, i.e. $O(M \log N/\sqrt{N})$ relative. There is no block-weight condition to check. [\[proved; the shortfall at the binding modulus \$q = 3\$ is measured as \$1.9 \cdot 10^{-3}, 5.8 \cdot 10^{-4}, 1.7 \cdot 10^{-4}, 4.8 \cdot 10^{-5}\$ at \$N = 2001, 8001, 32001, 128001\$, and deficit \$\cdot N/\(qM\)\$ is \$0.3057, 0.3140, 0.3199\$ at \$N = 8001, 32001, 128001\$ — constant, i.e. the measured deficit has exactly the derived shape \$qM/N\$. Errors \(i\), \(ii\), \(iii\) are additionally isolated and measured *one at a time* against their own bounds, over three sizes, eight moduli and three offsets, worst values \$1.1 \cdot 10^{-2}, 1.2 \cdot 10^{-2}\$ and \$8.5 \cdot 10^{-4}\$ against bounds of order \$10^{-1}\$; \[stab_r110\]\(#\), \[rate_r112\]\(#\), \[audit_r114\]\(#\)\]](#)

Step 2'' (the three moduli with $v \leq 2$, and the major arc). Corollary 3.2 gives $\frac{1}{2} \log 2$ at $q = 4$ ($v = 2$) and nothing at all at $q \in \{1, 2\}$ ($v = 1$). The last two are not prose; they are the following lemma, and $q = 4$ needs no more than the identity because $\frac{1}{2} \log 2$ is exactly the value the theorem is about.

Lemma 10.20 (the degenerate moduli). *Let $\phi(w) = -\log |\sin \pi w|$ and $L > 0$. Then $\frac{1}{L} \int_0^L \phi = \log 2 + \text{Cl}_2(2\pi\{L\})/(2\pi L)$, and*

$$\frac{1}{L} \int_0^L \phi > \frac{1}{2} \log 2 \quad \text{for every } L > 0, \quad \frac{1}{L} \int_0^L X \geq \log 2 - \frac{\text{Cl}_2(\pi/3)}{2\pi} = 0.5316 \quad \text{for } L \geq 1.$$

Consequently the mean over a $q = 2$ sweep exceeds $\frac{1}{2} \log 2$, and so does the mean over a $q = 1$ sweep once $L = N\|\theta\| \geq 1$ — which is exactly the definition of minor.

Proof. The closed form is $\int_0^x \phi = x \log 2 + \text{Cl}_2(2\pi x)/(2\pi)$ together with the periodicity of ϕ . For $L \leq \frac{1}{2}$: ϕ is convex on $(0, 1)$, so the average over $[0, L]$ is at least the midpoint value $-\log \sin(\pi L/2) \geq -\log \sin(\pi/4) = \frac{1}{2} \log 2$; at the endpoint $L = \frac{1}{2}$ the two-point midpoint bound $\frac{1}{2}[\phi(L/4) + \phi(3L/4)] = 0.51986$ shows the inequality is strict there. For $L \geq \frac{1}{2}$: $1/(2\pi L) \leq 1/\pi$ and $|\text{Cl}_2| \leq \text{Cl}_2(\pi/3)$, so the average is at least $\log 2 - \text{Cl}_2(\pi/3)/\pi = 0.3701 >$

$\frac{1}{2} \log 2$. For $q = 1$ the same computation with X in place of ϕ gives $\log 2 + \text{Cl}_2(2\pi L + \pi)/(2\pi L)$, and $L \geq 1$ makes the correction at most $\text{Cl}_2(\pi/3)/(2\pi) = 0.1615$. $\square$

The true $q = 2$ floor is 0.4945 at $L = 0.7908$ — above $\frac{1}{2} \log 2 + \delta$, not merely above $\frac{1}{2} \log 2$ — but that is measured, not proved, and the theorem does not need it. The smallest L at which the $q = 1$ average reaches $\frac{1}{2} \log 2 + \delta$ is 0.4528, so the radius $1/N$ in the definition of *minor* is comfortable rather than tight. [\[proved; the Clausen closed form checked against direct quadrature to \$4.5 \cdot 10^{-5}\$, and the convexity bound verified at 5000 values of \$L\$; rate_r112\]](#)

Step 3 (concentration). The summands are independent and lie in $[0, M]$, and there are $n = \lfloor (N-1)/2 \rfloor$ of them, so Hoeffding gives $\Pr[\sum \xi_m X_m^{(M)} < \mathbb{E} - t] \leq \exp(-2t^2/(nM^2))$. Take $t = \varepsilon \mathbb{E}|A|$ and use $\mathbb{E}|A| \asymp 2N/\log N$, $n \asymp N/2$, $Q = \lfloor \sqrt{N} \rfloor$ and hence $M \asymp \frac{1}{2} \log N$: the failure probability at one θ is

$$\exp(-c\varepsilon^2 N/\log^4 N), \quad c = 64.$$

Two readings of the same bound, and the theorem uses the second. With $\varepsilon = \frac{1}{2}\delta$ this is the fixed-margin statement, $\exp(-cN/\log^4 N)$ with $c = 16\delta^2 = 0.2135$. But the theorem only needs the mean to survive to within $o(1)$, so one may instead take $\varepsilon = N^{-1/4}$, which still leaves the exponent $\asymp \sqrt{N}/\log^4 N \rightarrow \infty$. *That is why $q \in \{1, 2\}$ need only clear $\frac{1}{2} \log 2$ and not $\frac{1}{2} \log 2 + \delta$* , and it is what makes Lemma 10.20 sufficient. No second moment of an exponential sum appears: the randomness is in the *sequence*, not in the *phases*.

Step 4 (net, then Borel–Cantelli). Where $X_m < M$ we have $|\cos(\pi m\theta)| \geq e^{-M}$, hence $|\partial_\theta X_m^{(M)}| = \pi m |\tan(\pi m\theta)| \leq \pi N e^M \leq 2eNQ \leq 2eN^{3/2}$, so a net of spacing $\varrho/(2eN^{3/2}b)$ suffices and has $O(N^{5/2}/\varrho)$ points. Since $N^{5/2} \exp(-c\varepsilon^2 N/\log^4 N)$ is summable in N for $\varepsilon = N^{-1/4}$, the union bound and Borel–Cantelli give the almost-sure statement.

Assembling. Steps 1–4 give, almost surely and uniformly over minor θ , $b^{-1} \sum_{a \in A} X_a(\theta) \geq \frac{1}{2} \log 2 - o(1)$, i.e. $\max_\theta \min_{\text{minor}} F_A(\theta)^{1/b} \leq 2^{-1/2}(1 + o(1))$: at $q = 4$ the floor is $\frac{1}{2} \log 2$ by Corollary 3.2, at every other $q \geq 3$ it is $\frac{1}{2} \log 2 + \delta$, at $q \in \{1, 2\}$ it is above $\frac{1}{2} \log 2$ by Lemma 10.20, and the error of Step 2' is $O(M \log N/\sqrt{N})$. The matching lower bound is free: $\theta = \frac{1}{4}$ is minor and $X_a(\frac{1}{4}) = \frac{1}{2} \log 2$ for *every* odd a , so $F_A(\frac{1}{4})^{1/b} = 2^{-1/2}$ identically. Hence the limit is $2^{-1/2}$. $\square$ [\[proved\]](#)

Remark 10.21 (what “proved” means here, and what it does not). Theorem 10.19 carries no hypothesis and every constant in its proof is written down. One qualification, stated here rather than left for a referee to raise. The three error terms of Step 2' are each reduced to a standard estimate — a Riemann-sum comparison for a function of bounded variation, and Abel summation against a monotone weight — which we state and apply in the form we need but do not prove; a referee would be within rights to ask for them in full, and doing so would lengthen Step 2' without changing any constant. *No hypothesis is not the same as written out in full*, a distinction Part II now also makes explicitly about its own main theorem, and the two should be read the same way.

Measured, at $\theta = r/q$ for $q \leq 12$ and eight instances at $b \approx 190$: the mean of $b^{-1} \sum_a -\log |\cos(\pi a/q)|$ agrees with $-\log M_{\text{odd}}(q)$ within the $O(b^{-1/2})$ fluctuation at every q , and at $q = 4$ it is 0.346574 with standard deviation 0.000000 — deterministic, as Definition 10.1 says it must be. [\[experimentally confirmed; mq4_r107, delta_r109\]](#)

Remark 10.22 ($\theta = \frac{1}{4}$ is the limit of the maxima, not a maximum). A first scan of $F_A(\theta)^{1/b}$ over a fine net asked whether anything exceeds $1/\sqrt{2}$, and found that something does. That is not a counterexample and the reason is worth stating, because it fixes the form of Theorem 10.19: the maximum is *approached* at $\frac{1}{4}$, and the statement is about the limit of $\max_\theta F_A^{1/b}$, not about where a maximum sits at finite b . Moving off $\frac{1}{4}$ by δ changes $\log F_A^{1/b}$ by $\delta b^{-1} \sum_a (\pm \pi a) + O(\delta^2)$, and the signs depend on $a \bmod 4$, so the first-order term is a

random walk of size $b^{-1}\sqrt{S_2}$. Optimising δ gives an excess of $O(1/b)$ in the exponent — a *bounded multiplicative factor* on F_A , not a geometric one. Measured over $b = 111, 190, 310, 542$: $\max_{\theta} F_A^{1/b} - 1/\sqrt{2}$ is $9.0 \cdot 10^{-4}, 1.8 \cdot 10^{-3}, 1.1 \cdot 10^{-4}, 1.2 \cdot 10^{-3}$, and $b \log(\sqrt{2} \max_{\theta} F_A^{1/b})$ stays in $[0.05, 0.89] = O(1)$, as derived. The argmax sits at $\theta = 0.24996, \dots, 0.25001$. [\[experimentally confirmed; mq4_r107\]](#)

11 Two questions from the transfer-function strand, settled

(i) *The c_d fold-in is settled, and it goes the transfer-function strand's way.* The share of the residual that c_d carries falls like $k^{-2\alpha}$ — measured exponents 1.98, 2.35, 2.99 against predicted 2.00, ≈ 2.4 , 3.00 on the odd numbers, the primes and $\alpha = \frac{3}{2}$, with the squares already below the numerical floor. So Φ as printed is right for every larger k and c_d belongs to E_k permanently; the transfer-function strand's 5.22 now says so. [\[experimentally confirmed, four profiles, \$k \leq 300\$; e8_r103\]](#)

(ii) *$c_A = c'_A$ is still open, but sharper.* Carried to $k = 300$ the ratio reaches 1.0064 (odd numbers), 1.0001 (primes), 0.9984 ($\alpha = \frac{3}{2}$), and the gap closes like $k^{-2.4}$ — faster than the $1/k$ law both constants sit on. [\[experimentally confirmed; e8_r103\]](#) **[TODO the mechanism is still missing, and a measured rate of convergence is evidence about a limit, not a proof of one. F45 stands: substituting s for λ makes the residual three times worse.]**

12 What is missing, stated once

Proposition 12.1 (the first correction to the annealed prediction). [\[derived, from the same local limit theorem the rest of the paper uses — so it inherits Appendix A's standing and no more. Verified against exact dynamic programming on three profile families: the ratio of measured to predicted relative error is 0.98 at \$k = 90\$ in all three, including the family for which \(H\) fails. \[correction_r150.\]\(#\)\]](#) At the centre $n = T/2$, write $L_d = \sum_{a \leq 2d} a$, $s_d = \sum_{a \leq 2d} a^2$, $\delta_d = d + L_d/2$ and $\sigma_d^2 = \sigma^2 - s_d/4$. Applying the local limit theorem to $B_d = \{a \in A : a > 2d\}$ layer by layer gives

$$\frac{\text{lm}_A(n)}{r_A(n)} = 1 + 2 \sum_{d \geq 1} 2^{-N_A(d)} \frac{\sigma}{\sigma_d} \exp\left(-\frac{\delta_d^2}{2\sigma_d^2}\right),$$

both signs of the offset contributing the same δ_d at the centre. Expanding to first order in σ^{-2} ,

$$\frac{\text{lm}_A(n)}{r_A(n)} = \Gamma(A) \left(1 - \frac{Q(0)}{\sigma^2}\right) + o(\sigma^{-2}), \quad Q(0) = \frac{1}{\Gamma(A)} \sum_{d \geq 1} 2^{-N_A(d)} (\delta_d^2 - \frac{s_d}{4}).$$

Proposition 12.2 (the correction at a general target). [\[derived, by the same layer expansion as Proposition 12.1 and inheriting its standing. Verified on three profile families at eight values of \$z\$: away from the crossing the ratio of measured to predicted deviation lies in \$\[0.87, 1.10\]\$. \[target_r151.\]\(#\)\]](#) Let $z = (n - \mu)/\sigma$. Then the layer prediction evaluates, with no further approximation, to

$$\frac{\text{lm}_A(n)}{r_A(n)} = 1 + 2 \sum_{d \geq 1} 2^{-N_A(d)} \frac{\sigma}{\sigma_d} \exp\left(-\frac{z^2 s_d/4 + \delta_d^2}{2\sigma_d^2}\right) \cosh\left(\frac{z\sigma\delta_d}{\sigma_d^2}\right),$$

and to first order in σ^{-2} ,

$$\frac{\text{lm}_A(n)}{r_A(n)} = \Gamma(A) \left(1 - (1 - z^2) \frac{Q(0)}{\sigma^2}\right) + o(\sigma^{-2}).$$

Corollary 12.3 (the ratio crosses Γ at one standard deviation). *[derived from Proposition 12.2; the crossing itself is measured, and it is the sharpest falsifiable consequence this expansion has. At $k = 90$ for the odd numbers the deviation is $+8.2 \cdot 10^{-5}$ at $z = 0$, $+1.7 \cdot 10^{-5}$ at $z = 0.9$, $+2.3 \cdot 10^{-6}$ at $z = 1$, and $-1.5 \cdot 10^{-5}$ at $z = 1.1$. target_r151.]* To first order the correction carries the factor $1 - z^2$. Hence $\text{lm}_A(n) < \Gamma(A) r_A(n)$ for $|z| < 1$, and $>$ for $|z| > 1$: the ratio sits below the annealed value inside one standard deviation of the centre and above it outside, crossing at $|n - \mu| = \sigma$.

Remark 12.4 (the first-order term is Hermite; the expansion is not). *[the identification $1 - z^2 = -\text{He}_2(z)$ is algebra. The failure of the natural second-order guess is measured, and is recorded as a caution and not as a claim: nothing in this paper depends on it. hermite_r153.]* Proposition 12.2 can be written

$$\frac{\text{lm}_A(n)}{r_A(n)} = \Gamma(A) \left(1 + \text{He}_2(z) \frac{Q(0)}{\sigma^2} \right) + o(\sigma^{-2}), \quad \text{He}_2(z) = z^2 - 1,$$

and the same Hermite polynomials organise the Edgeworth expansion of Appendix A. The natural guess is therefore that the next term is a multiple of He_4 at order σ^{-4} . *It is not.* Removing the He_2 term and rescaling by σ^4 leaves a residual whose ratio to $\text{He}_4(z)$ varies by 44–68% across z and whose mean does not settle with k (about $5.4 \cdot 10^3$, $9.9 \cdot 10^3$, $1.4 \cdot 10^4$ at $k = 24, 30, 36$). The guess is wrong and we say so rather than leaving the pattern to be extrapolated.

The obstruction is visible in Appendix A’s own budget, and it is worth naming because it explains why the *first* order works so well. The local limit theorem carries its own corrections of relative size $\beta = K^{(4)}/\sigma^4 \asymp k^{-1}$ and $\alpha^2 \asymp k^{-1}$, both *far larger* than the layer-geometry term $Q(0)/\sigma^2 \asymp k^{-3}$. They do not appear in the measured ratio because they cancel between lm and r — which is precisely the cancellation §12’s Step 4 analysis found. So the expansion is not in powers of σ^{-2} alone; the surviving σ^{-2} term is what is left after a cancellation of much larger ones, and there is no reason for the residue of that cancellation to be a single Hermite polynomial.

Remark 12.5 (the primes, where the theorem is unconditional). *[the crossing is derived from Proposition 12.2 and measured on the odd primes at $k = 40, 65, 90$, window fully occupied at every point; away from the crossing the ratio of measured to predicted deviation lies in $[0.94, 1.00]$ at $k = 90$. The value $Q(0) = \frac{61}{3}$ for the odd numbers is exact, summed in closed form and not fitted. primes_r152.]* The primes are the sharpest test available, because Part II [5] proves the theorem there at the centre with no hypothesis attached. The crossing appears exactly as predicted: at $k = 90$ the relative deviation runs $+3.24 \cdot 10^{-5}$ at $z = 0$, $+7.0 \cdot 10^{-6}$ at $z = 0.9$, $+9.8 \cdot 10^{-7}$ at $z = 1$, and $-5.8 \cdot 10^{-6}$ at $z = 1.1$. (For scale, Γ computed from the first 65 odd primes already agrees with $\Gamma(P)$ to $3 \cdot 10^{-15}$.)

Two constants are worth recording beside it. For the odd numbers the series can be summed: $N_d = d$, $L_d = d^2$, $s_d = d(4d^2 - 1)/3$ and $\delta_d = d + d^2/2$ give

$$\sum_{d \geq 1} 2^{-d} \left(\delta_d^2 - \frac{s_d}{4} \right) = 61 \quad \text{exactly,} \quad \text{so} \quad Q(0) = \frac{61}{3},$$

and with $\sigma^2 = k(4k^2 - 1)/12$ the relative error at the centre is asymptotically $61/k^3$. For the odd primes $Q(0) = 50.4369 \dots$

And then the comparison, which points the other way from the difficulty of the proofs. $Q(0)$ is only 2.5 times larger for the primes, while σ^2 is far larger at the same k — $p_k \sim k \log k$ against $2k$ — so the annealed prediction is **more accurate for the primes than for the odd numbers**, by a factor growing like $(\log k)^2/4$: measured, the ratio of the two relative errors is 1.80 at $k = 40$ and 4.74 at $k = 520$. The primes are the harder case to prove and

the easier case to approximate, and the two facts have the same cause — their elements are larger, which is what makes the exponential sums hard and the variance big.

Remark 12.6 (why the crossing is worth stating). Nothing in the annealed picture predicts a crossing, and nothing in the main theorem sees it: the limit is Γ at every target, so the whole phenomenon lives inside the error term. It is also the cleanest test the expansion admits — a *sign*, at a *stated place*, with no constant to fit. If a reader wants to falsify Proposition 12.1 cheaply, this is where to aim.

Remark 12.7 (what the correction says about (H)). $Q(0)$ is not a new quantity: it is the one already tabulated in §7 (20.3 for the odd numbers, 50.4 for the primes, 916 for the squares, $3.8 \cdot 10^5$ for the cubes). So the table that was there to show *how far out of reach* a profile is turns out to be, up to σ^{-2} , the error itself.

Three consequences, and the third is the point.

1. *The rate is explained rather than fitted.* For the odd numbers $Q(0) \rightarrow 20.33$ and $\sigma^2 \asymp k^3/3$, so the relative error is $\asymp 61k^{-3}$; for $a_i \asymp i^{3/2}$ the variance grows like k^4 and the error like k^{-4} . Both match the measurement.
2. *(H) is boundedness of the numerator, and the conclusion needs less.* What the limit requires is only $Q(0) = o(\sigma^2)$. (H) asks $Q(0) = O(1)$, which is strictly stronger.
3. *For power profiles the weaker condition always holds.* Distinctness forces $a_i \asymp k^{1-\alpha}i^\alpha$, hence $\sigma^2 \asymp k^3$ for every α , while $Q(0) \asymp k^{3(1-\alpha)}$; the ratio is $k^{-3\alpha} \rightarrow 0$ for every $\alpha > 0$. So Problem 12.9 is answered affirmatively for that family *by the correction formula* rather than by extrapolation.

[derived from Proposition 12.1; item 3 is a computation with the profile $a_i \asymp k^{1-\alpha}i^\alpha$ and inherits that proposition's standing. It does *not* settle Problem 12.9 in general — only for clean power profiles.]

Remark 12.8 (where this sits beside the rigorous work on number partitioning). [a positioning statement, not a result. The papers named are cited from their abstracts and titles; we have not worked through them, and say so rather than implying we have.] Subset-sum at the centre is the number partitioning problem, and that problem has a rigorous literature this programme should be read against. Borgs, Chayes and Pittel established the phase transition and finite-size scaling for random integer partitioning; Bauke, Franz and Mertens conjectured that the local energy statistics near a target are those of a random energy model, and Borgs, Chayes, Mertens and Nair proved that local REM conjecture, in two parts, with a threshold beyond which it fails.

The relation to Remark 13.1 is one of complement rather than overlap, and the difference is worth stating precisely because it is the whole of our claim to novelty here:

	local REM	this series
instance	<i>random</i> , and the statement is distributional	<i>deterministic</i> , e.g. the odd primes
observable	the spacings of the energy spectrum	the count of metastable states
output	convergence to a Poisson process	an <i>exact constant</i> $\Gamma(A)$

Both are annealed-type statements and both come with a threshold — theirs on the energy scale, ours the hypothesis (H) of Problem 12.9. Neither implies the other as far as we can see: a Poisson limit for the spectrum of a random instance says nothing about the ratio of metastable to ground states of one fixed instance, and our constant says nothing about spacings. *What we would want from a reader who knows that literature is whether the*

deterministic, exact-constant statement is new there, or whether it is a known corollary we have arrived at sideways. Given that the last three identities in this paper turned out to be classical transports, the prior should not be generous.

Problem 12.9 (is (H) necessary, or does it only set the rate?). [\[open, and stated here rather than asserted anywhere: the evidence below is measured, not proved, and by the standing rule of this project a measured claim goes into the open register or nowhere. breaks_r148.\]](#) Theorem 5.1 concludes under (H), and Proposition 7.1 cuts the power profiles at $\alpha = 1$. What is not known is whether the *conclusion* fails when (H) does. Computed exactly — no sampling — by the classification of Remark 13.1, with $\text{lm}_A(n)/r_A(n)$ averaged over 41 targets at the centre:

profile	(H)	relative error at $k = 90$	fitted p in Ck^{-p}
odd numbers ($\alpha = 1$)	holds	$8 \cdot 10^{-5}$	2.89
$a_i \asymp i^{3/2}$	holds	$1 \cdot 10^{-5}$	3.82
$a_i \asymp 4\sqrt{ki}$ ($\alpha = \frac{1}{2}$)	fails	$5.5 \cdot 10^{-4}$	1.7

Over $20 \leq k \leq 90$ the annealed prediction is approached in *all three* cases, including the one where (H) fails; what changes is the exponent, from about k^{-3} to about k^{-1} . *So the accessible evidence says (H) governs the rate rather than the truth* — and Proposition 12.1, found after this problem was posed, *predicts* the exponent as 3α instead of fitting it, which settles the clean power profiles (Remark 12.7) and leaves this problem standing for everything else. *The exponent quoted above for $\alpha = \frac{1}{2}$ was first reported as 0.98; that fit was contaminated by the smallest size, where the measured error changes sign, and over the remaining sizes it is about 1.7 against a predicted 1.5.* We do not conclude in general: $k \leq 90$ is small, the profile family is one family, and a convergence that is merely slow is indistinguishable at this range from one that stalls at a different limit. The question is therefore posed and not answered:

does $\text{lm}_A(n)/r_A(n) \rightarrow \Gamma(A)$ hold for every family of distinct odd integers, with (H) controlling only the rate?

A proof either way would sharpen Remark 13.1 into the statement the statistical-mechanics reader actually wants — *the annealed count is exact, full stop, and here is what the hypothesis buys* — or produce the first profile where the annealed count is provably wrong, which would be worth more.

A candidate witness, found after the problem was posed. Remark 12.7 says where to look: $Q(0)/\sigma^2$ must be forced up, and that needs a family growing faster than any power. Take $a_i = 2^i + 1$ — odd, distinct, and not super-increasing beyond $i = 3$, so the subset sums genuinely overlap. Then $\Gamma = k + 2$ grows linearly and the computed ratio does not follow it:

k	8	10	12	14	16	18
Γ	10	12	14	16	18	20
lm/r at the centre	7.18	7.45	7.73	7.33	6.92	7.08
ratio	0.718	0.621	0.552	0.458	0.385	0.354

lm/r stays near 7 while Γ grows, and $a_i = 2^i - 1$ behaves the same (0.713, 0.583, 0.473, 0.379 at $k = 8, 11, 14, 17$). The dynamic programme agrees with brute-force enumeration of all 2^k subsets at $k = 8, 10, 12$ with no disagreement, and the window occupancy is printed, so neither a coding error nor the support trap of §12 accounts for it. Note also that $Q(0)/\sigma^2 \approx 0.07$ here while the shortfall is 65%: *the first-order correction does not explain this failure either*, so these families lie outside the regime Proposition 12.1 describes.

The problem therefore splits in two, and only the first is the question we were asking.

(hrate-a) — *the convergent regime, and the real open problem.* Does $\text{lm}_A(n)/r_A(n) \rightarrow \Gamma(A)$ hold, without (H), for every family with $\limsup_k \Gamma(A_k) < \infty$? **Open, with no witness either way.** Everything the correction formula settles (Remark 12.7) lies inside this regime, and the geometric witness lies outside it.

(hrate-b) — *the divergent regime.* For $a_i = 2^i + 1$ the annealed series *itself* diverges: $N_A(d) \approx \log_2 d$, so the layer weights are $\approx 1/d$ and the partial sums grow like $\Gamma(D) = 3 + \log_2 D$. The measured ratio stays bounded while $\Gamma = k + 2$ grows, and over $10 \leq k \leq 16$ it sits at $\Gamma(D)$ for an effective depth $D_{\text{eff}} = 2$ — only the first two layers ever fire, because the deeper truncation targets are not representable. *Conjecturally the ratio stays bounded in this regime;* we do **not** claim it equals a fixed truncation, because extending to $k \leq 20$ moves the implied D_{eff} between 2 and 4 (median 4.0, 5.0, 4.5, 4.0, 5.0 at $k = 16, \dots, 20$), and a plateau seen on $k \leq 14$ alone is a plateau seen on too short a range.

A note on which number is being conjectured about, because two are in circulation. The r -weighted value $\sum_n \text{lm}(n)/\sum_n r(n)$ — the ratio at a typical *ground state* — runs 6.9 to 7.7; the median of $\text{lm}(n)/r(n)$ over representable targets — the ratio at a typical *target* — runs 4 to 5. Both are bounded, which is all (hrate-b) asserts, and they differ by a factor near 1.8 because in this family lm/r is larger exactly where r is larger, so weighting by r is not a neutral choice. `fverify_r163_reconcile`. It is a conjecture because $k \leq 18$ is a range and not a limit. Recording it here rather than asserting it anywhere is the third of the three outs the standing rule allows. [the witness is **measured**, exactly and with a brute-force control, over $6 \leq k \leq 18$ in two families; the failure of the limit is **conjectured** and lives in this problem's open register. `geom_r154`, `geomctl_r154`.]

Three statements of this Part were, for most of its life, conditional — and conditional on the *same* thing. Rather than attach a differently worded caveat to each, we stated the missing ingredient once, here, and let each consumer point at it. **It is now supplied, and this section is kept as the record of what was missing and of how it was closed**, because a reader who wants to know how much of this Part rests on a single argument is entitled to see that question answered in the place where it was originally posed.

Problem 12.10 (the shared missing ingredient). Write out the Edgeworth expansion for the tilted local limit theorem of Proposition 5.20 in region R1, uniformly over the window $|m - \mu_d(s)| \leq W(k)$, with explicit constants at the cumulant budget

$$\varepsilon = \frac{|K_d''''|}{(K_d'')^{3/2}} + \frac{K_d^{(4)}}{(K_d'')^2} + \frac{W(k)|K_d''''|}{(K_d'')^2}.$$

This is a classical local-limit computation under a product measure. It is given in §5.3 only to the level of its cumulant budget; what is absent is the explicit constants, not the method.

[**CLOSED.** Regions R2 and R3 of Proposition 5.20 were already proved — they quote Lemma 5.8 and Lemma 5.10 — and R1 is Proposition A.7 in Appendix A, with every constant proved rather than measured. By the standing rule of this project an argument written by one hand becomes “proved on paper” only after a second, independent reading, and that reading is now on record — in three parts, each covering the text as it then stood: the lemmas (r162), the repairs landed against them (r164), the restated proposition with the T_* construction and all five explicit constants (r171), each constant re-derived from scratch. The head of Appendix A says plainly what that does and does not amount to. *The problem is retained rather than deleted*: what a paper was missing, and what closed it, is part of what a reader is entitled to know about it.]

Who was waiting on it, and what closing it did. Exactly three *numbered statements* — and, transitively through Theorem 4.1, one unnumbered sentence: the algorithmic reading of §13, which until R1 closed had to say that the convergence there was a proof skeleton rather than a theorem. *We count environments and say so, because the justification below is about how far the damage would have reached, and damage does not respect environments.* Closing Problem 12.10 completed all of them at once:

- **Proposition 5.20** is unconditional; R2 and R3 were already proved and R1 is Proposition A.7.
- **Theorem 4.1** (rigidity of the gap series) is a theorem, with no conditional clause.
- **Theorem 5.1** (the transfer function) is a theorem, with no conditional clause.

No other statement of this Part depended on R1: the coset identity, the biased invariant, the modulus-four theorem and the random rate are independent of it, as is every entry of §14 not listed above. *That was worth checking before the closure and it is worth recording after it,* because the value of naming a single shared ingredient is that one can say exactly how far the damage would have reached had it failed.

13 Where this sits, and the algorithmic reading

Remark 13.1 (what the main theorem says to someone who works on disordered systems). [proved; the derivation is three lines from the classification of local minima in Part I [4] and is given in full here. The equivalence used in the first line was additionally checked by brute force on 1952 (set, target) pairs, and the identification of the two expressions for Γ on five families; annealed_r147.] The classification says: S is a strict local minimum for the target n if and only if, writing $D = \sigma(S) - n$, every $a \in S$ satisfies $a > 2D$ and every $a \notin S$ satisfies $a > -2D$. For a fixed offset $D = d > 0$ that condition is simply $S \subseteq \{a \in A : a > 2d\}$, and the number of elements it forbids is exactly $N_A(d) = \#\{a \in A : a \leq 2d\}$.

Now apply the *annealed* — that is, independence — heuristic: pretend the number of subsets of a subfamily $B \subseteq A$ hitting a prescribed target is proportional to $2^{|B|}$. Then the layer at offset d contributes a fraction $2^{-N_A(d)}$ of the ground-state count, the offsets $\pm d$ pair up, and $d = 0$ contributes 1. Summing,

$$\frac{\text{lm}_A(n)}{r_A(n)} \approx 1 + 2 \sum_{d \geq 1} 2^{-N_A(d)} = \Gamma(A).$$

So Γ is not merely computed by an argument that avoids the annealed approximation: Γ **is** the annealed prediction. What Theorem 5.1 and Part II establish is therefore

the annealed count of metastable states is asymptotically exact for this model,

under a hypothesis (H) that is checked rather than assumed, with §7 exhibiting a profile where it fails. That is the sentence we would offer to the statistical mechanics of disordered systems, where the annealed approximation is in constant use and is almost never controlled: here is a model with an NP-complete ground-state problem in which it is provably right, together with the condition that delimits it and an explicit case that violates the condition.

Two honest limits. The statement is about the *ratio* and not about each layer separately; and it is one model, so it is evidence about a heuristic rather than a theorem about it.

13.1 Iterating the coarse-graining

Lemma 3.1 is used once, in one direction, and that is not the only thing it says. Written as an operator, $(\mathcal{C}_v f)(t) = v^{-1} \sum_{k < v} f(t + k/v)$ acts on X by $\mathcal{C}_v X = (1 - 1/v) \log 2 + v^{-1} X(v \cdot + \tau_v)$, so X spans a one-dimensional affine invariant subspace of every $\mathcal{C}_v$ at once, and the constants compose: $\mathcal{C}_{v'} \mathcal{C}_v = \mathcal{C}_{vv'}$ on this subspace, with released constants adding as $(1 - 1/v) + v^{-1}(1 - 1/v') = 1 - 1/(vv')$. Three questions we have not attempted, listed so that the omission is deliberate rather than silent. (i) Which energies other than X have this property — is $-\log |\cos|$ characterised by being an exact fixed point of the family $\{\mathcal{C}_v\}$ up to affine terms? (ii) Is there a transfer-operator reading in which Φ of the transfer-function strand is the spectral object and $\mathcal{C}_v$ the renormalisation map? (iii) Parts I and II and the transfer-function strand’s landscape language was a source of questions; here it produced an identity. Does the same happen for the *biased* energy of §9, where the coset structure is weighted by q ? [\[conjectural directions; nothing in this paper depends on any of them\]](#)

13.2 Related work

Subset sums and number partitioning have a large landscape literature, and it is worth saying precisely where the present series does and does not meet it. That literature is overwhelmingly about *random instances*: the phase transition and barrier-tree statistics of the partitioning landscape, the local-REM description of the near-optimal spectrum, and more recently overlap-gap obstructions for local algorithms. Its objects are the energy spectrum and the dynamics; its randomness is in the instance; and its answers are laws for a typical instance.

Parts I and II and the transfer-function strand do something different in kind: they fix the sequence and compute an *arithmetic invariant of that set*, $\Gamma(A) = 1 + 2 \sum_d 2^{-\#\{a \in A: a \leq 2d\}}$, which the ratio lm/r converges to for *every* target. Nothing is averaged over instances, and the answer is a number one reads off A . The present paper’s §10 is the one place where the two meet — there the sequence is random — and even there the question is the ratio’s limit rather than the spectrum’s law.

One vocabulary warning, since the words collide. “Biased Bernoulli measure” is established terminology in the theory of Bernoulli convolutions, where it denotes a self-similar measure on $\mathbb{R}$; that is a different object from P_q here, and we do not claim the term. Within the subset-sum landscape literature we have not found the biased deformation $\Gamma^{(q)}$ of §8, nor the odd-residue extremal quantity M_{odd} of §10.1, under any name. [\[literature pass, not a proof of novelty: absence of evidence over the searches made\]](#)

13.3 The algorithmic reading

This programme began with local search on subset-sum instances, and the invariant answers a question one asks there. If a local-search procedure ends at a uniformly chosen strict local minimum, the expected number of attempts before landing on a ground state is lm/r , and Parts I and II and the transfer-function strand predict that this number converges to $\Gamma(A)$ — *computable from A alone, without solving a single instance, and independent of the target*. For the odd numbers it is 3; for the primes, 5.3492... One condition is attached to that sentence, and we attach it here rather than leaving it in §5: the uniformity of the terminal distribution is an assumption about the search, not a fact about the landscape. Until r171 two were attached, and they were of different kinds: the other was that the convergence here was a proof skeleton rather than a theorem, which is a defect of the paper and not a hypothesis of the sentence. Theorem 4.1 is now a theorem and that one is gone. **The remaining one is about the algorithm and not about the landscape**, and nothing in this series bears on

it: removing it would mean *deriving* terminal uniformity, which would be a theorem about a search procedure and not about Γ . We have no impossibility proof and claim none; we claim only that the assumption is of a different kind from everything else here, and that a reader should not expect the rest of the programme to discharge it.

Theorem 8.5 is the part of this subsection that has a proof, and it is a statement about *sampling*, not about search dynamics: drawing configurations with a biased coin instead of a fair one multiplies the ratio by $\Gamma^{(q)}/\Gamma \geq 1$, strictly for $q \neq \frac{1}{2}$. So a natural one-parameter family of sampling rules cannot beat the fair coin on this statistic — which is a claim about the statistic, and becomes a claim about a restart heuristic only under the uniformity assumption just named. Section 9 adds that the biased landscape also acquires a *first-order* sensitivity to the target that the fair one does not have.

Behind the uniformity assumption is the reason this is an outlook and not a theorem about algorithms: a count of local minima is not a distribution over basins of attraction, and if the basins are very unequal the ratio bounds the wrong quantity. Nor does anything here address dynamics — barrier heights, escape times, or the behaviour of any particular heuristic. What the invariant supplies is a target-independent constant that any such analysis has to be consistent with, together with a proved statement that one obvious way of trying to improve that constant does not improve it. **[TODO if this section is ever to make an algorithmic claim, the missing object is the basin-size distribution, not more landscape statistics.]**

14 Honest scope

14.1 The deformed and random strand

- *Lean-verified*: Lemma 8.3 in all three forms (Remark 8.4), including the equality case; and the maximisation step of Theorem 10.5 (Remark 10.6), the closed form being taken as a definition there; and the two peak facts of Remark 10.18 (Pnp/Theory/OddPeaks.lean); and the odd branch of Lemma 10.3 itself, by the double-angle route (Pnp/Theory/OddProd.lean, Remark 10.10). The canon also passes an independent kernel replay (lean4checker replays every constant of the import closure of the root module Pnp from the imports up, with a negative control). Each file named above is inside that closure; a file present in the source tree but outside it is *not* replayed, and the harness now fails if one exists, so this sentence cannot quietly become false.
- *Proved*: Lemma 8.3, Theorem 8.5 and Corollary 8.7; Lemma 3.1 (twice) with Corollaries 3.2 and 10.2; Theorem 10.5 and Lemma 10.3; Lemma 10.20; and Theorem 10.19, all constants written out. Also Remark 8.2’s derivation of Definition 8.1, modulo the biased flatness hypothesis, which is not yet stated.
- *Experimentally confirmed, with the range stated*: the point prediction (§8.3, two profiles, five q , three k), the stratum assignment (Remark 8.17), the linear response under bias (§8.4).
- *Derived and verified*: $\Phi^{(q)}$ of (5) and Proposition 9.1 — three exact internal checks, and the first-order point prediction confirmed against the naive alternative on 36 discriminating rows (§9.4); the random-ensemble floor of §10; the c_d verdict of §11.
- *Interpretation, with nothing deduced from it*: the renormalisation reading of Lemma 3.1 (Remark 3.5) and the three questions of §13.
- *Not yet written*: the passage from (H_q) to a rigidity statement.

14.2 The transfer-function strand

In the style of Part I [4] §7.3 and Part II [5] §11, here is what rests on what.

- *Formally verified*: the combinatorial layer, inherited from Part I, for every target.
- *Proved*: the bridge of §5.2 — Lemma 5.8 in full, Lemma 5.10 in all three parts (with (c) quoted verbatim from Part II’s Step 2 at $n = 1$, and the covering apparatus of that paper not needed here, see Remark 5.16), Lemma 5.17, and Proposition 7.1. The minor-arc input is m -independent, so it transfers to off-centre targets unchanged; Part II’s ineffectivity through the Siegel–Walfisz constant is inherited at the single point recorded in Remark 5.13.
- *Proved, and unconditional*: Theorems 4.1 and 5.1, together with Proposition 5.20. Regions R2 and R3 of that proposition quote the bridge; region R1 is Proposition A.7 of Appendix A, where the Edgeworth expansion is written out with every constant proved rather than measured. *This entry read “proof skeleton, with the analytic ingredients in place” until r171, and the change is recorded rather than made silently*: what closed the gap was the appendix plus a second, independent reading of it. Remark 5.23 records why those constants had to be derived rather than quoted from the Bernoulli-part-extraction literature, and it *proves* the point rather than asserting it: our summands have Bernoulli part exactly 0, so the quantity that literature’s error term is divided by is exactly zero and its bound reads $\leq \infty$; and the standard repair, blocking, fails outright because our weights make the variance grow like the cube of the number of summands. *The claim is about that source and is not a claim about every possible import*.
- *Experimentally confirmed, with the range stated*: rigidity (§4), the transfer function to 0.8–1.9% of the signal (Table 1), universality (§6), the profile-dependent constant $4(2\alpha + 1)/(\alpha + 1)$ verified on four profiles.
- *Not claimed*: the residual 0.8–1.9% of Table 1 is consistent with finite size and shrinks in k , but we do not claim it is zero.

Remark 14.1 (the single ineffective constant). Precisely: every estimate used here is effective except part (b) of Lemma 5.10, whose error term is the Siegel–Walfisz error, and Part II’s appendix identifies the same constant C_1 as the only ineffective one in that paper, entering there at one point as well. Part I is unconditional and effective throughout. **Ineffective is not conditional**: no statement here depends on an unproved hypothesis; what is not available is a numerical threshold beyond which the asymptotics take hold.

Remark 14.2 (approaches that did not work). As in Part II we keep the dead ends, with reasons. *The Γ_x family*: the idea that off-centre targets are described by a family of tilted gap series indexed by x . Rejected by the data — rewriting the observed collapse in the exact tilt parameter makes it worse, not better. *The multiplicative bridge*: $|\tilde{G}_B| \leq |G_B|^\kappa$, which would have carried every exponent of Part II across at once. Impossible for the reason in Remark 5.9, and the underlying inequality runs the other way. *The residual as a tilt-versus-slope artefact*: the constants of Remarks 2.5 and 5.7 agree to two digits on every profile, but substituting s for λ makes the residual three times worse — a shared driver, not a cause. *A half-integer exponent* reported in an early draft, which was an artefact of the Gaussian slope (Remark 2.3). *Restating (H) so that the cubes violate it*: the series is bounded for the cubes and unbounded for $\alpha < 1$, and the hypothesis follows the computation rather than the narrative.

A Region R1: the expansion, with every constant written out

[**proved, and independently read — in three parts.** This appendix closes Problem 12.10. The standing rule of this project is that an argument written by one hand counts as proved only after a second, independent reading, and *proved* and “proved on paper” name the same rung: the quotation marks distinguish a proof on paper from a proof machine-checked in Lean, not a weaker grade of proof. *The reading is on record in three parts, and it is worth saying why not as one.* The three lemmas of §A.1–§A.3 were read line by line in r162; the three repairs those readings prescribed were landed and confirmed in r164; the restated Proposition A.7, the T_* construction and all five explicit constants of ρ were read in r171, each constant re-derived from scratch and found to agree symbolically, including the reductions $16eC_T/\pi = e(2 + \pi)/\pi$ and $41580/82944 = 385/768$. Each part covered the text as it then stood, and between them they cover every line of this appendix in its present form — but **no single reading has covered the appendix as a whole**, because the text changed between the parts in response to the earlier ones. That is the honest description, and a reader who wants the stronger one is entitled to know it has not been supplied. Every constant below is proved, not measured; where a measurement exists it is quoted only to show how much slack the proof carries.]

The reduction to $s > 0$, stated once (R-a). Theorem 4.1 is about $\rho \in (0, 1)$ and this appendix works at $s > 0$, that is, below the centre. The two are the same statement: the complementation $S \mapsto A \setminus S$ gives $\text{lm}_A(n) = \text{lm}_A(T - n)$ and $r_A(n) = r_A(T - n)$ — proved in Part I [4] and verified in Lean — so we may assume $\rho \leq \frac{1}{2}$ throughout, hence $s \geq 0$; and $s = 0$ is covered because Lemma A.1 is stated for every $p \in (0, \frac{1}{2}]$ rather than for $p < \frac{1}{2}$. Both consumer theorems cite this paragraph.

Throughout, B is a finite set of odd positive integers, $p_a = (1 + e^{sa})^{-1}$ with $s > 0$ so that $p_a < \frac{1}{2}$ for every a , $S = \sum_{a \in B} aX_a$ with X_a independent Bernoulli(p_a), and

$$\mu = \sum a p_a, \quad \sigma^2 = K'' = \sum a^2 p_a q_a, \quad N = \max B, \quad z = \frac{m - \mu}{\sigma}, \quad \alpha = \frac{K'''}{\sigma^3}, \quad \beta = \frac{K^{(4)}}{\sigma^4}.$$

A.1 The two-point remainder

Lemma A.1. [*proved*] Let $p \in (0, \frac{1}{2}]$, $q = 1 - p$, and let $P_4(v) = ipv - \frac{pq}{2}v^2 - \frac{i}{6}\kappa_3v^3 + \frac{1}{24}\kappa_4v^4$ with $\kappa_3 = pq(1 - 2p)$ and $\kappa_4 = pq(1 - 6pq)$. Then for every real v with $|v| \leq 1$,

$$|\log(1 - p + pe^{iv}) - P_4(v)| \leq C_T |v|^5, \quad C_T = \frac{2(2 + \pi)}{2^5} = 0.32136\dots$$

Proof. Write $K(t) = \log(q + pe^t)$, so that $\log(1 - p + pe^{iv}) = K(iv)$ and $\kappa_j = K^{(j)}(0)$.

Analyticity. $q + pe^t = 0$ forces $e^t = -q/p$, hence $\text{Im } t \equiv \pi \pmod{2\pi}$ and $\text{Re } t = \log(q/p)$. The nearest such point to the origin is at distance $|\log(q/p) + i\pi| \geq \pi$. So K is analytic on $|t| < \pi$, for every p . And the branch is the principal one throughout (R-c): $q + pe^t$ meets the cut $(-\infty, 0]$ only where $\text{Im}(pe^t) = 0$ with $q + pe^t \leq 0$, which again forces $\text{Im } t \equiv \pi \pmod{2\pi}$ and so cannot happen on $|t| \leq 2$. Hence $\log$ is the principal logarithm on the whole disc used below and $|\text{Im } K| < \pi$ is rigorous rather than a convention — the one step where a reader is entitled to pause.

A lower bound on $|q + pe^t|$ for $|t| = 2$. Write $t = x + iy$ with $x^2 + y^2 = 4$, so $|y| \leq 2$ and $\cos y \geq \cos 2$. With $u = pe^x > 0$,

$$|q + pe^t|^2 = q^2 + 2qu \cos y + u^2 \geq q^2 + 2qu \cos 2 + u^2 \geq q^2(1 - \cos^2 2) = q^2 \sin^2 2,$$

the last step by minimising the quadratic in u at $u = -q \cos 2$. Hence $|q + pe^t| \geq q|\sin 2| \geq \frac{1}{2}|\sin 2| = 0.4546\dots$ using $q \geq \frac{1}{2}$.

An upper bound on $|K|$ for $|t| = 2$. Also $|q+pe^t| \leq q+pe^2 \leq e^2$. Since $|\log w| \leq |\log |w|| + \pi$ for $w \neq 0$, and $|\log |q+pe^t|| \leq \max\{2, |\log 0.4546|\} = 2$, we get $M_2 := \max_{|t|=2} |K(t)| \leq 2 + \pi$.

Cauchy, and the tail. By Cauchy's estimate on $|t| = 2$, $|\kappa_j| \leq j! M_2 / 2^j$. Therefore, for $|v| \leq 1$,

$$|K(iv) - P_4(v)| = \left| \sum_{j \geq 5} \kappa_j \frac{(iv)^j}{j!} \right| \leq M_2 \sum_{j \geq 5} \left(\frac{|v|}{2} \right)^j = M_2 \frac{(|v|/2)^5}{1 - |v|/2} \leq 2M_2 \left(\frac{|v|}{2} \right)^5,$$

which is the claim. $\square$ $\square$

Remark A.2 (how much slack the proof carries). The proved constant is 0.3214; the maximum of $|K(iv) - P_4(v)|/|v|^5$ over $p \leq \frac{1}{2}$ and $|v| \leq 1$ is measured at 0.00114, so the proof is loose by a factor 282. That costs nothing: only the exponent of $|v|$ enters the scaling of §A.2, and the constant enters once, linearly. [\[experimentally confirmed; r1proof_r141\]](#). The same script checks the two hand-derived inequalities of the proof: the singularity is never closer than π over $p \in (0, \frac{1}{2}]$, and $|q + pe^t| \geq |q| \sin 2|$ holds on $|t| = 2$ with no violation and with equality attained.]

A.2 The summed remainder, estimated under the integral

Lemma A.3. [\[proved\]](#) Put $T_1 = \sigma/N$ and, for $|t| \leq T_1$,

$$R_5(t) = \log \psi(t/\sigma) - \left(-\frac{t^2}{2} - \frac{i}{6}\alpha t^3 + \frac{1}{24}\beta t^4 \right), \quad \psi(u) = \mathbb{E} e^{iu(S-\mu)}.$$

Then

$$\int_{|t| \leq T_1} e^{-t^2/2} |R_5(t)| dt \leq 16 C_T \frac{\sum_{a \in B} a^5}{\sigma^5}.$$

Proof. $\log \psi(t/\sigma) = \sum_{a \in B} [K(iat/\sigma) - ip_a t/\sigma]$, and the phase of the a -th factor is $|at/\sigma| \leq NT_1/\sigma = 1$, so Lemma A.1 applies to every factor. Summing it, $|R_5(t)| \leq C_T \sum_a a^5 |t/\sigma|^5$, and $\int_{\mathbb{R}} e^{-t^2/2} |t|^5 dt = 16$. $\square$ $\square$

Remark A.4 (why the estimate is made under the integral, and not uniformly). The natural move is to bound $|R_5|$ by its supremum over the whole of R1. That bound cannot work: measured, $\sup_{R1} |R_5|$ is of order 10^2 and *does not decrease with k* , because at the outer edge of R1 the cubic term of the expansion is itself of order 10^2 . The integral is dominated by $|t| \lesssim 1$, where the expansion is excellent, and the error at the edge is never integrated against anything: the $|\hat{G}|$ -weighted remainder is $9 \cdot 10^{-4}$ at $k = 100$ and $9 \cdot 10^{-5}$ at $k = 450$, decaying like $k^{-3/2}$ — *faster* than the leading terms of the budget it is a remainder for. [\[experimentally confirmed; r1remainder_r134, r1weighted_r134\]](#) A general form worth stating once: **a supremum taken over a region where the integrand is already negligible charges the whole region at its worst point. When an estimate is going into an integral, make it under the integral.**

Remark A.5 (the threshold, and where it comes from). Cutting at $T_1 = \sigma/N$ rather than at the LCLT-natural radius is what makes every phase at most 1. Beyond T_1 no Taylor bound is needed: on $T_1 \leq |t| \leq \pi\sigma$ the quadratic estimate of Remark 5.15 gives $|\psi(t/\sigma)| \leq \exp(-ct_{\min}\sigma^2/N^2)$ with c absolute, and $\sigma^2/N^2 \rightarrow \infty$ under (H), so that contribution is smaller than any power of k once k exceeds a threshold k_0 depending on the profile through σ/N alone. *This is a derived threshold and not a hypothesis:* (H) supplies it for every admissible profile. For the odd numbers and the odd primes, $\sigma^2/N^2 \asymp k$, and the threshold is of order a few thousand — immaterial to statements about $k \rightarrow \infty$, and worth printing rather than hiding.

A.3 The inversion

Lemma A.6 (Hermite inversion). *[proved] For every $j \geq 0$ and every real z , $\frac{1}{2\pi} \int_{\mathbb{R}} e^{-t^2/2} (it)^j e^{-izt} dt = \text{He}_j(z) \varphi(z)$, where φ is the standard normal density and He_j the j -th Hermite polynomial.*

Proof. $\frac{1}{2\pi} \int e^{-t^2/2} e^{-izt} dt = \varphi(z)$; differentiating j times in z brings $(-it)^j$ inside, and $\varphi^{(j)} = (-1)^j \text{He}_j \varphi$. $\square$

Proposition A.7 (the expansion on R1). *[proved; independently read — this proposition, the T_* construction and all five constants of ρ in r171, the lemmas it uses line by line in r162; see the head of this appendix for the grain of each part] With the notation above, on the extraction window $|m - \mu| \leq W$ with $Z = W/\sigma$, and for all k above the threshold of Remark A.5,*

$$\mathbb{P}[S = m] = \frac{\varphi(z)}{\sigma} (1 + \theta), \quad |\theta| \leq \varepsilon^*(Z),$$

the Gaussian-relative form, which is what Remark A.8 and both consumer theorems actually use. The Hermite terms are then not part of the statement but a description of where the budget goes: the leading behaviour of θ is $\frac{\alpha}{6} \text{He}_3(z) + \frac{\beta}{24} \text{He}_4(z) + \frac{\alpha^2}{72} \text{He}_6(z)$, and ε^* bounds that leading part together with every higher term of the exponentiation. Explicitly, writing $S_5 = \sum_{a \in B} a^5$,

$$\varepsilon^*(Z) = \frac{|\alpha|}{6} \min(3Z, 2 + Z^3) + \frac{|\beta|}{24} (3 + 6Z^2 + Z^4) + \frac{\alpha^2}{72} (15 + 45Z^2 + 15Z^4 + Z^6) + \sqrt{2\pi} e^{Z^2/2} \rho,$$

$$\begin{aligned} \rho &= \frac{e(2 + \pi)}{\pi} \cdot \frac{S_5}{\sigma^5} + \frac{1}{3\pi} |\alpha| |\beta| + \frac{105}{1152\sqrt{2\pi}} \beta^2 + \frac{32e}{27\pi} |\alpha|^3 + \frac{385e}{768\sqrt{2\pi}} |\beta|^3 + \rho_\infty \\ &= 4.4488 \frac{S_5}{\sigma^5} + 0.1061 |\alpha| |\beta| + 0.03636 \beta^2 + 1.0255 |\alpha|^3 + 0.5436 |\beta|^3 + \rho_\infty, \end{aligned}$$

where ρ is an absolute bound on the residual at the scale of $\sigma\mathbb{P}$ and ρ_∞ is the contribution of $|t| \geq T_*$, smaller than any power of k above the threshold of Remark A.5 — the same bucket, and the same device, that remark already prints for the region beyond T_1 .

Two features of this display are there because a first version of it was wrong, and both are worth naming. First, the conversion from an absolute residual to a relative one costs the factor $e^{Z^2/2}$, and that factor is not a constant: the relative size of an absolute residual ρ is $\rho\sigma/\varphi(z)$, whose supremum on $|z| \leq Z$ is $\rho\sigma\sqrt{2\pi} e^{Z^2/2}$. Under (H), $Z \rightarrow 0$ and the factor tends to 1, so nothing downstream moves — but an earlier form of this proposition carried the constants $2/\varphi(0)$ and a bare $16C_{\text{T}}S_5/\sigma^5$ in its place, and those are valid only for $Z \leq \sqrt{2\log 2} = 1.1774$ and $Z \leq \sqrt{\log(\pi/2)} = 0.6720$ respectively. Both restrictions hold eventually and neither was stated. A bound that becomes true only in the limit, printed without its condition, is a bound whose reader cannot tell which one they have. Second, ρ carries no $O(\cdot)$: an O is not a constant, and the banner of this appendix admits no exceptions. *[the five constants of ρ are proved below; the four Gaussian moments, the cube split, the quadratic estimate and the two validity thresholds quoted here are additionally verified numerically, and the assembled ε^* is checked against the true relative error by exact dynamic programming at sixteen (profile, k , x) points with no failure. r11ift_r164.]*

Proof. Invert ψ over $[-\pi\sigma, \pi\sigma]$ in the variable t , so that $\mathbb{P}[S = m] = \frac{1}{2\pi\sigma} \int \psi(t/\sigma) e^{-izt} dt$. Within T_1 every phase $|at/\sigma|$ is at most 1, so with

$$X(t) = -\frac{i}{6}\alpha t^3 + \frac{1}{24}\beta t^4, \quad Q(t) = 1 + X(t) - \frac{1}{72}\alpha^2 t^6,$$

Lemma A.3 gives $\psi(t/\sigma) = e^{-t^2/2}e^{X(t)+R_5(t)}$, and Q is what Lemma A.6 sends to the three Hermite terms, since $(it)^3 = -it^3$, $(it)^4 = t^4$ and $(it)^6 = -t^6$. Those three are bounded on $|z| \leq Z$ by $\sup |\text{He}_3| \leq \min(3Z, 2 + Z^3)$, $\sup |\text{He}_4| \leq 3 + 6Z^2 + Z^4$ and $\sup |\text{He}_6| \leq 15 + 45Z^2 + 15Z^4 + Z^6$, each from the coefficients of the polynomial and, for He_3 , from $\text{He}_3(0) = 0$.

Where the Taylor bounds are cut, and why not at T_1 . Put

$$T_* = \min\left\{T_1, (3/|\alpha|)^{1/3}, (12/|\beta|)^{1/4}, (\sigma^5/(C_T S_5))^{1/5}\right\},$$

the largest radius on which $|X| \leq 1$ and $|R_5| \leq 1$ are both guaranteed by bounds already proved. **The natural move — keep T_1 and secure $|X| \leq 1$ by a threshold in k — is not available, and fails in the direction opposite to the one a threshold repairs:** $|X(T_1)| \leq |\alpha|T_1^3/6 + |\beta|T_1^4/24$ and $|\alpha|T_1^3 \asymp k^{-1/2} \cdot k^{3/2} = k$, so that condition holds for *small* k and is lost for large. This is Remark A.4's lesson one subsection later and in a sharper form: on $[T_*, T_1]$ the integrand is already negligible, and a Taylor bound evaluated at the far edge charges the whole window at a point that never contributes. Under (H), $|\alpha| \asymp k^{-1/2}$ gives $T_* \asymp k^{1/6} \rightarrow \infty$. [measured: $|X(T_1)|$ runs 0.13, 0.26, 0.51, 1.02 at $k = 32, 64, 128, 256$ for the odd numbers — doubling as k doubles — while $T_*/k^{1/6}$ stays within a factor 1.4 over the same range. `r1lift_r164`.]

Inside T_* . There $|X| \leq 1$ and $|R_5| \leq 1$, so $|e^X| \leq e$ and $|e^{R_5} - 1| \leq 2|R_5|$, whence

$$|e^{X+R_5} - Q| \leq \underbrace{2e|R_5|}_{\text{Lemma A.3}} + \underbrace{\frac{e}{6}|X|^3}_{|e^x-1-x-\frac{x^2}{2}| \leq \frac{|x|^3 e^{|x|}}{6}} + \underbrace{\frac{|\alpha||\beta|}{144}|t|^7 + \frac{\beta^2}{1152}t^8}_{\text{the unkept part of } X^2/2},$$

the last bracket because $X^2/2 = -\frac{\alpha^2}{72}t^6 - \frac{i}{144}\alpha\beta t^7 + \frac{\beta^2}{1152}t^8$ and only its first term is in Q . Extend every integral to $\mathbb{R}$, which dominates, and use $\int e^{-t^2/2}|t|^5 dt = 16$, $\int |t|^7 = 96$, $\int t^8 = 105\sqrt{2\pi}$, $\int |t|^9 = 768$ and $\int t^{12} = 10395\sqrt{2\pi}$ — the odd ones being $2 \cdot 2^m m!$ and the even ones $(2m-1)!!\sqrt{2\pi}$ — together with $(a+b)^3 \leq 4(a^3+b^3)$ applied to $|X|^3$. Division by 2π gives the five constants of ρ exactly as displayed: $16eC_T/\pi = e(2+\pi)/\pi$; then $96/(144 \cdot 2\pi) = 1/(3\pi)$; then $105\sqrt{2\pi}/(1152 \cdot 2\pi) = 105/(1152\sqrt{2\pi})$; then $(e/6) \cdot 4 \cdot 768/(216 \cdot 2\pi) = 32e/(27\pi)$; then $(e/6) \cdot 4 \cdot 10395\sqrt{2\pi}/(13824 \cdot 2\pi) = 385e/(768\sqrt{2\pi})$.

Outside T_* . For $|v| \leq 1$, $1 - \cos v \geq (1 - \cos 1)v^2$ because $(1 - \cos v)/v^2$ decreases on $(0, \pi)$; hence $|q_a + p_a e^{iv}|^2 = 1 - 2p_a q_a(1 - \cos v) \leq e^{-2p_a q_a(1 - \cos 1)v^2}$, and multiplying over $a \in B$ at $v = at/\sigma$ gives $|\psi(t/\sigma)| \leq e^{-(1 - \cos 1)t^2}$ on the whole of $|t| \leq T_1$, with $1 - \cos 1 = 0.4597$. So the contribution of $T_* \leq |t| \leq T_1$ — from ψ and from the Gaussian–polynomial comparison alike — and the extension of the Hermite integrals from T_* to $\mathbb{R}$ are each bounded by $e^{-(1 - \cos 1)T_*^2}$ times a constant, which by $T_* \asymp k^{1/6}$ is $e^{-ck^{1/3}}$: smaller than any power of k . That, with the region beyond T_1 , is ρ_∞ , and it is exactly the class of term Remark A.5 already prints. $\square \square$

Remark A.8 (what the budget is, in the size that matters). Every term of ε^* is $O(k^{-1})$ or smaller: $|\beta| \asymp k^{-1}$, $\alpha^2 \asymp k^{-1}$, $|\alpha|Z \asymp k^{-1.8}$ since $Z = W/\sigma \rightarrow 0$ by (H), $S_5/\sigma^5 \asymp k^{-3/2}$, and the five terms of ρ fall in the classes $k^{-3/2}$, $k^{-3/2}$, k^{-2} , $k^{-3/2}$, k^{-3} . The conversion factor $\sqrt{2\pi}e^{Z^2/2}$ tends to $\sqrt{2\pi}$ because $Z \rightarrow 0$, so it is bounded and does not disturb the rates. Hence $\varepsilon^* \rightarrow 0$, which is all Theorems 4.1 and 5.1 require of it. *The price of stating the conversion honestly is paid at small k , where Z is not small:* at the sixteen test points Z reaches 2.86 and $e^{Z^2/2}$ reaches 60, so the assembled bound is looser than the earlier, conditionally valid form by one to two orders of magnitude. It is loose in the region where the theorems say nothing, and it is correct everywhere. [experimentally confirmed: ε^* bounds the true relative error of the Gaussian, computed by exact dynamic programming over the window, at all sixteen test

points tried, with ratio 68 to 997; and $k\varepsilon^*$ decreases along $k = 100, 200, 300, 450$, so ε^* decays at least as fast as k^{-1} . `r1lift_r164`; the earlier, conditionally valid form is `r1step3_r136`, `r1hermite_r137`, where the ratio was 1.7 to 6.5. *The check covers the five explicit constants of ρ and not ρ_∞ , which is an asymptotic statement of the same kind as Remark A.5 and is not numerically small at the k reachable by exact enumeration.*]

Tool and computational resource disclosure

This section follows recommendation 01 for individual mathematicians of the *Leiden Declaration on Artificial Intelligence and Mathematics* (June 2026, DOI:10.5281/zenodo.20302944), which the author has signed.

Tools. This work was carried out by the author working with AI language models (Anthropic’s Claude, models `opus-5` and `fable-5`) as tools, under the author’s direction. They were used to fill in Lean proofs, to run and tabulate the numerical experiments, and to draft sections which the author then checked. Formal verification uses Lean 4 (toolchain `leanprover/lean4:v4.32.2`) with Mathlib, and the kernel replay uses `lean4checker`. Type-setting uses Xe $\text{\LaTeX}$; the numerical work uses Python 3 with no libraries beyond the standard one.

Computational resources. All of it runs on a single personal computer. No cluster, no accelerator, and no compute grant were used; the longest single computation in this Part is a kernel replay of the whole Lean development, which takes about three minutes. This is recorded because recommendation 10 of the Declaration asks whether smaller systems suffice, and here the honest answer is mixed: the verification and the experiments are small enough for anyone’s laptop, while the language models are commercial, proprietary and energy-intensive, and the author found no smaller system adequate for that role.

Responsibility and attribution. The choice of problem, the design decisions, the completeness of the citations, and the responsibility for every claim here are the author’s alone; credit belongs to no automated system. Where the author has been unable to attribute an idea satisfactorily, the text says so.

Why the apparatus exists. Work produced this way cannot reasonably be taken on the author’s word, and the apparatus described under Data availability exists for that reason: every settled theorem is machine-checked in Lean and replayed through the kernel by an independent checker that must first reject three deliberately corrupted modules; every number quoted in the text is produced by a committed script that writes a log beside itself; every theorem, proposition and lemma declares its status where it is stated; and a suite of mechanical checks enforces all of this before each commit. The repository records, in full, the mistakes this process has made — which is the part the author would argue for hardest, since those failures are precisely the ones no machine check can see.

Data availability

Every number in this paper is produced by a script that writes a log beside itself, as in Parts I and II and the transfer-function strand. The scripts, the logs, the Lean development and the checks that enforce all of this are public at <https://github.com/mathlab0911/arithmetic-landscapes>.

References

- [1] OpenAI, *Ten advances in mathematics and theoretical computer science*, 2026, with Lean 4 certificates at github.com/openai/ten-proofs. Cited here only as a precedent for the genre; no mathematical result of that work is used.
- [2] R. Giuliano and M. Weber, *Approximate local limit theorems with effective rate and application to random walks in random scenery*. Bernoulli **23** (2017), 3268–3310; arXiv:1412.3980. Cited only for the definition and the applicability condition quoted in Remark 5.23.
- [3] D. S. Kubert and S. Lang, *Modular units*. Grundlehren der mathematischen Wissenschaften **244**, Springer, 1981; and D. S. Kubert, *The universal ordinary distribution*. Bull. Soc. Math. France **107** (1979), 179–202. Cited for the distribution relation $\sum_{k < v} \log |2 \sin \pi(t+k/v)| = \log |2 \sin \pi vt|$, of which Lemma 3.1 is the half-shifted form.
- [4] K. Amauchi, *Arithmetic landscapes I: the gap series*. Manuscript, 2026.
- [5] K. Amauchi, *Arithmetic landscapes II: asymptotic flatness of subset-sum landscapes of primes*. Manuscript, 2026.